

DEVELOPER TRAINING FOR APACHE SPARK

Chapter 1

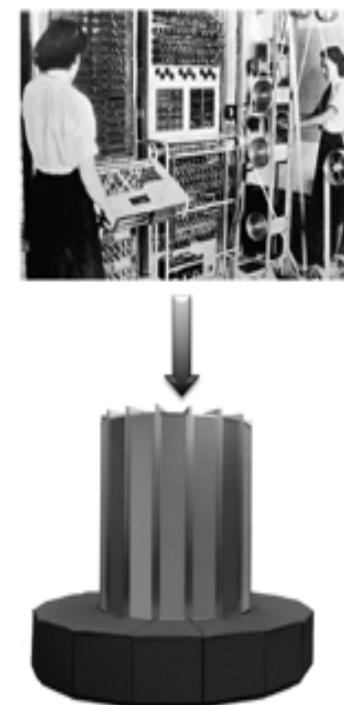
WHY SPARK?

Why Spark?

- In this chapter you will learn
- What problems exist with traditional large-scale computing systems
- How Spark addresses those issues
- Some typical big data questions Spark can be used to answer

Traditional Large-Scale Computation

- Traditionally, computation has been processor bound
 - Relatively small amounts of data
 - Lots of complex processing
- The early solution: bigger computers
 - Faster processor, more memory
 - But even this couldn't keep up

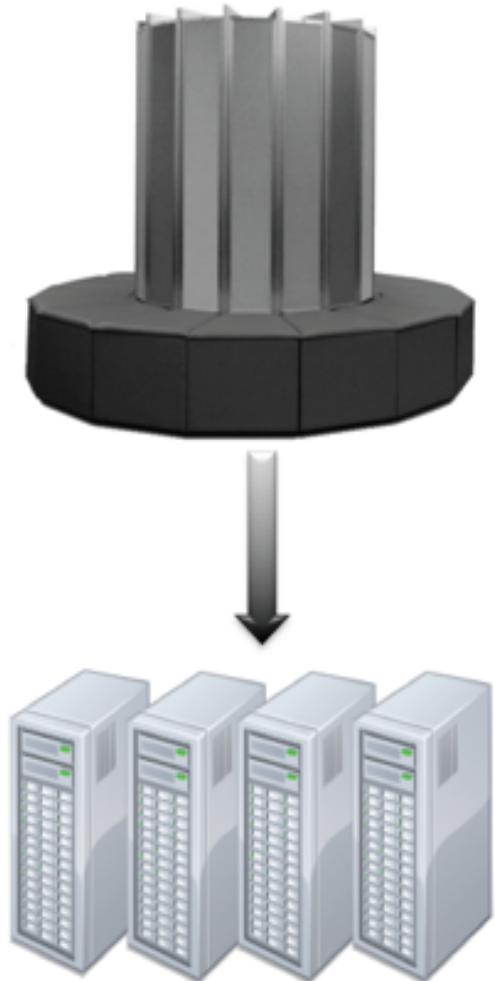


Distributed Systems

- The better solution: more computers
 - Distributed systems - use multiple machines for a single job

“In pioneer days they used oxen for heavy pulling, and when one ox couldn’t budge a log, we didn’t try to grow a larger ox. We shouldn’t be trying for bigger computers, but for more systems of computers.”

-Grace Hopper

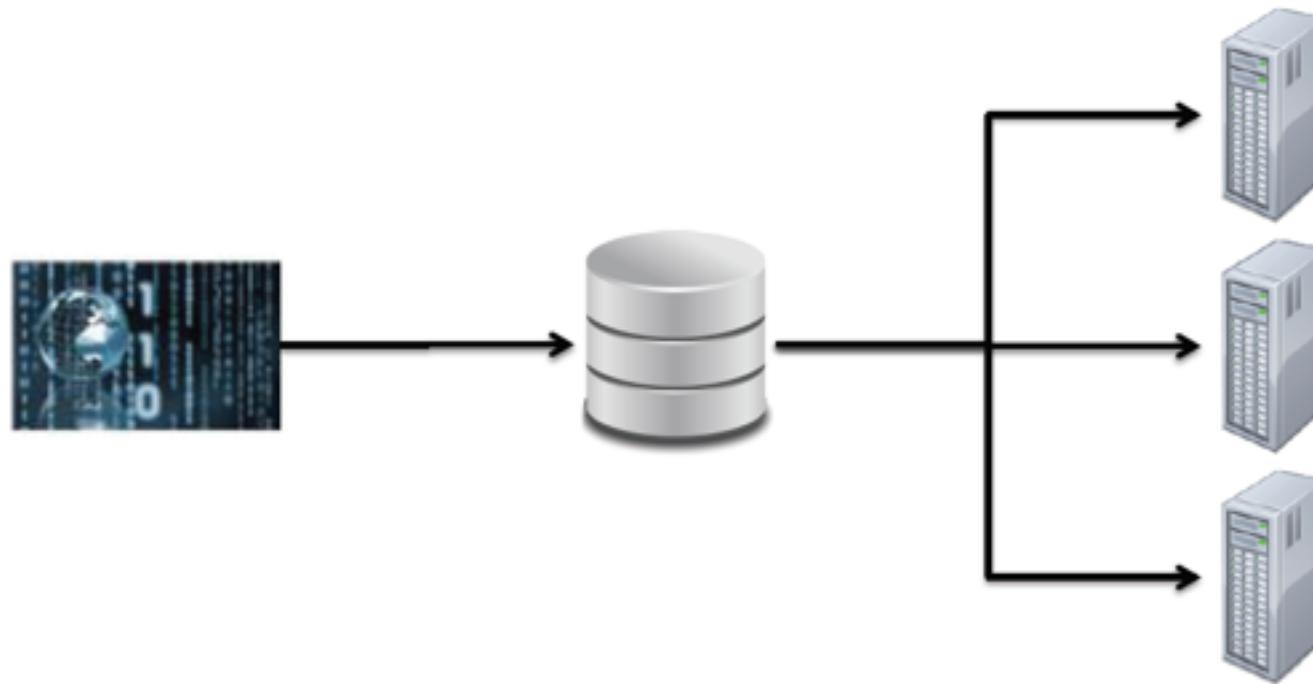


Distributed Systems: Challenges

- Challenges with distributed systems
 - Programming complexity
 - Keeping data and processes in sync
 - Finite bandwidth
 - Partial failures

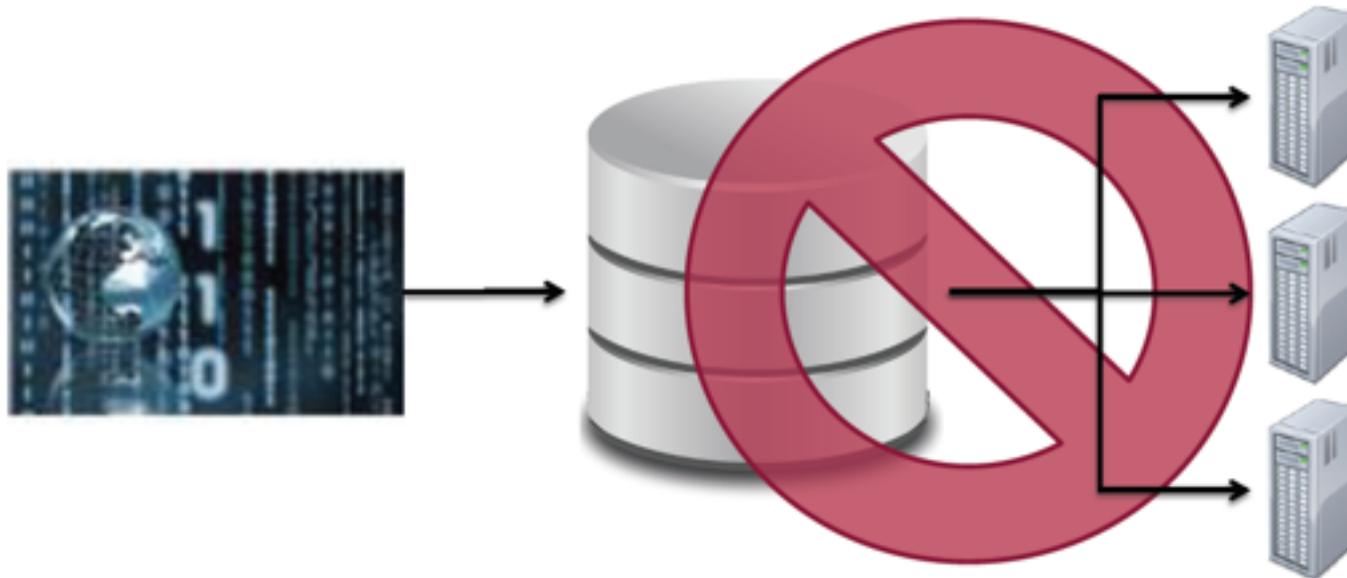
Distributed Systems: The Data Bottleneck

- Traditionally, data is stored in a central location
- Data is copied to processors at runtime
- Fine for limited amounts of data



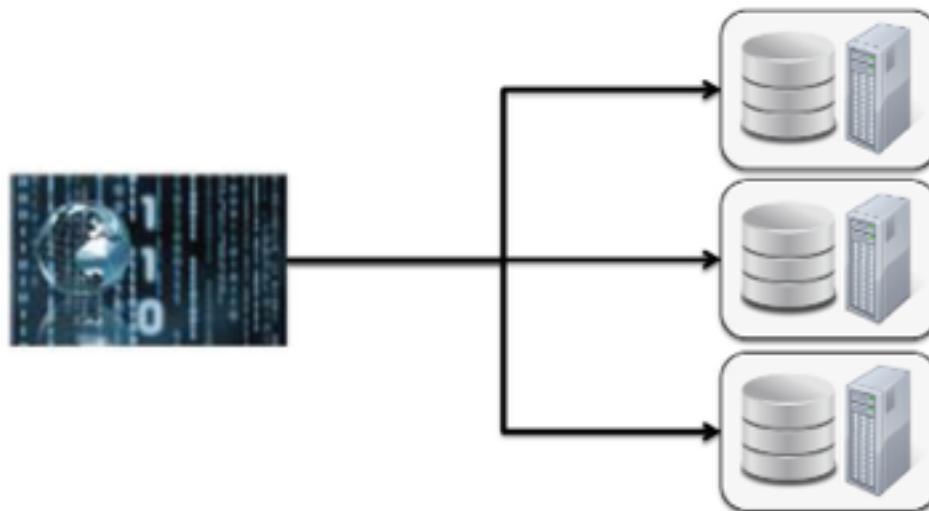
Distributed Systems: The Data Bottleneck

- Modern systems have much more data
 - terabytes+ a day
 - petabytes+ total
- We need a new approach...

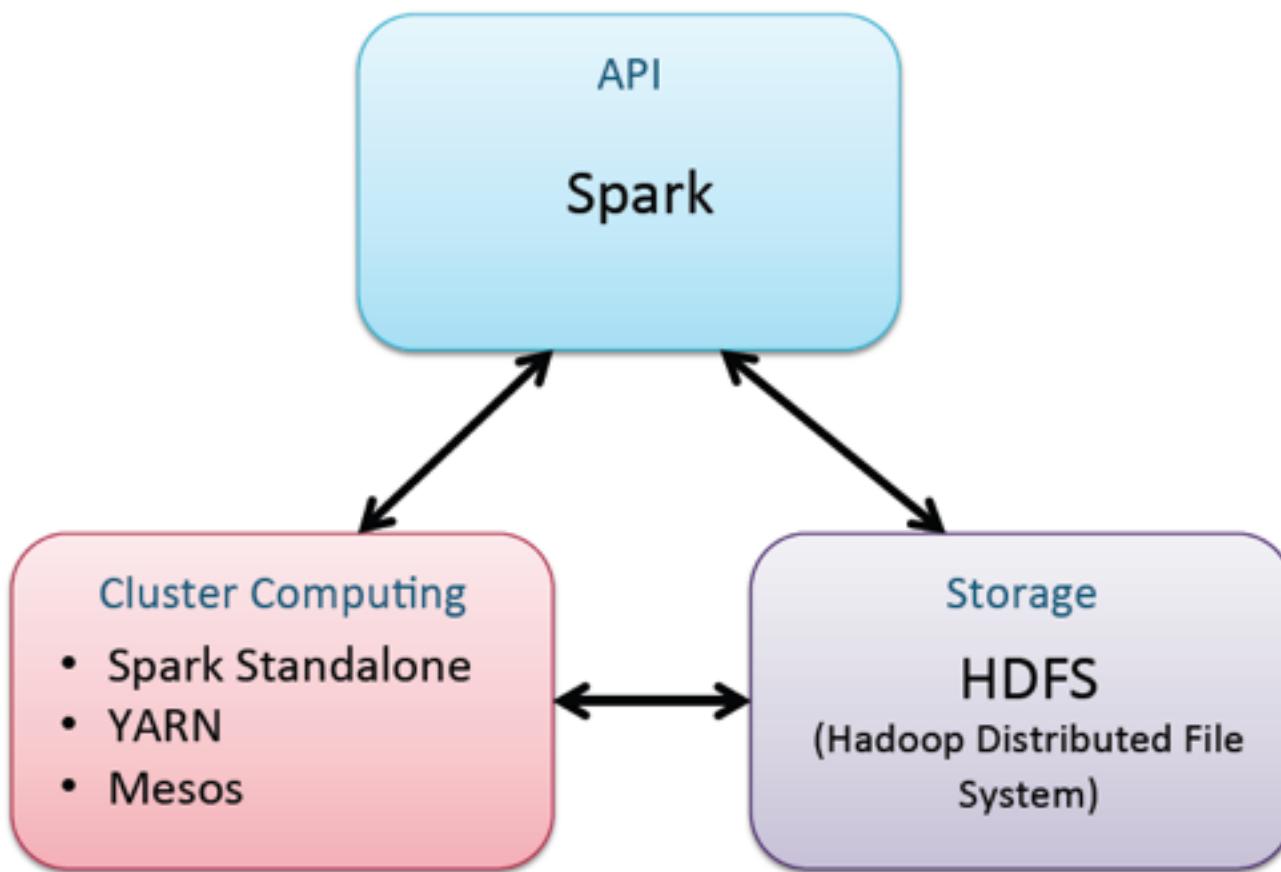


Big Data Processing

- Hadoop introduced a radical new approach based on two key concepts
 - Distribute data when the data is stored
 - Run computation where the data is
- Spark takes this new approach to the next level
 - Data is distributed in memory



Distributed Processing with the Spark Framework

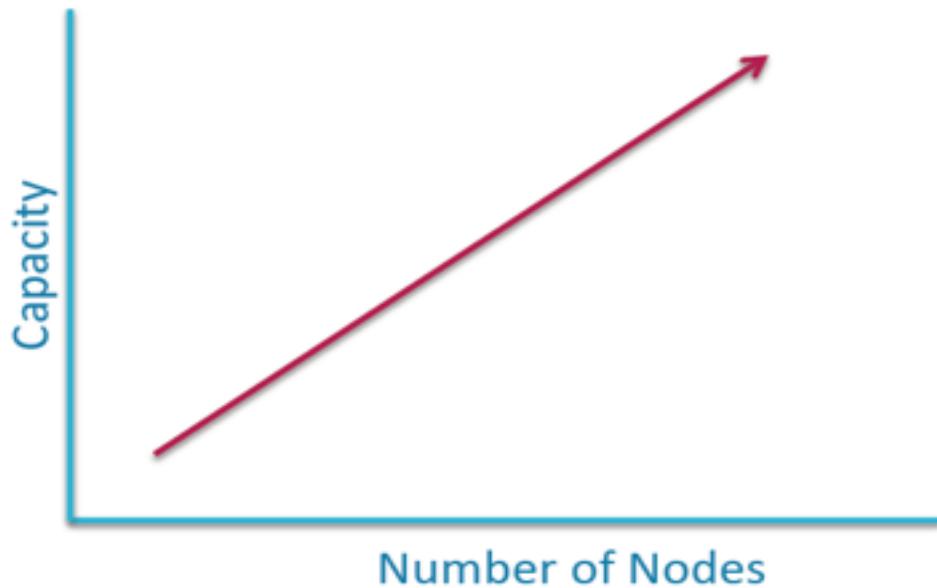


Advantages of Spark

- High level programming framework
 - Programmers can focus on logic, not plumbing
- Cluster computing
 - Application processes are distributed across a cluster of worker nodes
 - Managed by a single “master”
 - Scalable and fault tolerant
- Distributed storage
 - Data is distributed when it is stored
 - Replicated for efficiency and fault tolerance
 - “Bring the computation to the data”
- Data in memory
 - Configurable caching for efficient iteration

Scalability

- Increasing load results in a graceful decline in performance
 - Not failure of the system
- Adding nodes adds capacity proportionally



Fault Tolerance

- Node failure is inevitable
- What happens?
 - System continues to function
 - Master re-assigns tasks to a different node
 - Data replication= no loss of data
 - Nodes which recover rejoin the cluster automatically

Who Uses Spark?

- **Yahoo!**
 - Personalization and ad analytics
- **Conviva**
 - Real-time video stream optimization
- **Technicolor**
 - Real-time analytics for telco clients
- **Ooyala**
 - Cross-device personalized video experience
- **Plus...**
 - Intel, Groupon, TrendMicro, Autodesk, Nokia, Shopify, ClearStory, Technicolor, and many more...

Common Spark Use Cases

- Extract/Transform/Load (ETL)
- Text mining
- Index building
- Graph creation and analysis
- Pattern recognition
- Collaborative filtering
- Prediction models
- Sentiment analysis
- Risk assessment

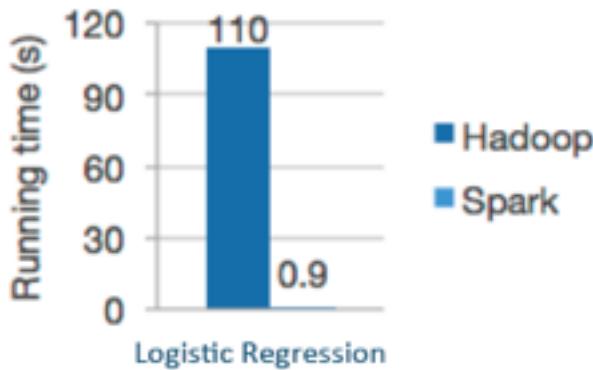
Benefits of Spark

- Previously impossible or impractical analysis
- Lower cost
- Less time
- Greater flexibility
- Near-linear scalability



Spark v. Hadoop MapReduce

- Spark takes the concepts of MapReduce to the next level
- Higher level API = faster, easier development
 - Low latency = near real-time processing
 - In-memory data storage = up to 100x performance improvement



```
sc.textFile(file) \
    .flatMap(lambda s: s.split()) \
    .map(lambda w: (w,1)) \
    .reduceByKey(lambda v1,v2: v1+v2)
    .saveAsTextFile(output)
```



```
public class WordCount {  
    public static void main(String[] args) throws  
        JobException {  
        Job job = new Job();  
        job.setMapperClass(WordMapper.class);  
        job.setReducerClass(WordReducer.class);  
        job.setInputFormat(new TextInputFormat());  
        job.setOutputFormat(new TextOutputFormat());  
        job.setMapOutputKeyClass(Text.class);  
        job.setMapOutputValueClass(IntWritable.class);  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
        boolean success = job.waitForCompletion(true);  
        System.exit(success ? 0 : 1);  
    }  
  
    public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable> {  
        public void map(LongWritable key, Text value,  
            Context context) throws IOException, InterruptedException {  
            String line = value.toString();  
            for (String word : line.split("\\W+")) {  
                if (word.length() > 0)  
                    context.write(new Text(word), new IntWritable(1));  
            }  
        }  
    }  
  
    public class WordReducer extends Reducer<Text, IntWritable, Text, IntWritable> {  
        public void reduce(Text key, Iterable<IntWritable>  
            values, Context context) throws IOException, InterruptedException {  
            int wordCount = 0;  
            for (IntWritable value : values) {  
                wordCount += value.get();  
            }  
            context.write(key, new IntWritable(wordCount));  
        }  
    }  
}
```



Chapter 2

SPARK BASICS

What is Apache Spark?

- Apache Spark is a fast and general engine for large scale data processing
- Written in Scala
 - Functional programming language that runs in a JVM
- Spark Shell
 - Interactive- for learning or data exploration
 - Python or Scala
- Spark Applications
 - For large scale data processing
 - Python, Scala , or Java



Spark Shell

- The Spark Shell provides interactive data exploration (REPL)
 - Writing Spark applications without the shell will be covered later
 -
 - Python Shell: `pyspark` Scala Shell: `spark-shell`

```
$ pyspark

Welcome to

$$\begin{array}{c} \diagup \diagdown \\ \diagdown \diagup \end{array} \begin{array}{c} \diagup \diagdown \\ \diagdown \diagup \end{array}$$
 version 0.9.1

Using Python version 2.6.6 (r266:84292, Jan
22 2014 09:42:36)
Spark context available as sc.

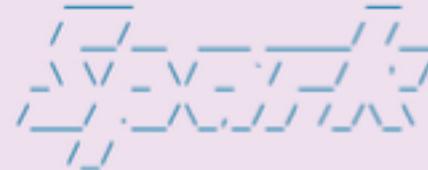
>>>
```

REPL: Read/Evaluate/Print Loop

Scala Shell: spark-shell

```
$ spark-shell
```

Welcome to



version 0.9.1

```
Using Scala version 2.10.3 (Java HotSpot(TM)  
64-Bit Server VM, Java 1.7.0_51)  
Created spark context..  
Spark context available as sc.
```

scala>

Spark Context

- Every Spark application requires a Spark Context
 - The main entry point to the Spark API
- Spark Shell provides a preconfigured Spark Context called sc

```
Using Python version 2.6.6 (r266:84292, Jan 22 2014 09:42:36)
Spark context available as sc.
```

```
>>> sc.appName
u'PySparkShell'
```

```
Using Scala version 2.10.3 (Java HotSpot(TM) 64-Bit Server VM, Java
1.7.0_51)
```

```
Created spark context..
```

```
Spark context available as sc.
```

```
scala> sc.appName
res0: String = Spark shell
```

RDD (Resilient Distributed Dataset)

- **RDD (Resilient Distributed Dataset)**
 - Resilient - if data in memory is lost, it can be recreated
 - Distributed - stored in memory across the cluster
 - Dataset - initial data can come from a file or be created programmatically
- RDDs are the fundamental unit of data in Spark
- Most Spark programming consists of performing operations on RDDs

Creating an RDD

- Three ways to create an RDD
 - From a file or set of files
 - From data in memory
 - From another RDD

File-Based RDDs

- For file-based RDDS, use `Spark Context.textFile`
 - Accepts a single file, a wildcard list of files, or a comma-separated list of files
 - Examples
 - `sc.textFile("myfile.txt")`
 - `sc.textFile("mydata/*.log")`
 - `sc.textFile("myfile1.txt,myfile2.txt")`
 - Each line in the file(s) is a separate record in the RDD
- Files are referenced by absolute or relative URI
 - Absolute URI: `file:/home/training/myfile.txt`
 - Relative URI (uses default file system): `myfile.txt`

Example: A File-based RDD

```
> mydata = sc.textFile("purplecow.txt")
...
14/01/29 06:20:37 INFO storage.MemoryStore:
  Block broadcast_0 stored as values to
  memory (estimated size 151.4 KB, free 296.8
  MB)

> mydata.count()
...
14/01/29 06:27:37 INFO spark.SparkContext: Job
finished: take at <stdin>:1, took
0.160482078 s
```

4

File: purplecow.txt

I've never seen a purple cow.
I never hope to see one;
But I can tell you, anyhow,
I'd rather see than be one.

RDD: mydata

I've never seen a purple cow.
I never hope to see one;
But I can tell you, anyhow,
I'd rather see than be one.

RDD Operations

- Two types of RDD operations
 - Actions - return values
 - Transformations- define a new RDD based on the current one(s)
- Pop quiz:
 - Which type of operation is count()?

RDD Operations: Actions

- Some common actions
 - `count()` - return the number of elements
 - `take(n)` - return an array of the first n elements
 - `collect()`- return an array of all elements
 - `saveAsTextFile(file)` - save to text file(s)

```
> mydata =  
  sc.textFile("purplecow.txt")  
  
> mydata.count()  
4  
  
> for line in mydata.take(2):  
    print line  
I've never seen a purple cow.  
I never hope to see one;
```

```
> val mydata |=  
  sc.textFile("purplecow.txt")  
  
> mydata.count  
4  
  
> for (line <- mydata.take(2))  
    println(line)  
I've never seen a purple cow.  
I never hope to see one;
```

RDD Operations: Transformations

- Transformations create a new RDD from an existing one
- RDDs are immutable
 - Data in an RDD is never changed
 - Transform in sequence to modify the data as needed
- Some common transformations
 - `map(function)` - creates a new RDD by performing a function on each record in the base RDD
 - `filter(function)` - creates a new RDD by including or excluding each record in the base RDD according to a boolean function

Example: map and filter Transformations

```
I've never seen a purple cow.  
I never hope to see one;  
But I can tell you, anyhow,  
I'd rather see than be one.
```

```
map(lambda line: line.upper())
```

```
map(line => line.toUpperCase())
```

```
I'VE NEVER SEEN A PURPLE COW.  
I NEVER HOPE TO SEE ONE;  
BUT I CAN TELL YOU, ANYHOW,  
I'D RATHER SEE THAN BE ONE.
```

```
filter(lambda line: line.startswith('I'))
```

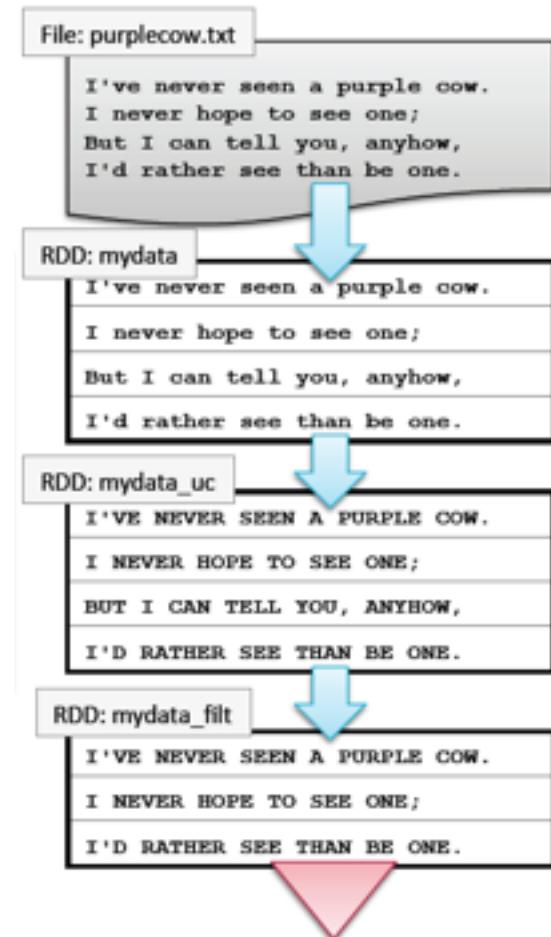
```
filter(line => line.startsWith("I"))
```

```
I'VE NEVER SEEN A PURPLE COW.  
I NEVER HOPE TO SEE ONE;  
I'D RATHER SEE THAN BE ONE.
```

Lazy Execution

- Data in RDDs is not processed until an action is performed

```
> mydata = sc.textFile("purplecow.txt")
> mydata_uc = mydata.map(lambda line:
   line.upper())
> mydata_filt = \
   mydata_uc.filter(lambda line: \
   line.startswith('I'))
> mydata_filt.count()
3
```



Chaining Transformations

- Transformations may be chained together

```
> mydata = sc.textFile("purplecow.txt")
> mydata_uc = mydata.map(lambda line: line.upper())
>         = mydata_uc.filter(lambda line: line.startswith('I'))
> mydata_filt.count()
3
```

is exactly equivalent to

```
> sc.textFile("purplecow.txt").map(lambda line: line.upper()) \
  .filter(lambda line: line.startswith('I')).count()
3
```

Functional Programming in Spark

- Spark depends heavily on the concepts of functional programming
 - Functions are the fundamental unit of programming
 - Functions have input and output only
 - No state or side effects
- Key concepts
 - Passing functions as input to other functions
 - Anonymous functions

Passing Functions as Parameters

- Many RDD operations take functions as parameters
- Pseudocode for the RDD map operation
 - Applies function fn to each record in the RDD

```
RDD {  
    map (fn (x)) {  
        foreach record in rdd  
        emit fn(record)  
    }  
}
```

Example: Passing Named Functions

```
> def toUpper(s: String): String =  
  { s.toUpperCase }  
> val mydata = sc.textFile("purplecow.txt")  
> mydata.map(toUpper).take(2)
```

Anonymous Functions

- Functions defined in-line without an identifier
 - Best for short, one-off functions
- Supported in many programming languages

Example: Passing Anonymous Functions

```
> mydata.map(line => line.toUpperCase()).take(2)
```

OR

```
> mydata.map(_.toUpperCase()).take(2)
```

Example: Java

Java 7

```
...
JavaRDD<String> lines = sc.textFile("file");
JavaRDD<String> lines_uc = lines.map(
    new MapFunction<String, String>() {
        public String call(String line) {
            return line.toUpperCase();
        }
    }
...
...
```

Java 8

```
...
JavaRDD<String> lines = sc.textFile("file");
JavaRDD<String> lines_uc = lines.map(
    line -> line.toUpperCase());
...
...
```

Now, please do the following three Hands-On Exercises

1. Viewing the Spark Documentation
2. Using the Spark Shell
3. Getting Started with RDDs

HANDS-ON EXERCISES

Chapter 3

WORKING WITH RDDS

RDDs

- RDDs can hold any type of element
 - Primitive types: integers, characters, booleans, etc.
 - Sequence types: strings, lists, arrays, tuples, dicts, etc. (including nested data types)
 - Scala/Java Objects (if serializable)
 - Mixed types
- Some types of RDDs have additional functionality
 - Pair RDDs
 - RDDs consisting of Key-Value pairs
 - Double RDDs
 - RDDs consisting of numeric data

Creating RDDs From Collections

- You can create RDDs from collections instead of files

```
> randomnumlist = \
    [random.uniform(0,10) for _ in xrange(10000)]
> randomrdd = sc.parallelize(randomnumlist)
> print "Mean: %f" % randomrdd.mean()
```

Some Other General RDD Operations

- Transformations
 - distinct - filter out duplicates
 - union - add all elements of two RDDs into a single new RDD
 - zip - pair the elements of one RDD with another
- Other RDD operations
 - first - return the first element of the RDD
 - for each - apply a function to each element in an RDD
 - top(n) - return the largest n elements using natural ordering
- Sampling operations
 - sample(percent) - create a new RDD with a sampling of elements
 - takeSample (percent) - return an array of sampled elements
- Double RDD operations
 - Statistical functions, e.g. mean, sum, variance, stdev

Pair RDDs

- Pair RDDs are a special form of RDD
 - Each element must be a key-value pair (a two-element tuple)
 - Keys and values can be any type
- Why?
 - Use with Map Reduce algorithms
 - Many additional functions are available for common data processing needs
 - e.g., sorting, joining, grouping, counting, etc.

Pair RDD
(key1,value1)
(key2,value2)
(key3,value3)
...

Creating Pair RDDs

- The first step in most workflows is to get the data into key/value form
 - What should the RDD should be keyed on?
 - What is the value?
- Commonly used functions to create Pair RDDs
 - map
 - flatMap
 - keyBy

Example: A Simple Pair RDD

```
> users = sc.textFile(file) \
    .map(lambda line: line.split('\t')) \
    .map(lambda fields: (fields[0],fields[1]))
```

```
> val users = sc.textFile(file) \
    .map(line => line.split('\t')) \
    .map(fields => (fields(0),fields(1)))
```

user001 Fred Flintstone
user090 Bugs Bunny
user111 Harry Potter
...



(user001,Fred Flintstone)
(user090,Bugs Bunny)
(user111,Harry Potter)
...

Example: Keying Web Logs by User ID

```
> sc.textFile(logfile) \
    .keyBy(lambda line: line.split(' ')[2])
```

```
> sc.textFile(logfile) \
    .keyBy(line => line.split(' ')(2))
```

User ID	
56.38.234.188 -	99788
56.38.234.188 -	99788
203.146.17.59 -	25254
...	



(99788, 56.38.234.188 - 99788 "GET /KBDOC-00157.html...")

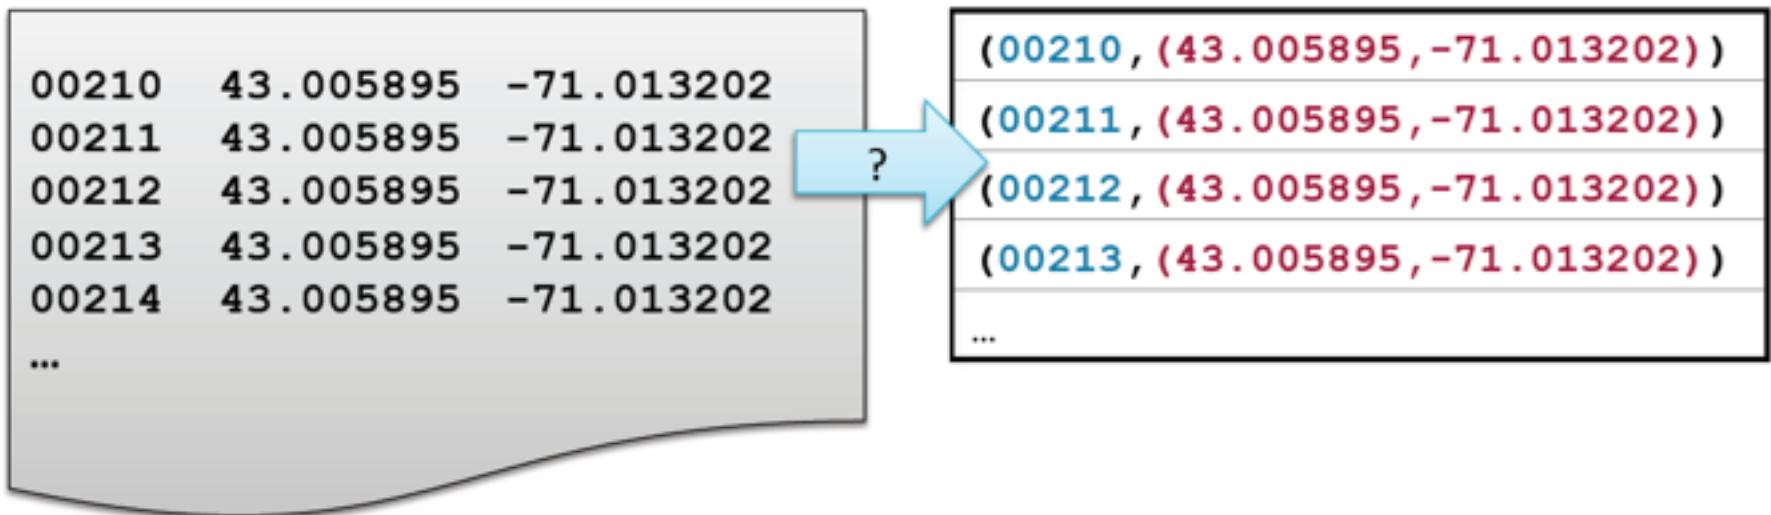
(99788, 56.38.234.188 - 99788 "GET /theme.css...")

(25254, 203.146.17.59 - 25254 "GET /KBDOC-00230.html...")

...

Pairs With Complex Values

- How would you do this?
 - Input: a list of postal codes with latitude and longitude
 - Output: postal code (key) and lat/long pair (value)



Pairs With Complex Values

```
> sc.textFile(file) \
    .map(lambda line: line.split()) \
    .map(lambda fields: (fields[0], (fields[1], fields[2])))
```

00210 43.005895 -71.013202
00211 43.005895 -71.013202
00212 43.005895 -71.013202
00213 43.005895 -71.013202
00214 43.005895 -71.013202
...



(00210, (43.005895, -71.013202))
(00211, (43.005895, -71.013202))
(00212, (43.005895, -71.013202))
(00213, (43.005895, -71.013202))
...

Mapping Single Rows to Multiple Pairs

- How would you do this?
 - Input: order numbers with a list of SKUs in the order
 - Output: order (key) and sku (value)



Mapping Single Rows to Multiple Pairs

- Hint: map alone won't work

```
00001  sku010:sku933:sku022  
00002  sku912:sku331  
00003  sku888:sku022:sku010:sku594  
00004  sku411
```



```
(00001, (sku010, sku933, sku022))  
(00002, (sku912, sku331))  
(00003, (sku888, sku022, sku010, sku594))  
(00004, (sku411))
```

Answer 2: Creating a Pair RDD

```
> sc.textFile(file) \
    .map(lambda line: line.split('\t')) \
    .flatMap(lambda fields: (fields[0], fields[1].split(':')))
```

00001 sku010:sku933:sku022

00002 sku912:sku331

00003 sku888:sku022:sku010:sku594

00004 sku411

(00001,sku010:sku933:sku022)

(00002,sku912:sku331)

(00003,sku888:sku022:sku010:sku594)

(00004,sku411)

(00001,sku010)

(00001,sku933)

(00001,sku022)

(00002,sku912)

(00002,sku331)

(00003,sku888)

...

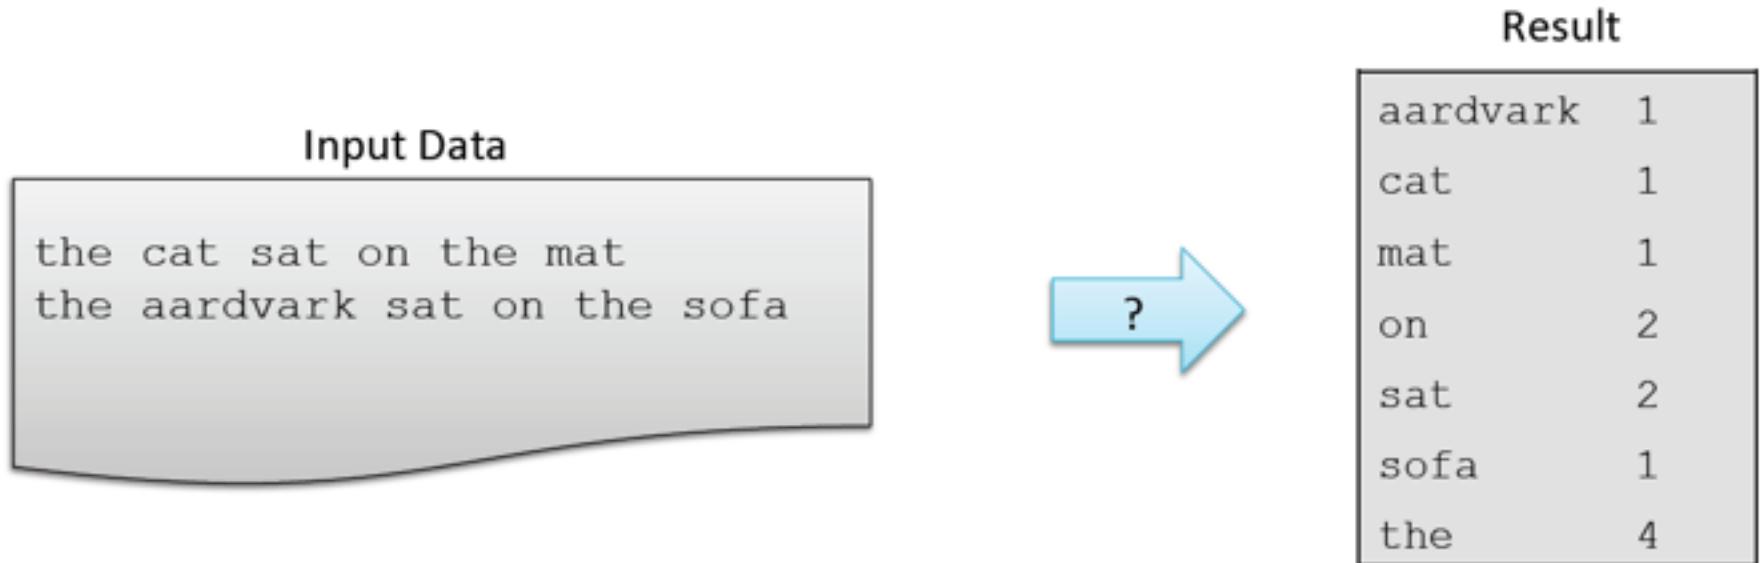
MapReduce

- MapReduce is a common programming model
 - Easily applicable to distributed processing of large data sets
- Hadoop MapReduce is the major implementation
 - Somewhat limited
 - Each job has one Map phase, one Reduce phase
 - Job output is saved to files
- Spark implements MapReduce with much greater flexibility
 - Map and Reduce functions can be interspersed
 - Results stored in memory
 - Operations can easily be chained

MapReduce in Spark

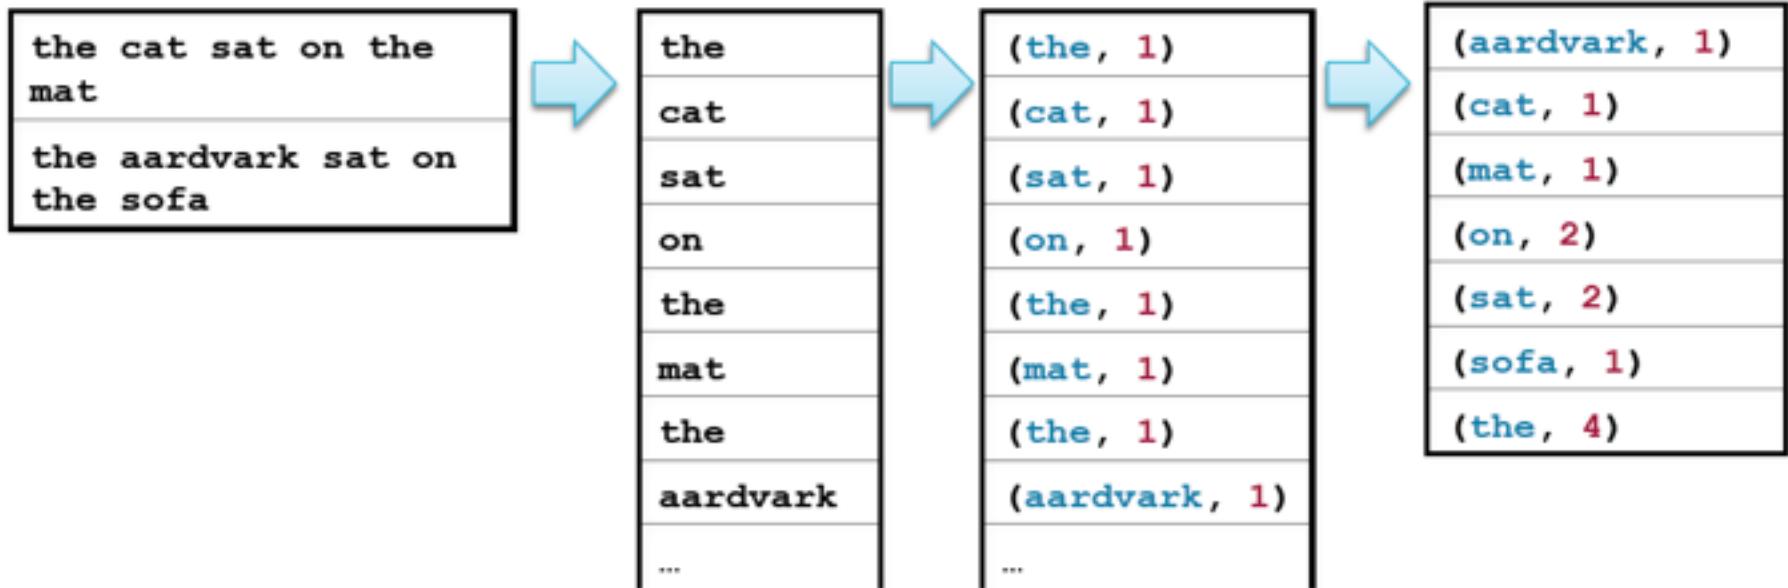
- MapReduce in Spark works on Pair RDDs
- Map
 - Operates on one record at a time
 - “Maps” each record to one or more new records
 - map and flatMap
- Reduce
 - Works on Map output
 - Consolidates multiple records
 - reduceByKey

MapReduce Example: Word Count



Example: Word Count

```
> counts = sc.textFile(file) \
    .flatMap(lambda line: line.split()) \
    .map(lambda word: (word,1)) \
    .reduceByKey(lambda v1,v2: v1+v2)
```

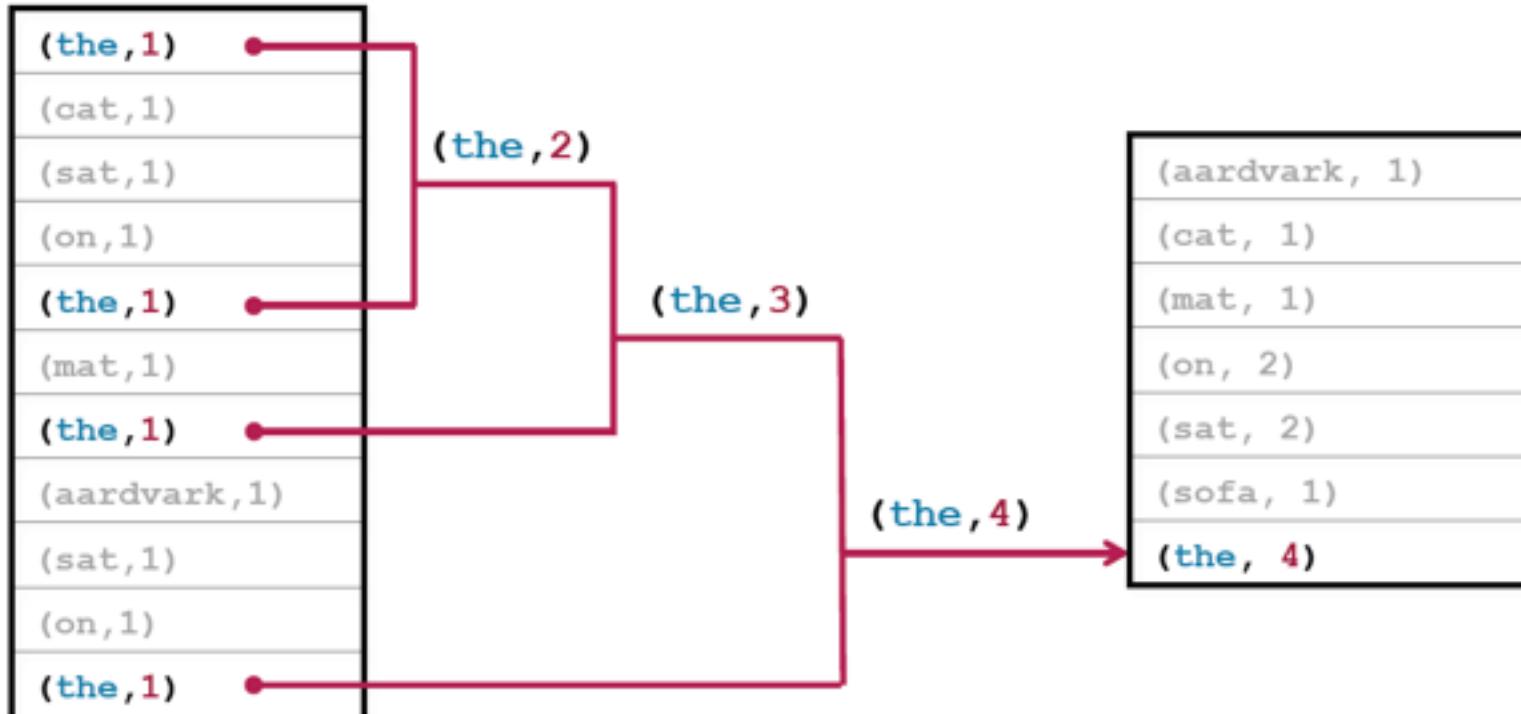


ReduceByKey

ReduceByKey functions must be

- Binary – combines values from two keys
- Commutative – $x+y = y+x$
- Associative – $(x+y)+z = x+(y+z)$

```
> counts = sc.textFile(file) \
    .flatMap(lambda line: line.split()) \
    .map(lambda word: (word,1)) \
    .reduceByKey(lambda v1,v2: v1+v2)
```



Word Count Recap (the Scala Version)

```
> val counts = sc.textFile(file) \  
  .flatMap(line => line.split("\\W")) \  
  .map(word => (word,1)) \  
  .reduceByKey((v1,v2) => v1+v2)
```

OR

```
> val counts = sc.textFile(file) \  
  .flatMap(_.split("\\W")) \  
  .map(_ ,1)) \  
  .reduceByKey(_+_)
```

Why Do We Care About Counting Words?

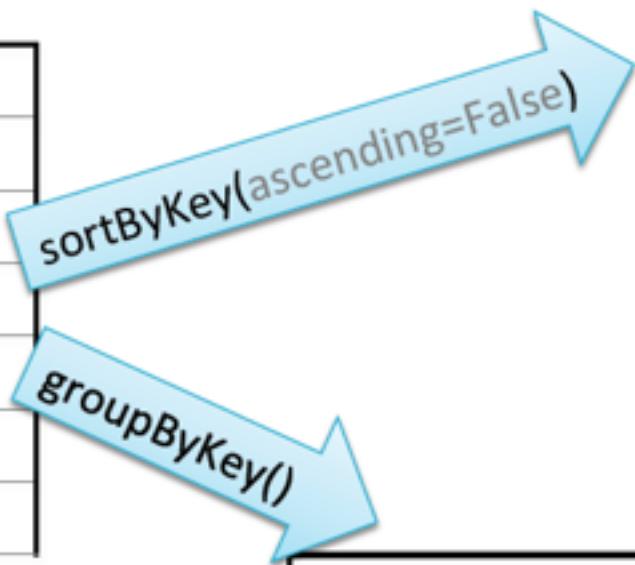
- Word count is challenging over massive amounts of data
 - Using a single compute node would be too time-consuming
 - Number of unique words could exceed available memory
- Statistics are often simple aggregate functions
 - Distributive in nature
 - e.g., max, min, sum, count
- MapReduce breaks complex tasks down into smaller elements which can be executed in parallel
- Many common tasks are very similar to word count
 - e.g., log file analysis

Pair RDD Operations

- In addition to map and reduce functions, Spark has several operations specific to Pair RDDs
- Examples
 - countByKey - return a map with the count of occurrences of each key
 - groupByKey - group all the values for each key in an RDD
 - sortByKey - sort in ascending or descending order
 - join - return an RDD containing all pairs with matching keys in two RDDs

Example: Pair RDD Operations

(00001, sku010)
(00001, sku933)
(00001, sku022)
(00002, sku912)
(00002, sku331)
(00003, sku888)
...

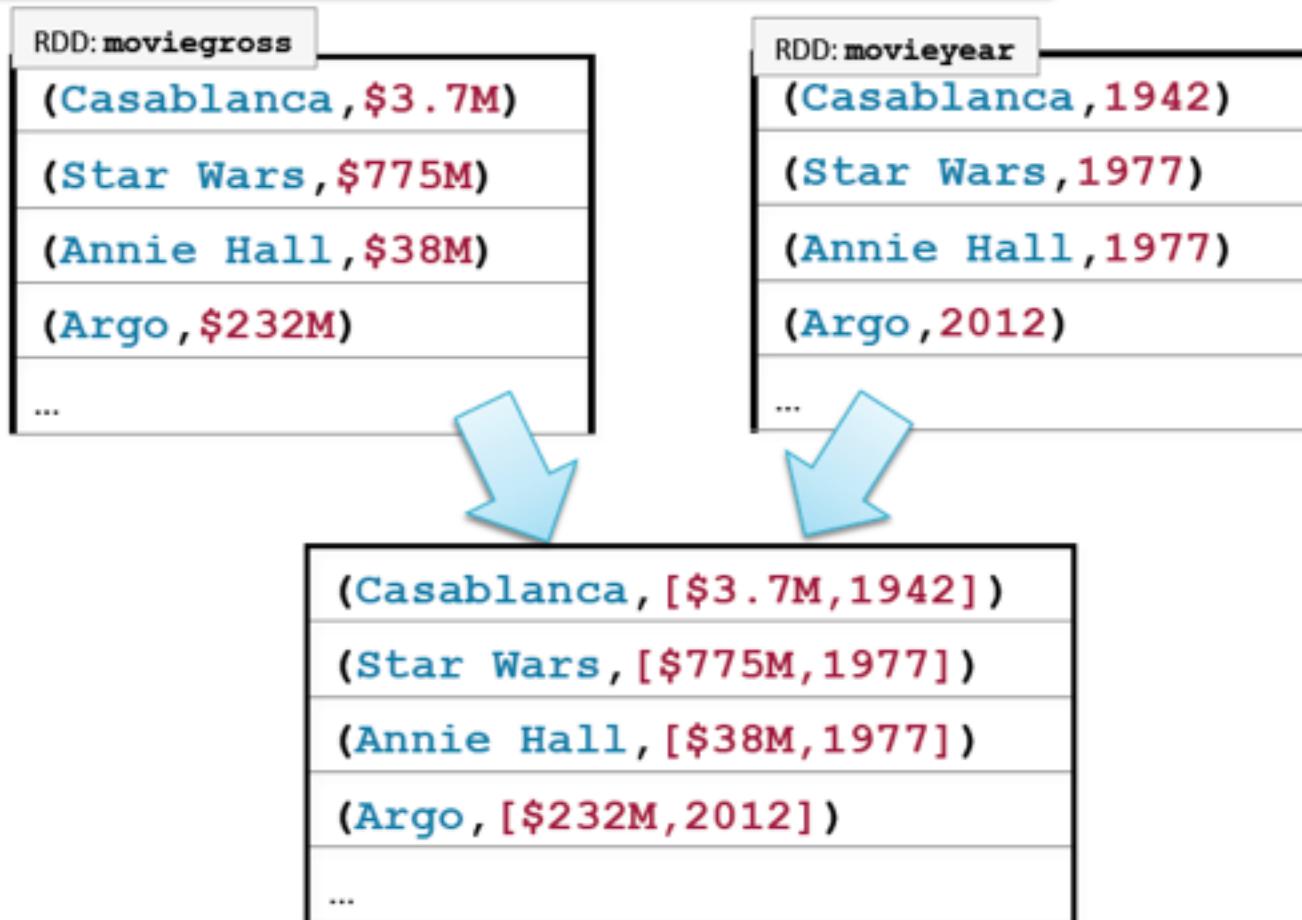


(00004, sku411)
(00003, sku888)
(00003, sku022)
(00003, sku010)
(00003, sku594)
(00002, sku912)
...

(00001, [sku010, sku933, sku022])
(00002, [sku912, sku331])
(00003, [sku888, sku022, sku010, sku594])
(00004, [sku411])

Example: Joining by Key

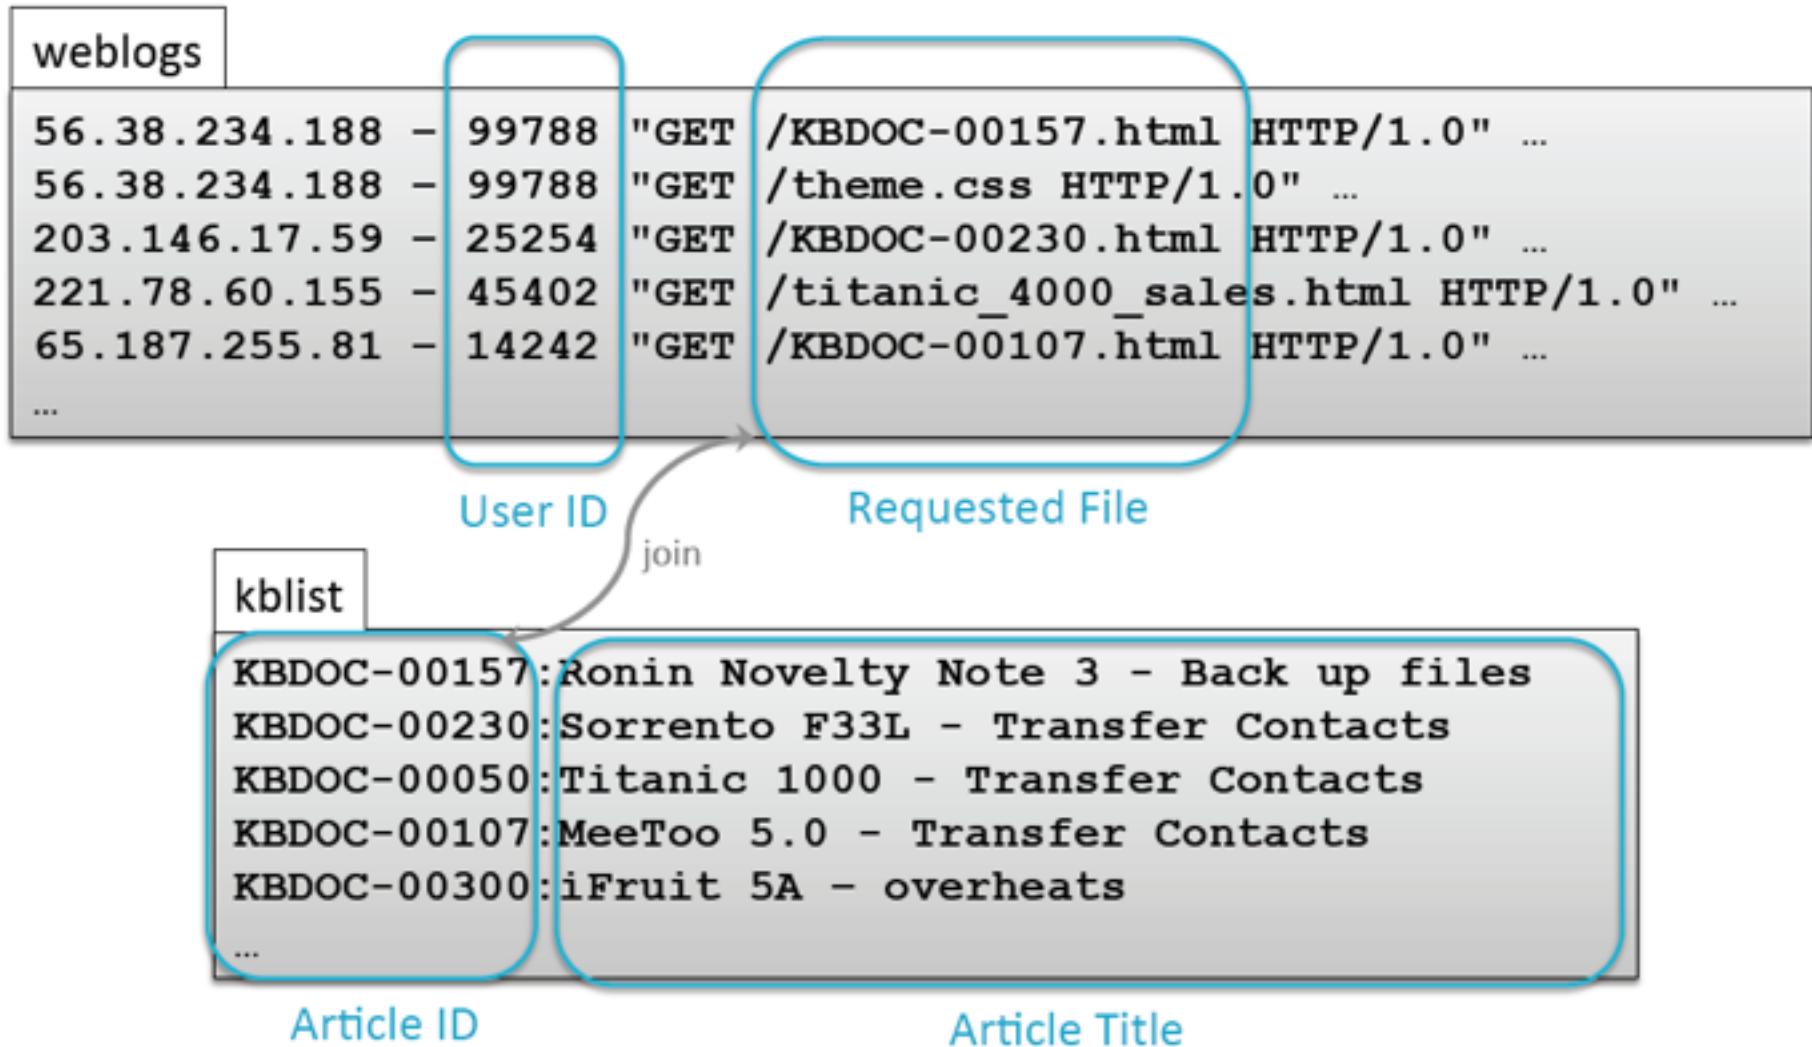
```
> movies = moviegross.join(movieyear)
```



Using Join

- A common programming pattern
 - 1. Map separate datasets into key-value Pair RDDs
 - 2. Join by key
 - 3. Map joined data into the desired format
 - 4. Save, display, or continue processing...

Example: Join Web Log With Knowledge Base Articles (1)



Example: Join Web Log With Knowledge Base Articles (2)

- Steps
 - 1. Map separate datasets into key-value Pair RDDs
 - a. Map web log requests to (docid,userid)
 - b. Map KB Doc index to (docid,title)
 - 2. Join by key: docid
 - 3. Map joined data into the desired format: (userid,title)
 - 4. Further processing: Group titles by userid

Step1a: Map Web Log Requests to (docid,userid)

```
> import re
> def getRequestDoc(s):
    return re.search(r'KBDOC-[0-9]*',s).group()

> kbreqs = sc.textFile(logfile) \
    .filter(lambda line: 'KBDOC-' in line) \
    .map(lambda line: (getRequestDoc(line),line.split(' ')[2])) \
    .distinct()
```

```
56.38.234.188 - 99788 "GET /KBDOC-00157.html HTTP/1.0" ...
56.38.234.188 - 99788 "GET /theme.css HTTP/1.0" ...
203.146.17.59 - 25254 "GET /KBDOC-00230.html HTTP/1.0"
221.78.60.155 - 45402 "GET /titanic_4000_sales.html I
65.187.255.81 - 14242 "GET /KBDOC-00107.html HTTP/1.
...

```



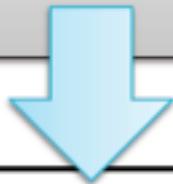
kbreqs
(KBDOC-00157, 99788)
(KBDOC-00203, 25254)
(KBDOC-00107, 14242)
...

Step1b: Map KB Index to (docid,userid)

```
> kblist = sc.textFile(kblistfile) \
    .map(lambda line: line.split(':')) \
    .map(lambda fields: (fields[0],fields[1]))
```

```
KBDOC-00157:Ronin Novelty Note 3 - Back up files  
KBDOC-00230:Sorrento F33L - Transfer Contacts  
KBDOC-00050:Titanic 1000 - Transfer Contacts  
KBDOC-00107:MeeToo 5.0 - Transfer Contacts  
KBDOC-00206:iFruit 5A - overheats  
...
```

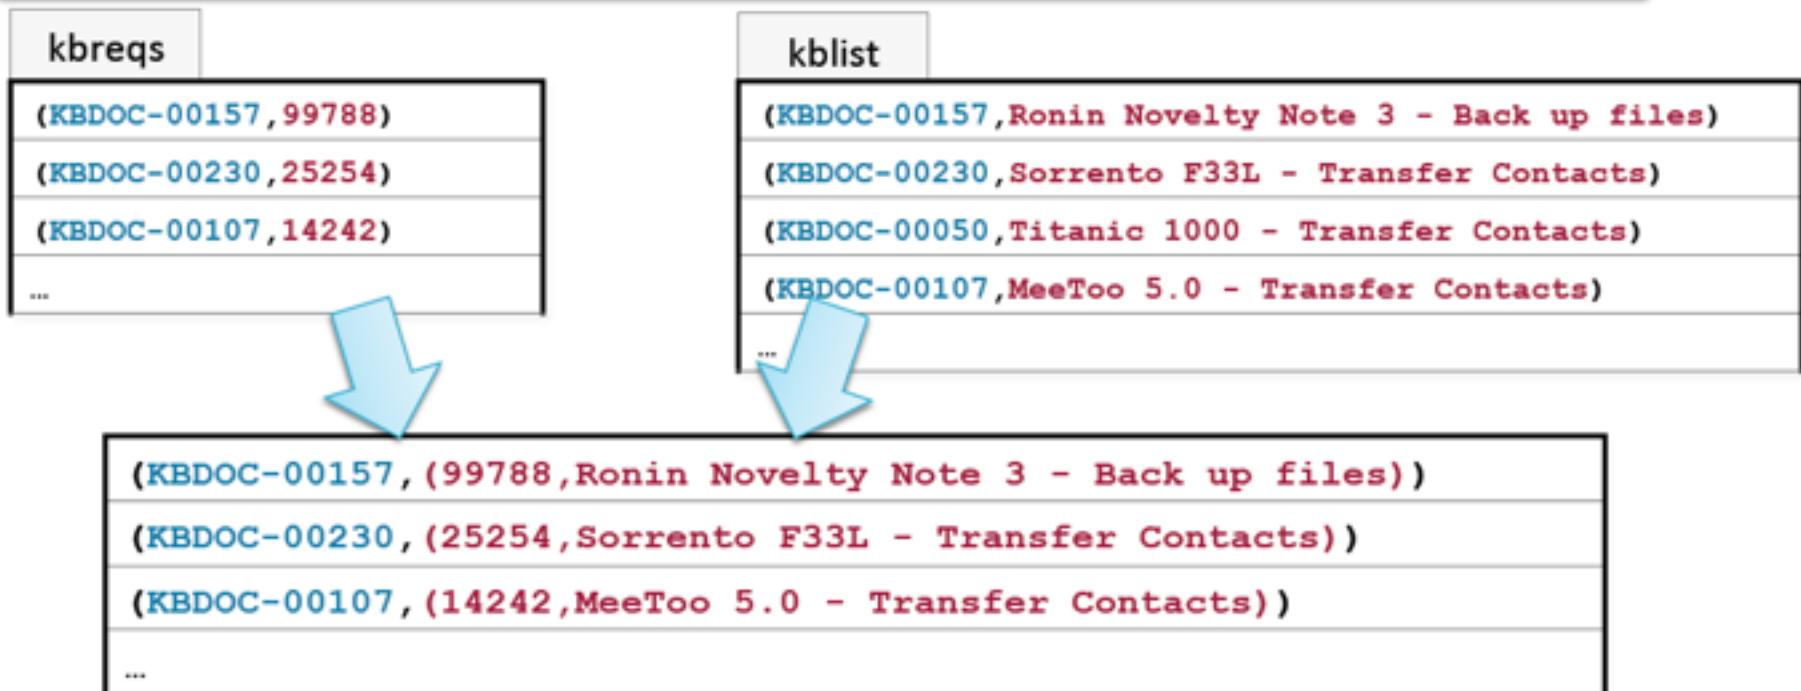
kblist



(KBDOC-00157 ,Ronin Novelty Note 3 - Back up files)
(KBDOC-00230 ,Sorrento F33L - Transfer Contacts)
(KBDOC-00050 ,Titanic 1000 - Transfer Contacts)
(KBDOC-00107 ,MeeToo 5.0 - Transfer Contacts)
...

Step 2: Join By Key docid

```
> titlereqs = kbreqs.join(kblist)
```



Step 3: Map Result to Desired Format (userid,title)

```
> titlereqs = kbreqs.join(kblist) \
    .map(lambda (docid, (userid, title)): (userid, title))
```

```
(KBDOC-00157, (99788, Ronin Novelty Note 3 - Back up files))
(KBDOC-00230, (25254, Sorrento F33L - Transfer Contacts))
(KBDOC-00107, (14242, MeeToo 5.0 - Transfer Contacts))
...
```

```
(99788, Ronin Novelty Note 3 - Back up files)
(25254, Sorrento F33L - Transfer Contacts)
(14242, MeeToo 5.0 - Transfer Contacts)
...
```

Step 4: Continue Processing- Group by User ID

```
> titlereqs = kbreqs.join(kblist) \
    .map(lambda (docid, (userid,title)): (userid,title)) \
    .groupByKey()
```

(99788, Ronin Novelty Note 3 - Back up files)

(25254, Sorrento F33L - Transfer Contacts)

(14242, MeeToo 5.0 - Transfer Contacts)

...

(99788, [Ronin Novelty Note 3 - Back up files,
Ronin S3 - overheating])

(25254, [Sorrento F33L - Transfer Contacts])

(14242, [MeeToo 5.0 - Transfer Contacts,
MeeToo 5.1 - Back up files,
iFruit 1 - Back up files,
MeeToo 3.1 - Transfer Contacts])

...

Example Output

```
> for (userid,titles) in titlereqs.take(10):  
    print 'user id: ',userid  
    for title in titles: print '\t',title
```

```
user id: 99788  
    Ronin Novelty Note 3 - Back up files  
    Ronin S3 - overheating  
user id: 25254  
    Sorrento F33L - Transfer Contacts  
user id: 14242  
    MeeToo 5.0 - Transfer Contacts  
    MeeToo 5.1 - Back up files  
    iFruit 1 - Back up files  
    MeeToo 3.1 - Transfer Contacts  
...
```

```
(99788,[Ronin Novelty Note 3 - Back up files,  
         Ronin S3 - overheating])  
(25254,[Sorrento F33L - Transfer Contacts])  
(14242,[MeeToo 5.0 - Transfer Contacts,  
         MeeToo 5.1 - Back up files,  
         iFruit 1 - Back up files,  
         MeeToo 3.1 - Transfer Contacts])  
...
```

Aside: Anonymous Function Parameters

Python

```
> map(lambda (docid,(userid,title)): (userid,title))
```

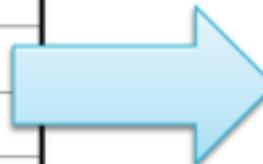
Scala

```
> map(pair => (pair._2._1,pair._2._2))
```

OR

```
> map{case (docid,(userid,title)) => (userid,title)}
```

(KBDOC-00157, (99788,...title...))
(KBDOC-00230, (25254,...title...))
(KBDOC-00107, (14242,...title...))
...



(99788,...title...)
(25254,...title...)
(14242,...title...))
...

Other Pair Operations

- Some other pair operations
 - keys - return an RDD of just the keys, without the values
 - values - return an RDD of just the values, without keys
 - lookup(key)
 - leftOuterJoin, rightOuterJoin - join, including keys defined only in the left or right RDDs respectively
 - mapValues, flatMapValues - execute a function on just the values, keeping the key the same
- See the `PairRDDFunctions` class Scaladoc for a full list

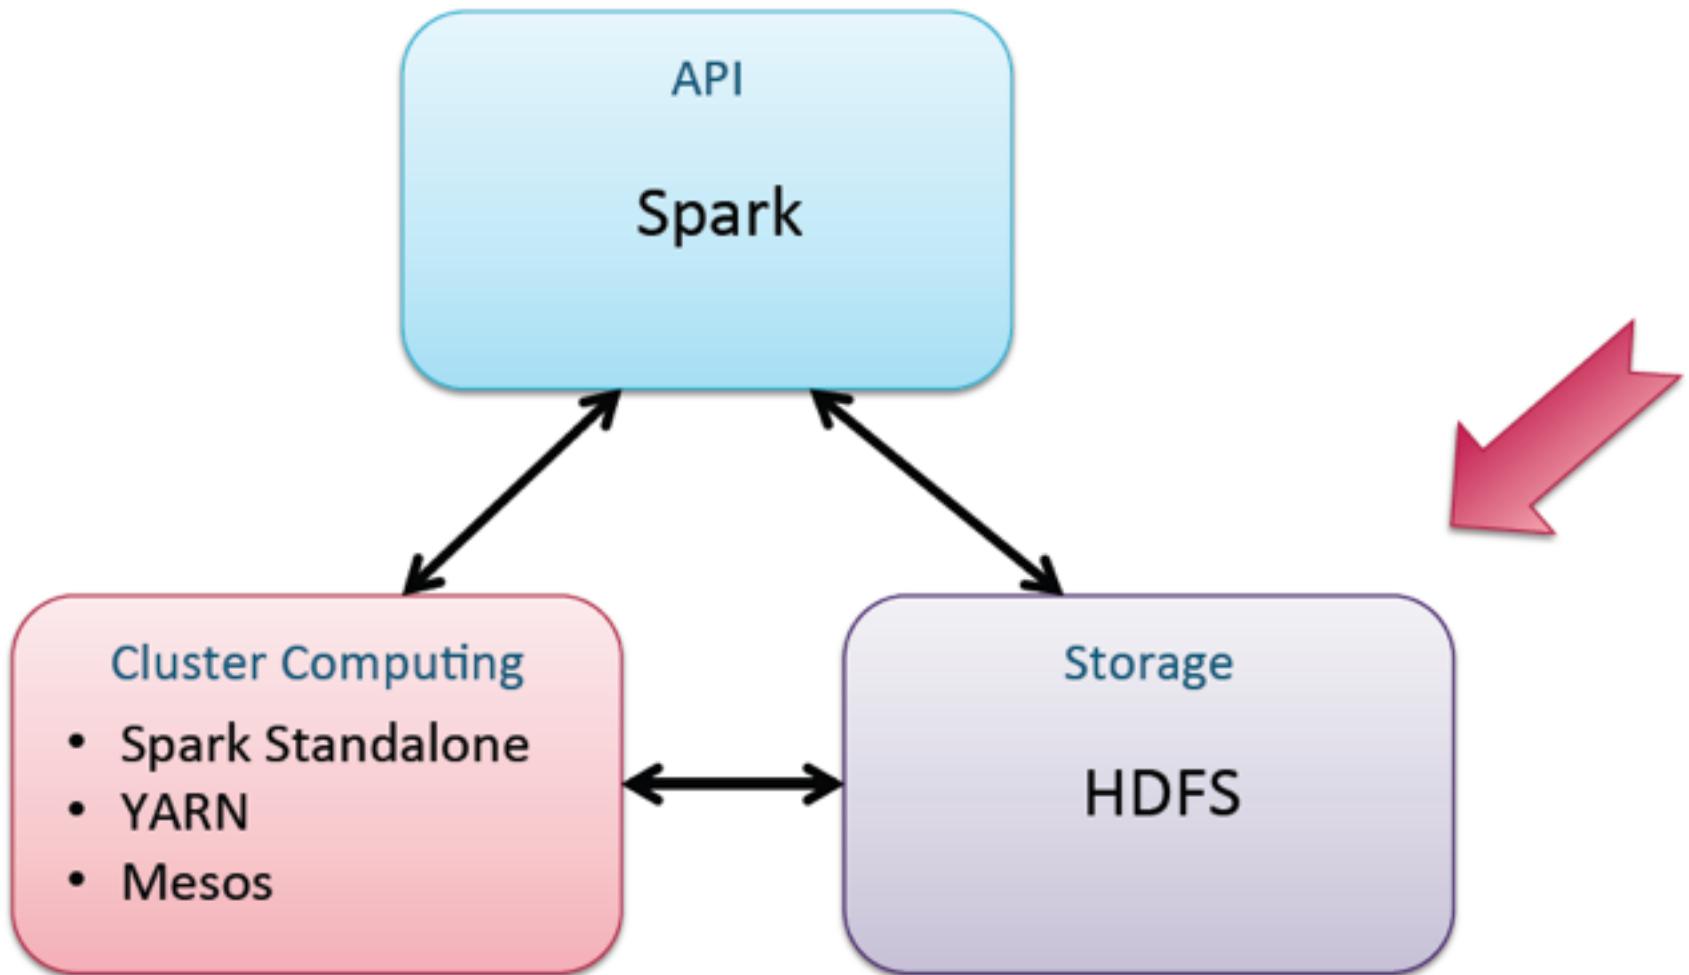
Hands-On Exercise: Working with Pair RDDs

HANDS-ON EXERCISE: WORKING WITH PAIR RDDS

Chapter 4

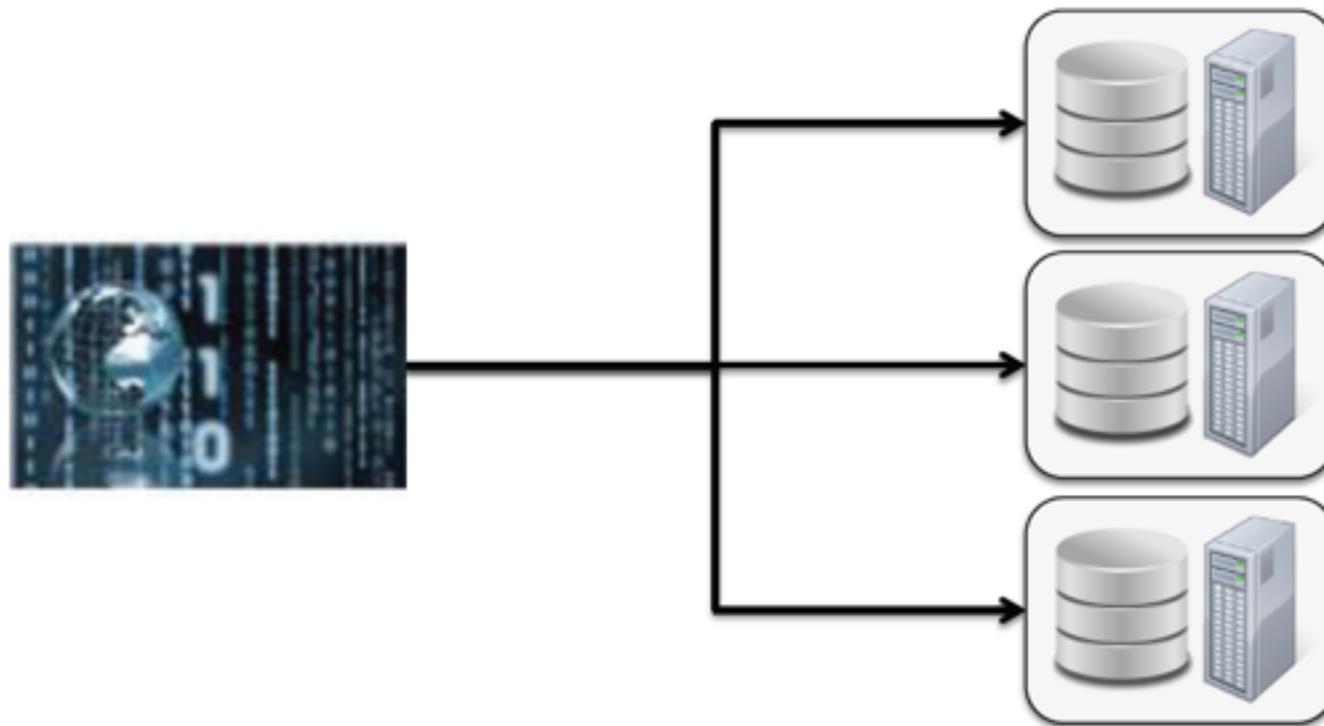
THE HADOOP DISTRIBUTED FILE SYSTEM (HDFS)

Distributed Processing with the Spark Framework



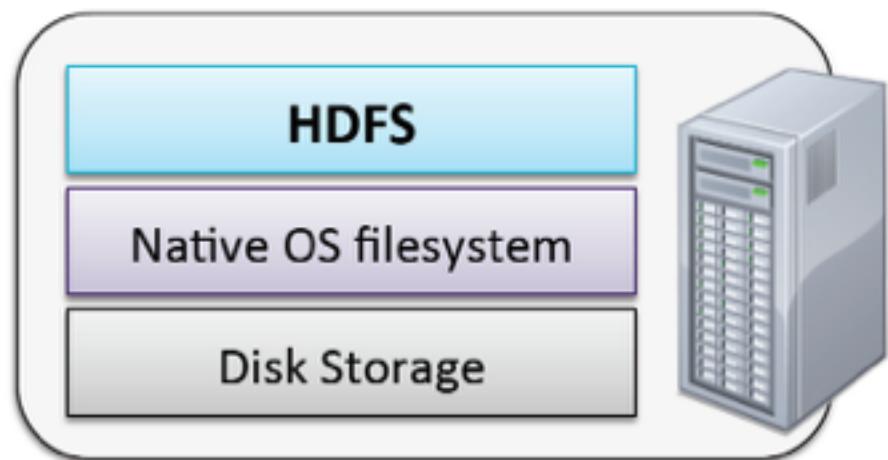
Big Data Processing with Spark

- Three key concepts
 - Distribute data when the data is stored - HDFS
 - Run computation where the data is - HDFS and Spark
 - Cache data in memory - Spark



HDFS Basic Concepts

- HDFS is a filesystem written in Java
 - Based on Google's GFS
- Sits on top of a native filesystem
- Provides redundant storage for massive amounts of data
 - Using readily-available, industry-standard computers

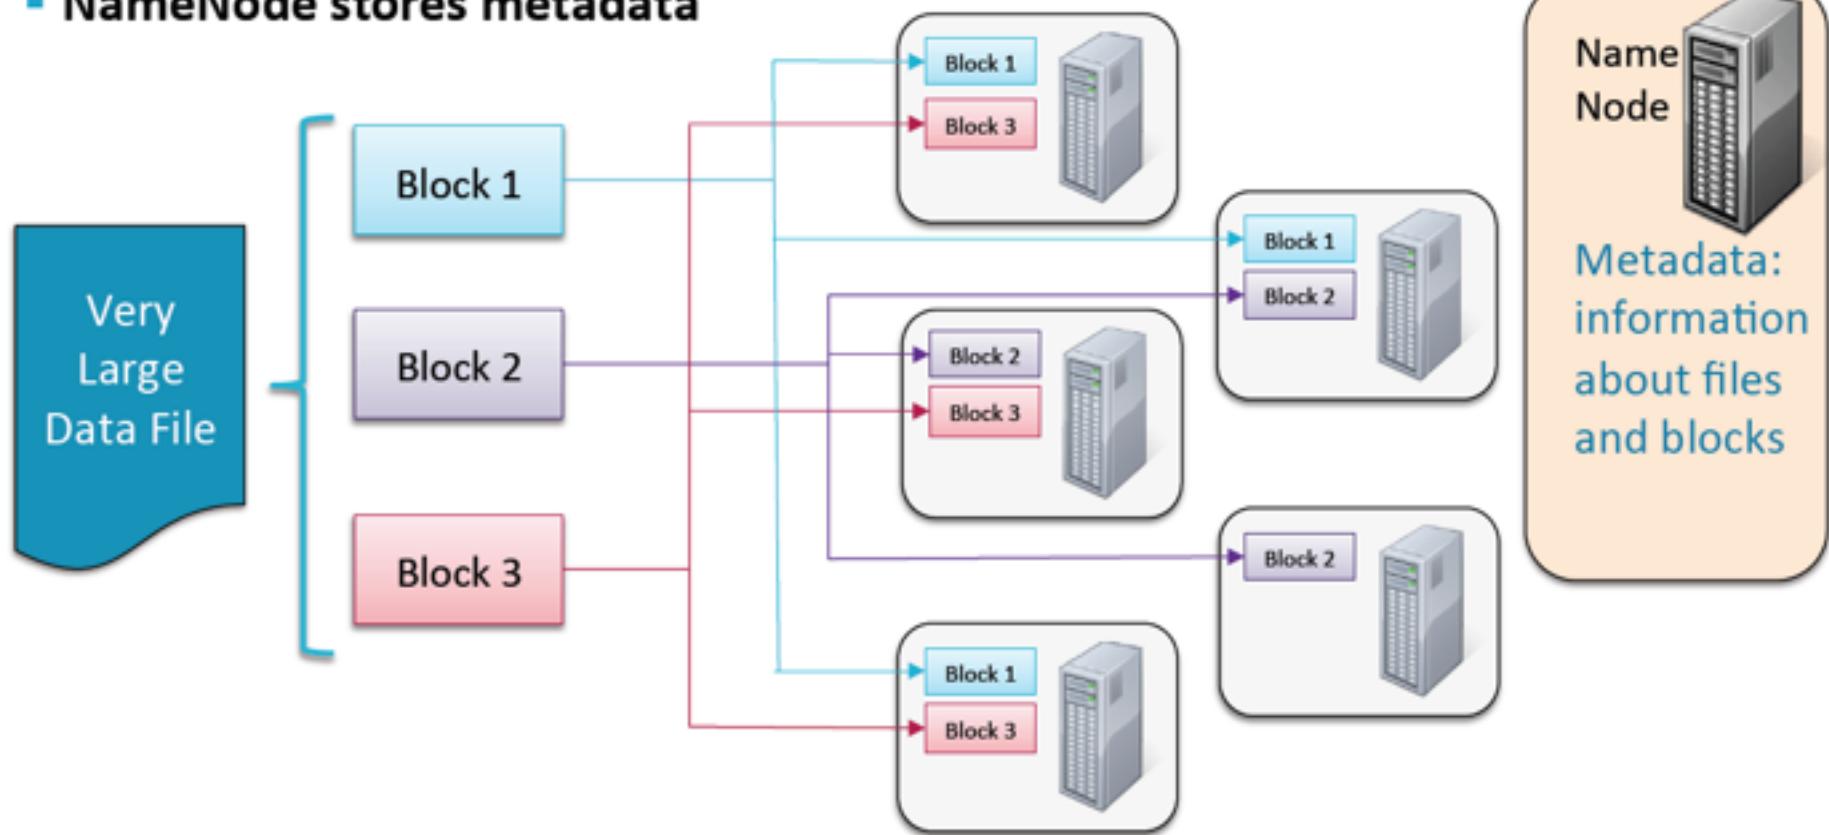


HDFS Basic Concepts

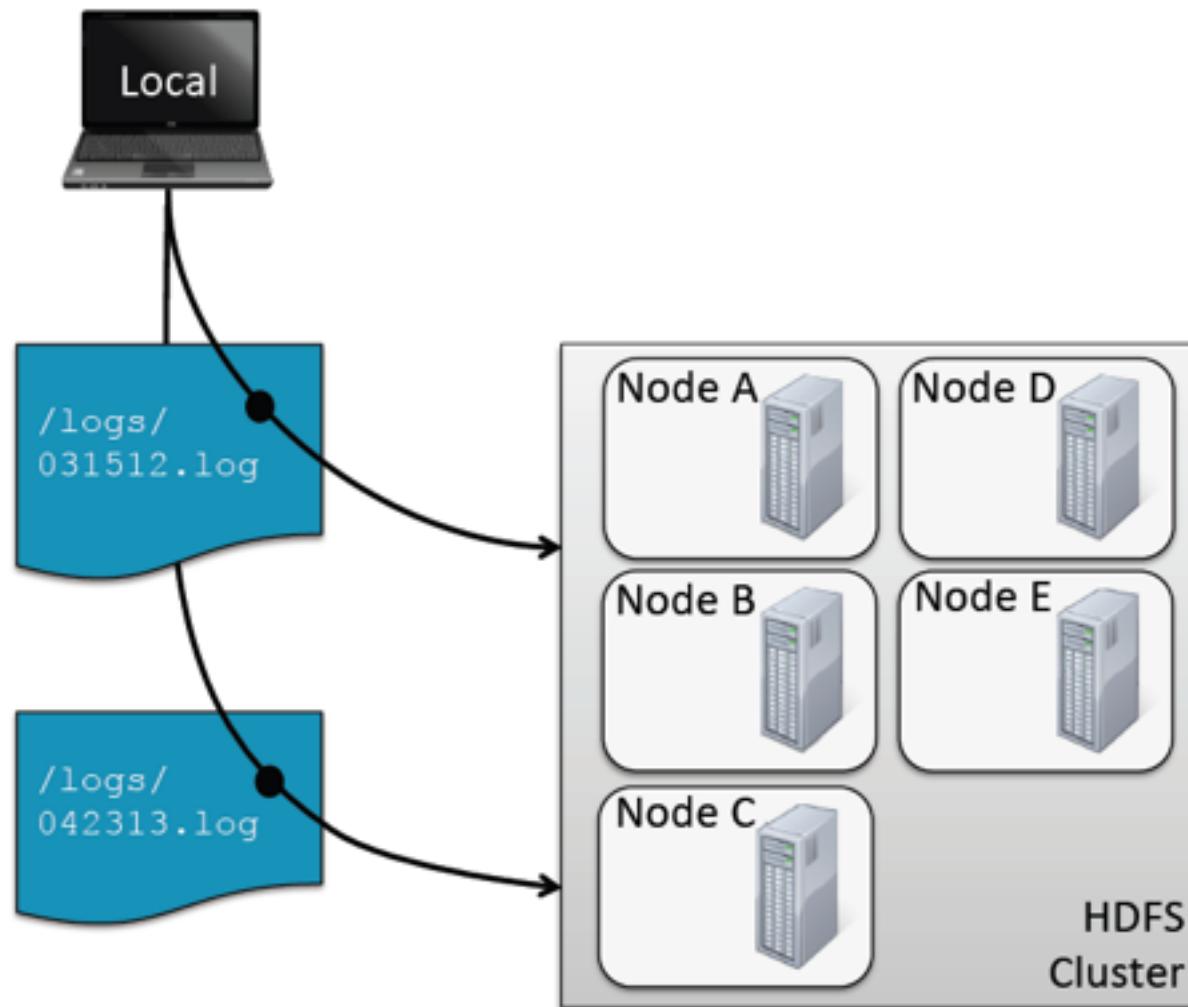
- HDFS performs best with a ‘modest’ number of large files
 - Millions, rather than billions, of files
 - Each file typically 100MB or more
- Files in HDFS are ‘write once’
 - No random writes to files are allowed
- HDFS is optimized for large, streaming reads of files
 - Rather than random reads

How Files Are Stored

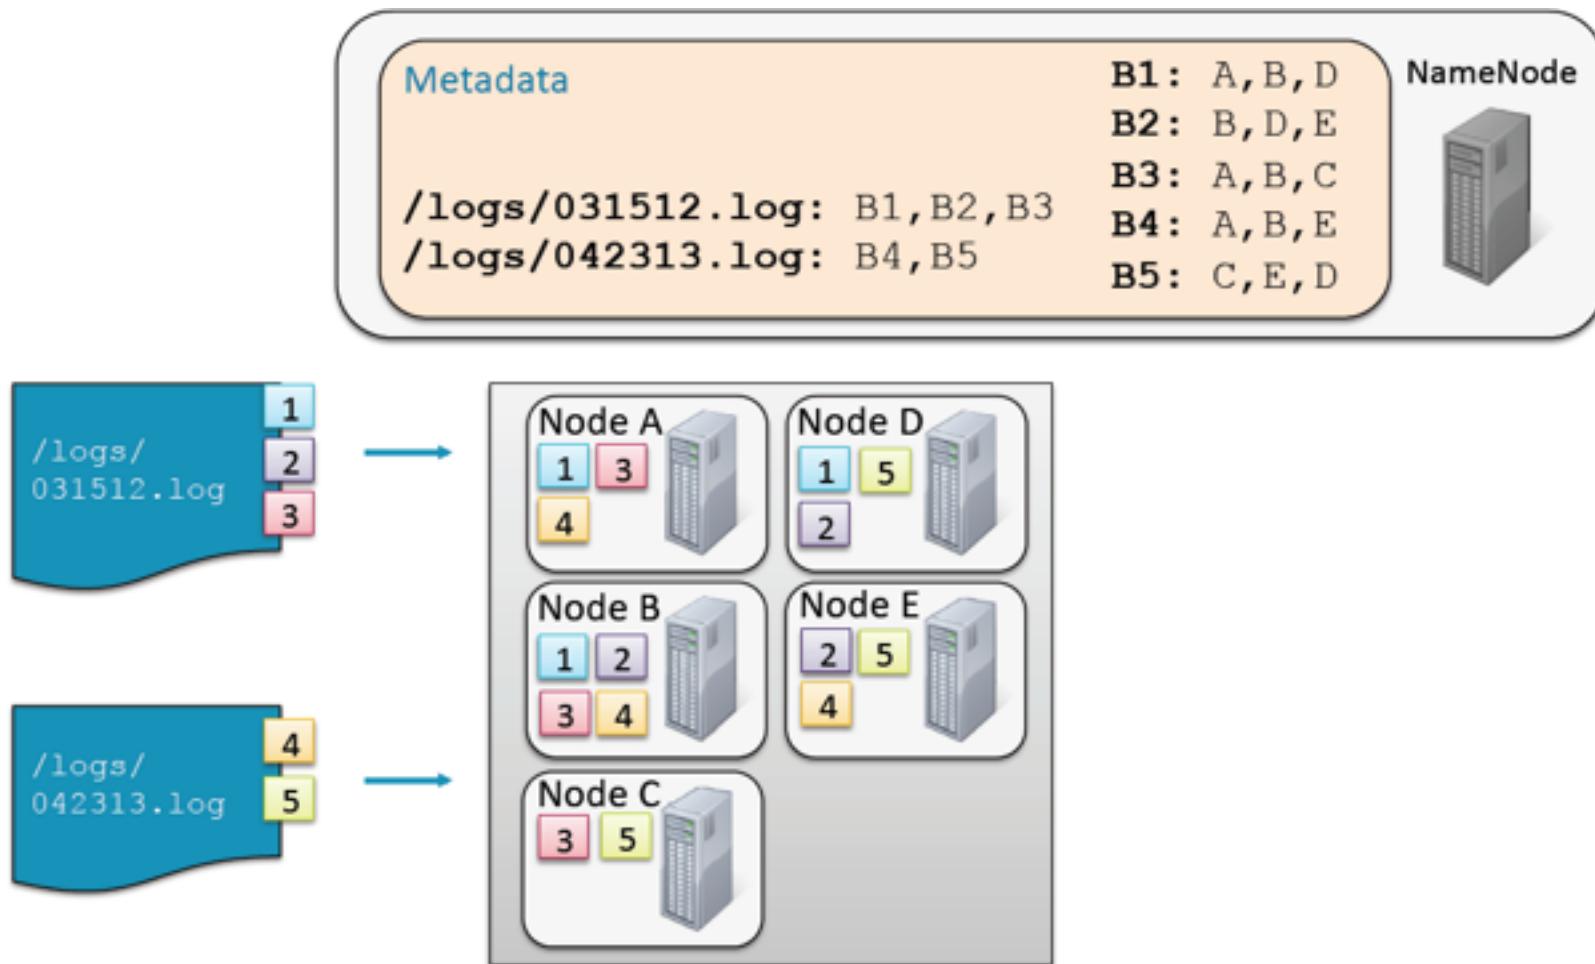
- NameNode stores metadata



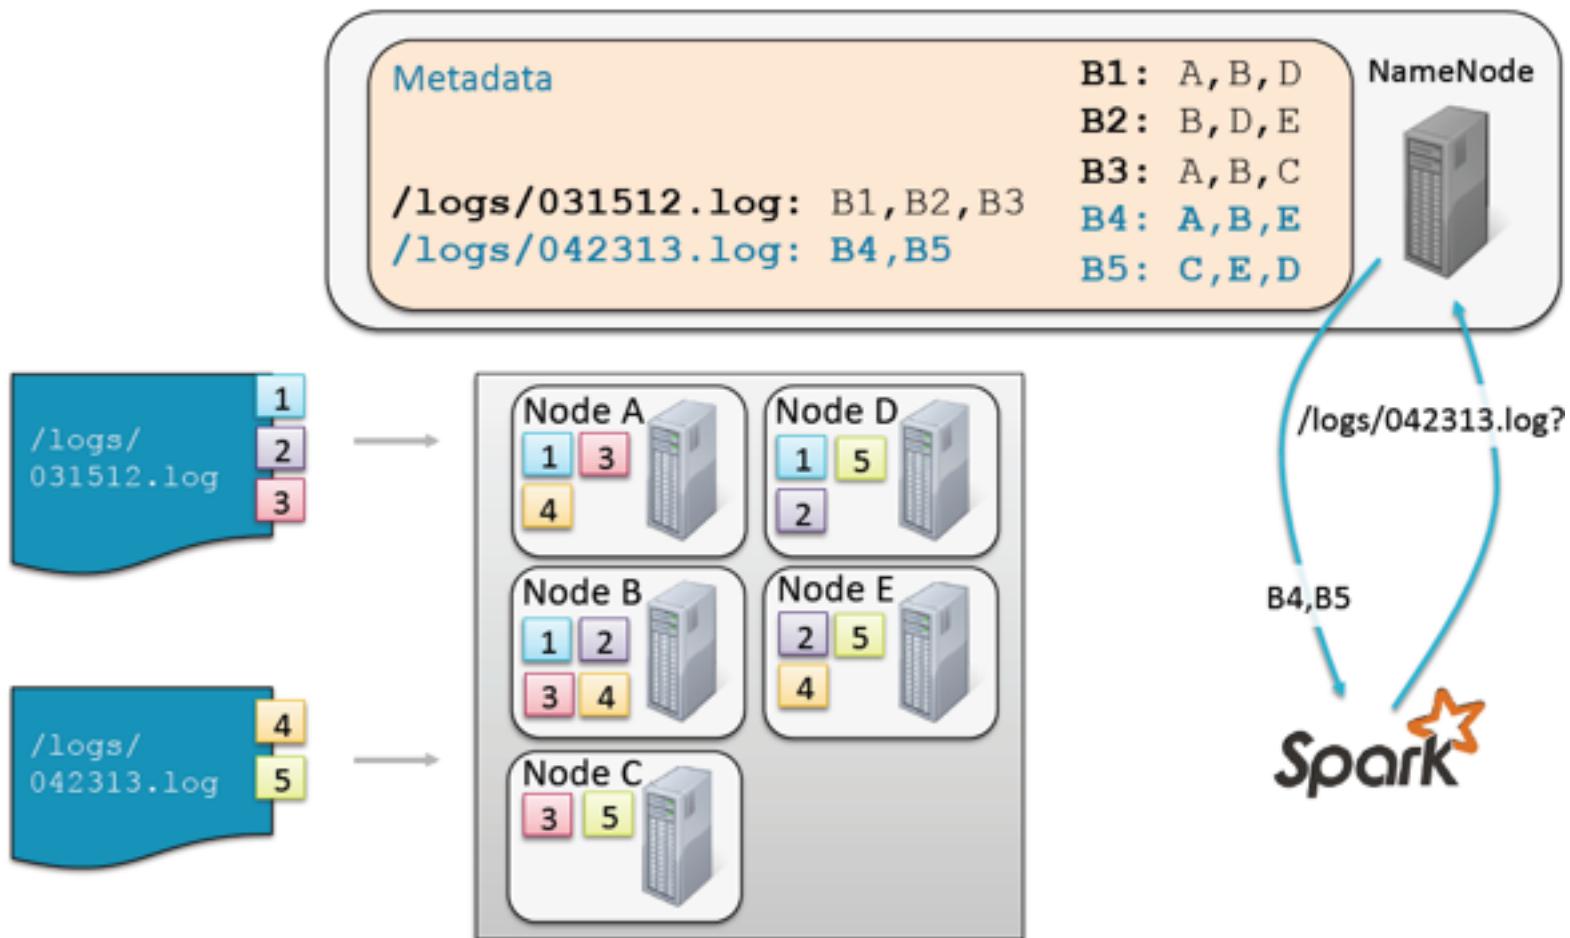
Storing and Retrieving Files



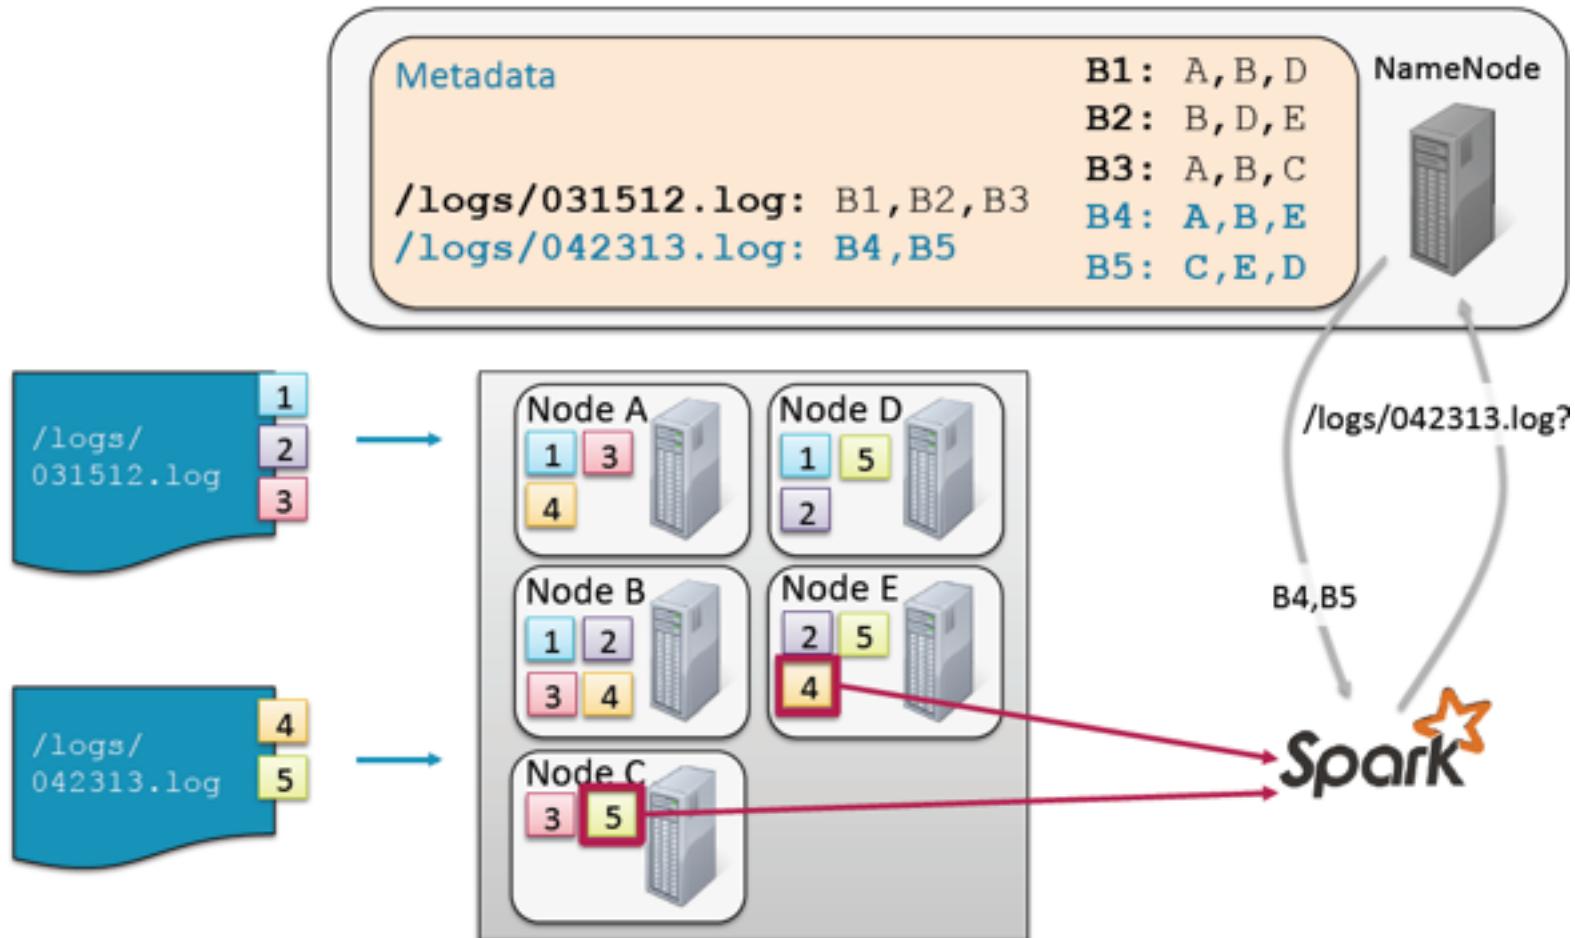
Storing and Retrieving Files



Storing and Retrieving Files



Storing and Retrieving Files



HDFS NameNode Availability

- The NameNode daemon must be running at all times
 - If the NameNode stops, the cluster becomes inaccessible
- HDFS is typically set up for High Availability
 - Two NameNodes: Active and Standby"
- Small clusters may use ‘Classic mode’
 - One NameNode
 - One “helper” node called the Secondary NameNode
 - Bookkeeping, not backup

Options for Accessing HDFS

- From the command line
 - FsShell: `hdfs dfs`"
- In Spark
 - By URI, e.g. `hdfs://host:port/file...`
- Other programs
 - Java API
 - Used by Hadoop MapReduce,
 - Hue, Sqoop,
 - Flume, etc.
 - RESTful interface

hdfs dfs Examples

```
$ hdfs dfs -put foo.txt foo.txt
```

- This will copy the file to /user/username/foo.txt
- **Get a directory listing of the user's home directory in HDFS**

```
$ hdfs dfs -ls
```

- **Get a directory listing of the HDFS root directory**

```
$ hdfs dfs -ls /
```

hdfs dfs Examples

```
$ hdfs dfs -cat /user/fred/bar.txt
```

- Copy that file to the local disk, named as **baz.txt**

```
$ hdfs dfs -get /user/fred/bar.txt baz.txt
```

- Create a directory called **input** under the user's home directory

```
$ hdfs dfs -mkdir input
```

Note: copyFromLocal is a synonym for put; copyToLocal is a synonym for get

Example: HDFS in Spark

- Specify HDFS files in Spark by URI
 - `hdfs://hdfs-host[:port]/path`
 - Default port is 8020

```
> mydata = sc.textFile \
  ("hdfs://hdfs-host:port/user/training/purplecow.txt")

> mydata.map(lambda s: s.upper()).\
  saveAsTextFile \
  ("hdfs://hdfs-host:port/user/training/purplecowuc")
```

Using HDFS By Default

- If Hadoop configuration files are on Spark's classpath, Spark will use HDFS by default
 - e.g. /etc/hadoop/conf
- Paths are relative to user's home HDFS directory

```
> mydata = sc.textFile("purplecow.txt")
```

hdfs://hdfs-host:port/user/training/purplecow.txt

Hands On Exercise: Using HDFS

HANDS-ON EXERCISE: USING HDFS

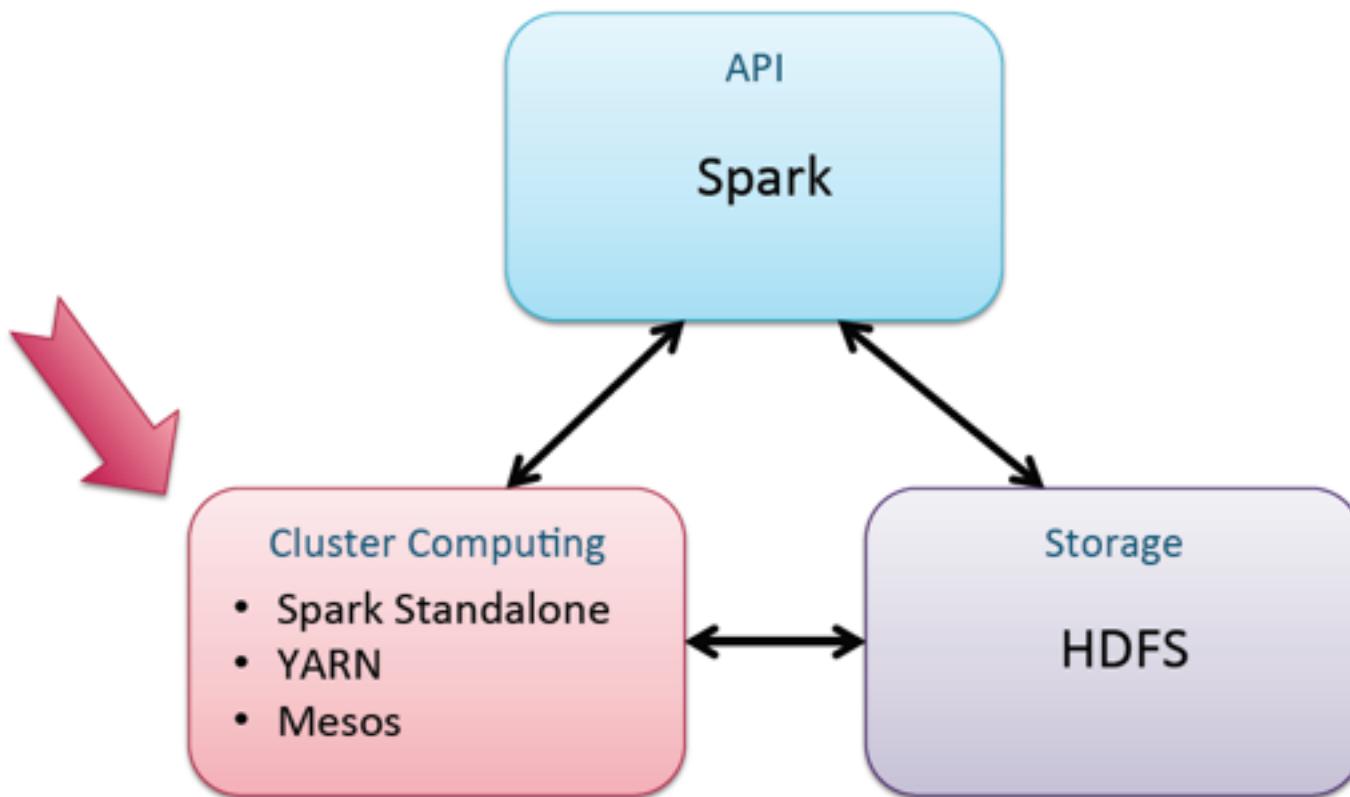
Chapter 5

SPARK ON A CLUSTER

Spark Cluster Options

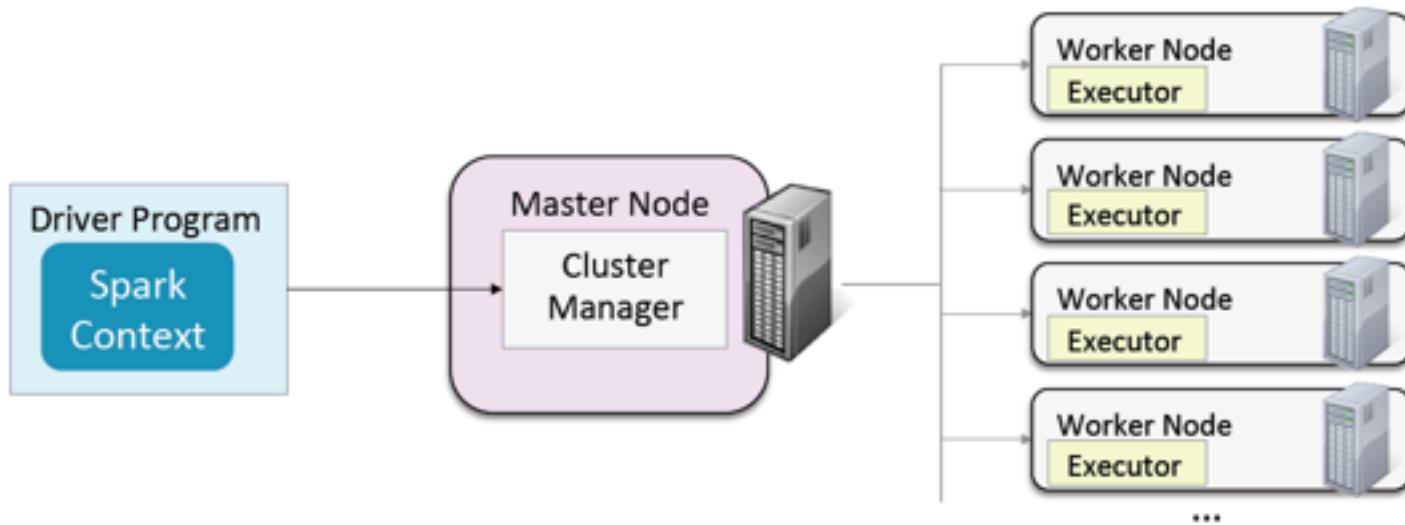
- Spark can run
 - Locally
 - No distributed processing
 - Locally with multiple worker threads
 - On a cluster
 - Spark Standalone
 - Apache Hadoop YARN (Yet Another Resource Negotiator)

Distributed Processing with the Spark Framework



The Spark Driver Program

- A Spark Driver
 - The “main” program
 - Either the Spark Shell or a Spark application
 - Creates a Spark Context configured for the cluster
 - Communicates with Cluster Manager to distribute tasks to executors



Starting the Spark Shell on a Cluster

- Set the Spark Shell master to
 - url - the URL of the cluster manager
 - local[*] - run with as many threads as cores (default)
 - local[n] - run locally with n worker threads
 - local - run locally without distributed processing
- This configures the `SparkContext.master` property

Python

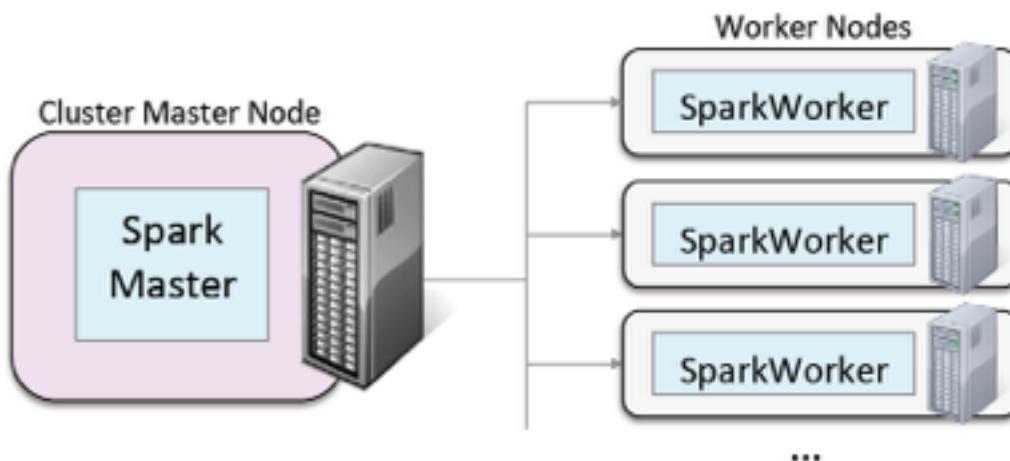
```
$ MASTER=spark://masternode:7077 pyspark
```

Scala

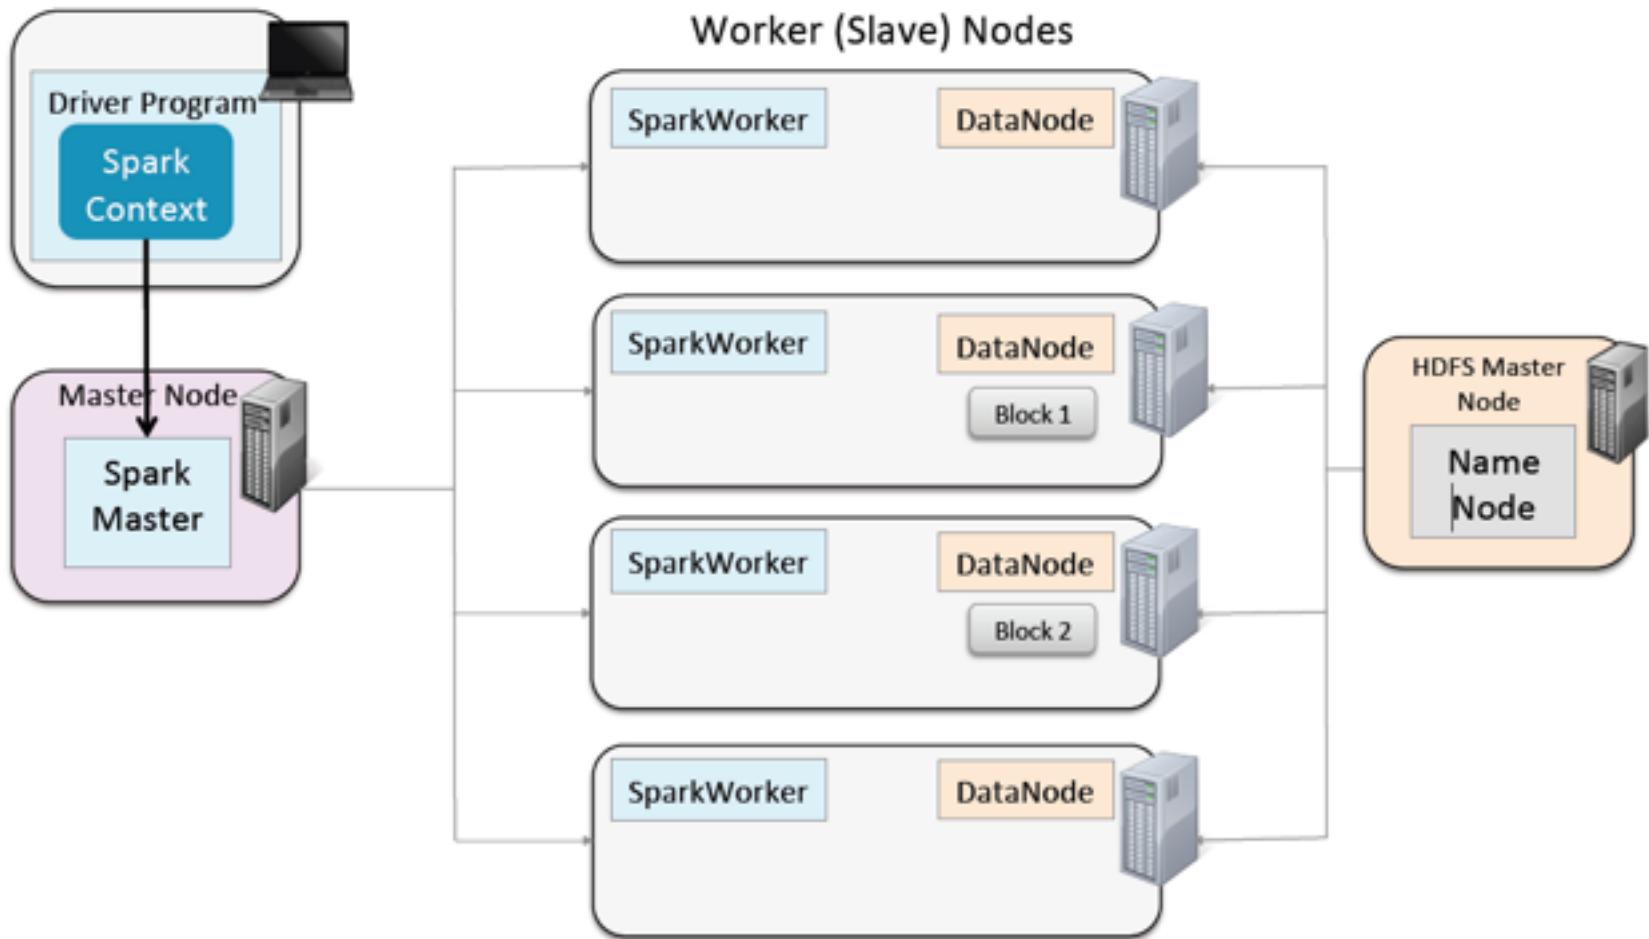
```
$ spark-shell --master spark://masternode:7077
```

Spark Standalone Daemons

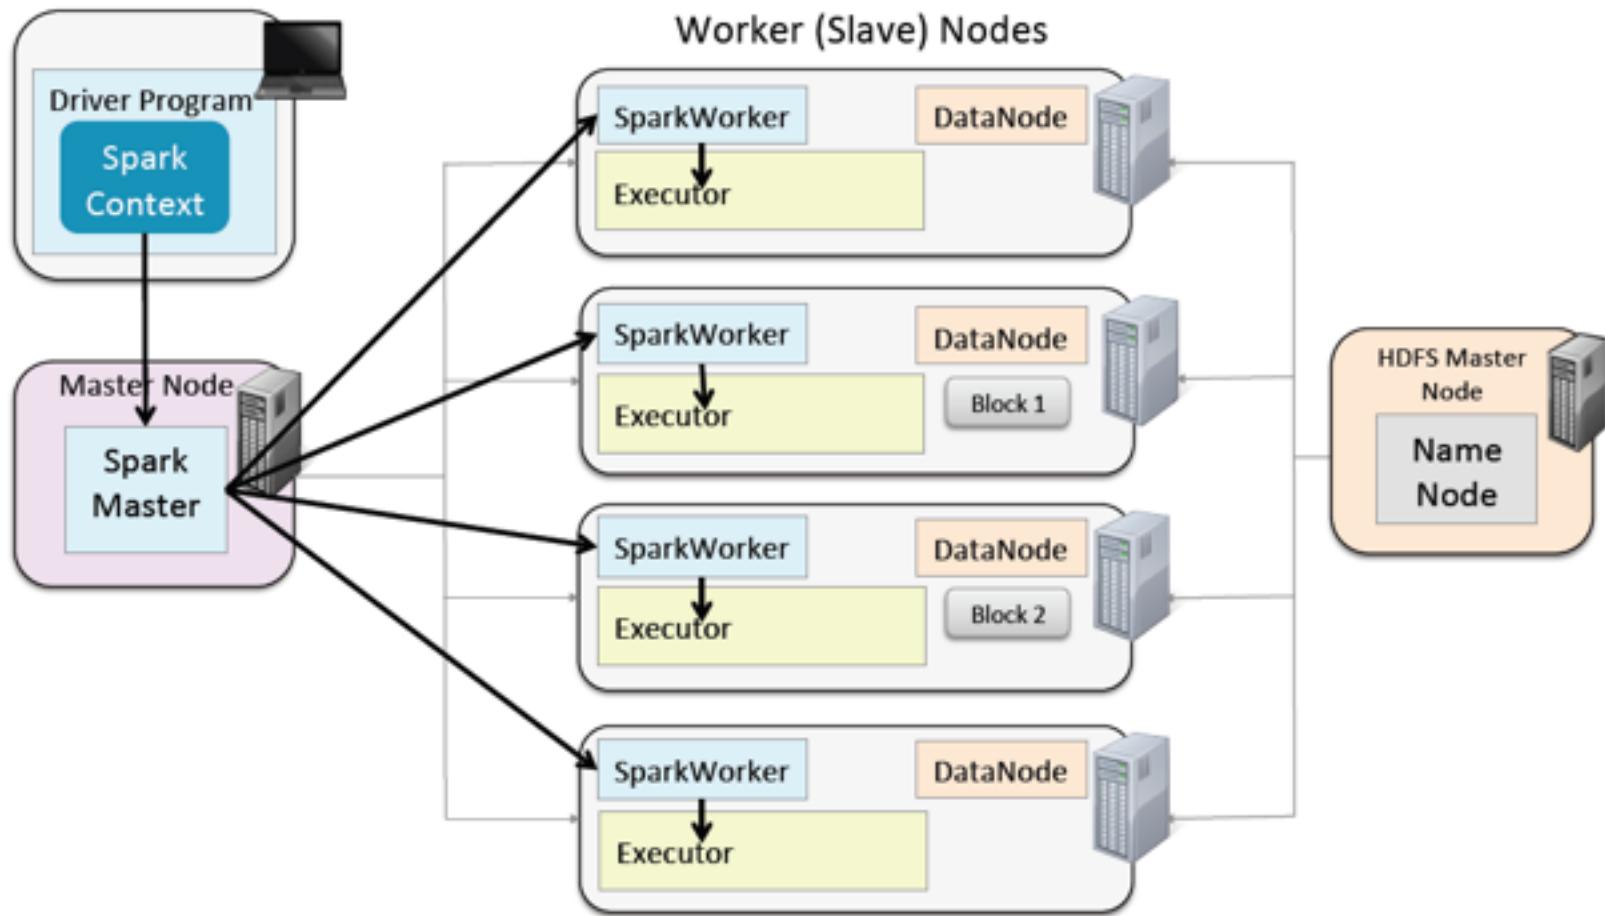
- Spark Standalone daemons
 - Spark Master
 - One per cluster
 - Manages applications, distributes individual tasks to Spark Workers
 - Spark Worker
 - One per worker node
 - Starts and monitors Executors for applications



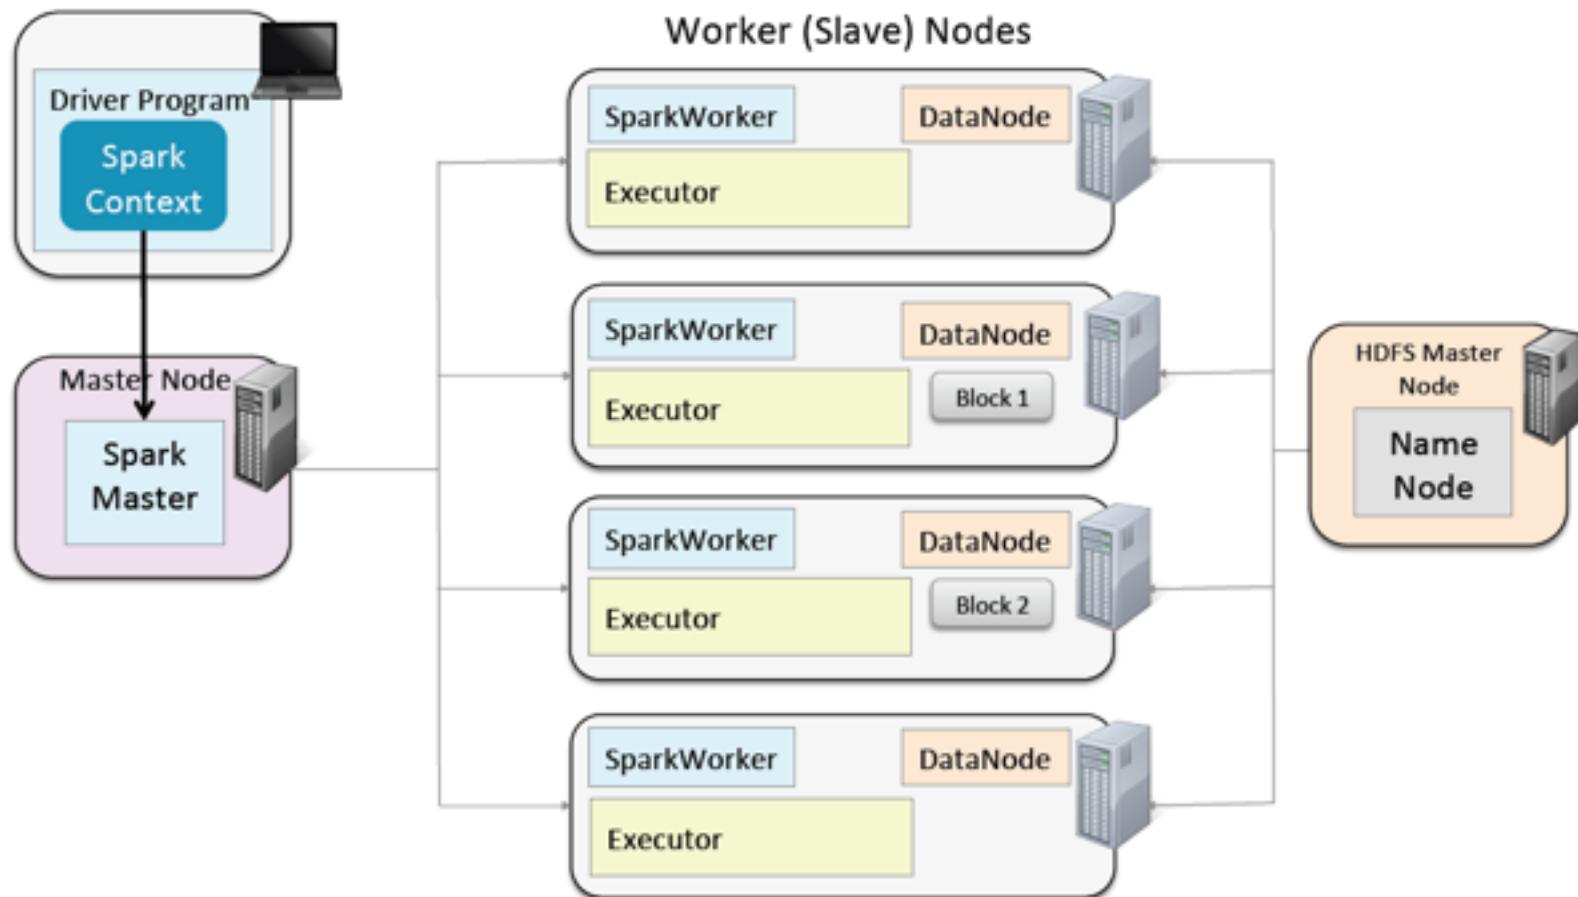
Running Spark on a Standalone Cluster



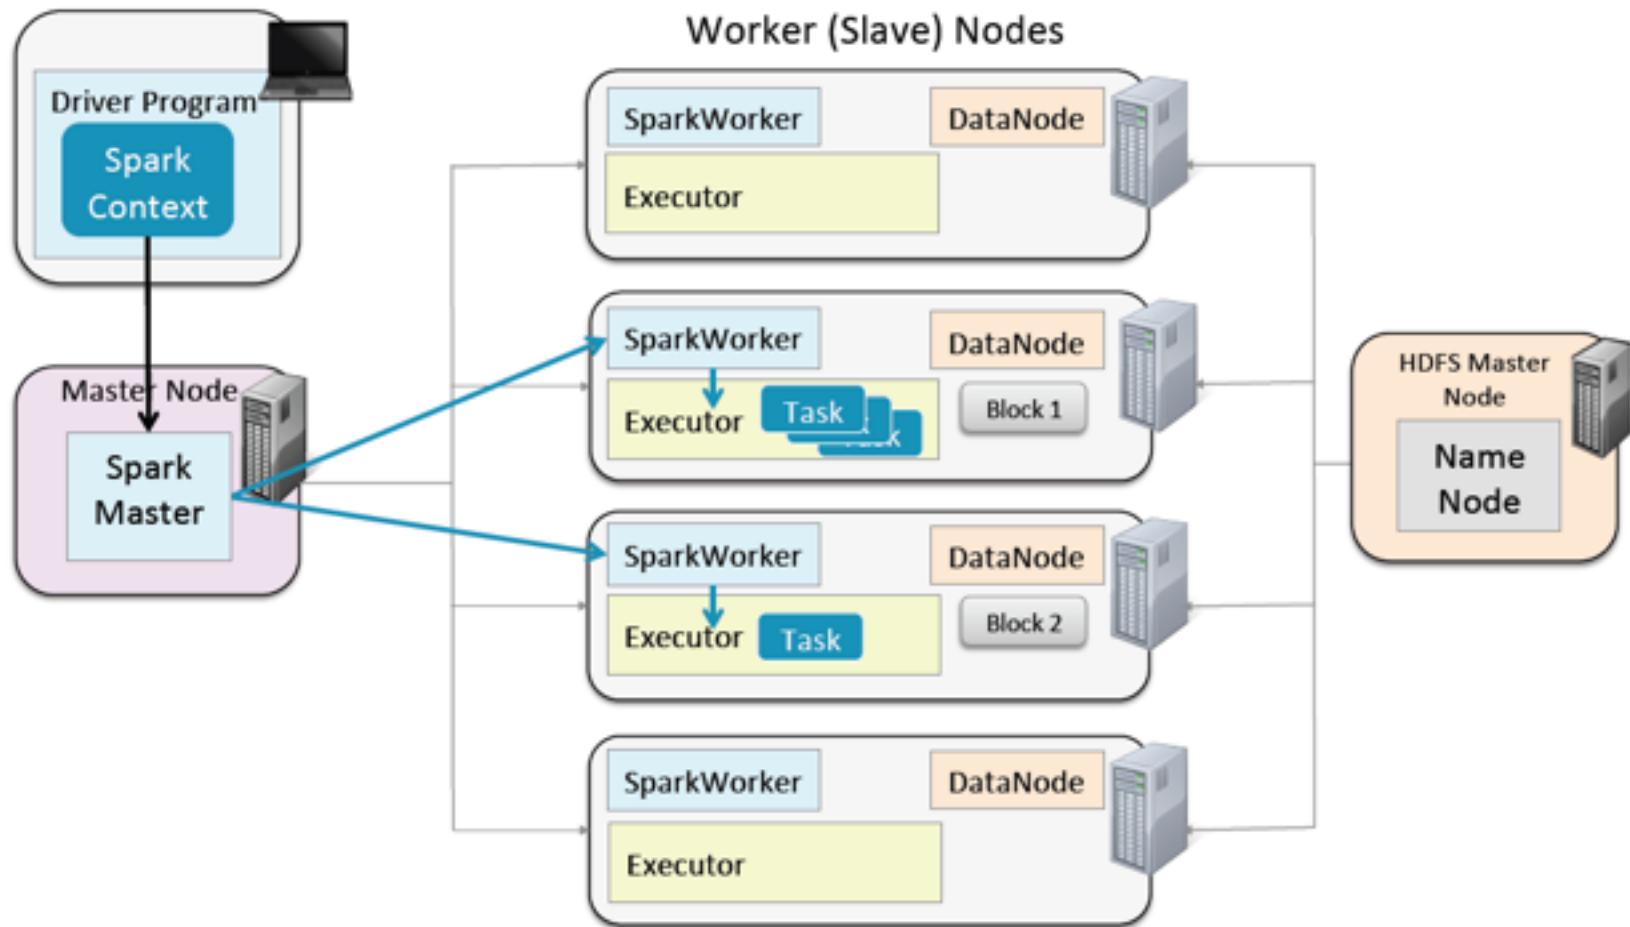
Running Spark on a Standalone Cluster



Running Spark on a Standalone Cluster



Running Spark on a Standalone Cluster



Spark Standalone Web UI

Spark Standalone clusters offer a Web UI to monitor the cluster

- `http://masternode:uiport`
- – e.g., in our class environment, `http://localhost:18080`

 Spark Master at spark://ec2-23-20-24-104.compute-1.amazonaws.com:7077

URL: spark://ec2-23-20-24-104.compute-1.amazonaws.com:7077
Workers: 5
Cores: 20 Total, 20 Used
Memory: 68.2 GB Total, 63.2 GB Used
Applications: 1 Running, 2 Completed

Master URL

Worker Nodes

ID	Address	State	Cores	Memory
worker-20140121065745-ip-10-236-129-42.ec2.internal:50105	ip-10-236-129-42.ec2.internal:50105	ALIVE	4 (4 Used)	13.6 GB (12.6 GB Used)
worker-20140121065747-ip-10-137-18-53.ec2.internal:54087	ip-10-137-18-53.ec2.internal:54087	ALIVE	4 (4 Used)	13.6 GB (12.6 GB Used)
worker-20140121065747-ip-10-138-3-46.ec2.internal:50661	ip-10-138-3-46.ec2.internal:50661	ALIVE	4 (4 Used)	13.6 GB (12.6 GB Used)
worker-20140121065747-ip-10-236-151-85.ec2.internal:60016	ip-10-236-151-85.ec2.internal:60016	ALIVE	4 (4 Used)	13.6 GB (12.6 GB Used)
worker-20140121065748-ip-10-238-128-41.ec2.internal:42252	ip-10-238-128-41.ec2.internal:42252	ALIVE	4 (4 Used)	13.6 GB (12.6 GB Used)

Running Applications

ID	Name	Cores	Memory per Node	Submitted Time	User	State	Duration
app-20140121215958-0002	PageRank	20	12.6 GB	2014/01/21 21:59:58	root	RUNNING	13 s

Completed Applications

ID	Name	Cores	Memory	Submitted Time	User	State	Duration
app-20140121215822-0001	SparkPI	20	12.6 GB	2014/01/21 21:58:22	root	FINISHED	10 s
app-20140121215016-0000	Spark shell	20	12.6 GB	2014/01/21 21:50:16	root	FINISHED	1.2 min

Applications

Spark Standalone Web UI: Application Overview

Running Applications		
ID	Name	Cores
app-20140121215958-0002	PageRank	20

Completed Applications

Link to Spark Application UI

Spark Application: PageRank

ID: app-20140121220952-0006
Name: PageRank
User: root
Cores: Unlimited (20 granted)
Executor Memory: 12.6 GB
Submit Date: Tue Jan 21 22:09:52 UTC 2014
State: RUNNING
[Application Detail UI](#)

Executor Summary

ExecutorID	Worker	Cores	Memory	State	Logs
3	worker-20140121065747-ip-10-138-3-46.ec2.internal-50661	4	12936	RUNNING	stdout stderr
4	worker-20140121065745-ip-10-236-129-42.ec2.internal-60105	4	12936	RUNNING	stdout stderr
1	worker-20140121065748-ip-10-238-128-41.ec2.internal-42252	4	12936	RUNNING	stdout stderr
2	worker-20140121065747-ip-10-236-151-85.ec2.internal-60016	4	12936	RUNNING	stdout stderr
0	worker-20140121065747-ip-10-137-18-53.ec2.internal-54087	4	12936	RUNNING	stdout stderr

Executors for this application

Spark Standalone Web UI: Worker Detail

Workers

ID	Address
worker-20140121065745-ip-10-236-129-42.ec2.internal-60105	ip-10-236-129-42.ec2.internal:60105
worker-20140121065747-ip-10-137-18-53.ec2.internal-54087	ip-10-137-18-53.ec2.internal:54087
worker-20140121065747-ip-10-138-3-45.ec2.internal-50661	ip-10-138-3-45.ec2.internal:50661

Spark Worker at ip-10-236-129-42.ec2.internal:60105

ID: worker-20140121065745-ip-10-236-129-42.ec2.internal-60105
Master URL: spark://ec2-23-20-24-104.compute-1.amazonaws.com:7077
Cores: 4 (4 Used)
Memory: 13.6 GB (12.6 GB Used)

[Back to Master](#)

Running Executors 1

ExecutorID	Cores	Memory	Job Details	Logs
4	4	12.6 GB	ID: app-20140121220135-0003 Name: PageRank User: root	stdout stderr

Finished Executors

ExecutorID	Cores	Memory	Job Details	Logs
4	4	12.6 GB	ID: app-20140121215522-0001 Name: SparkPi User: root	stdout stderr
4	4	12.6 GB	ID: app-20140121215958-0002 Name: PageRank User: root	stdout stderr

All executors on this node

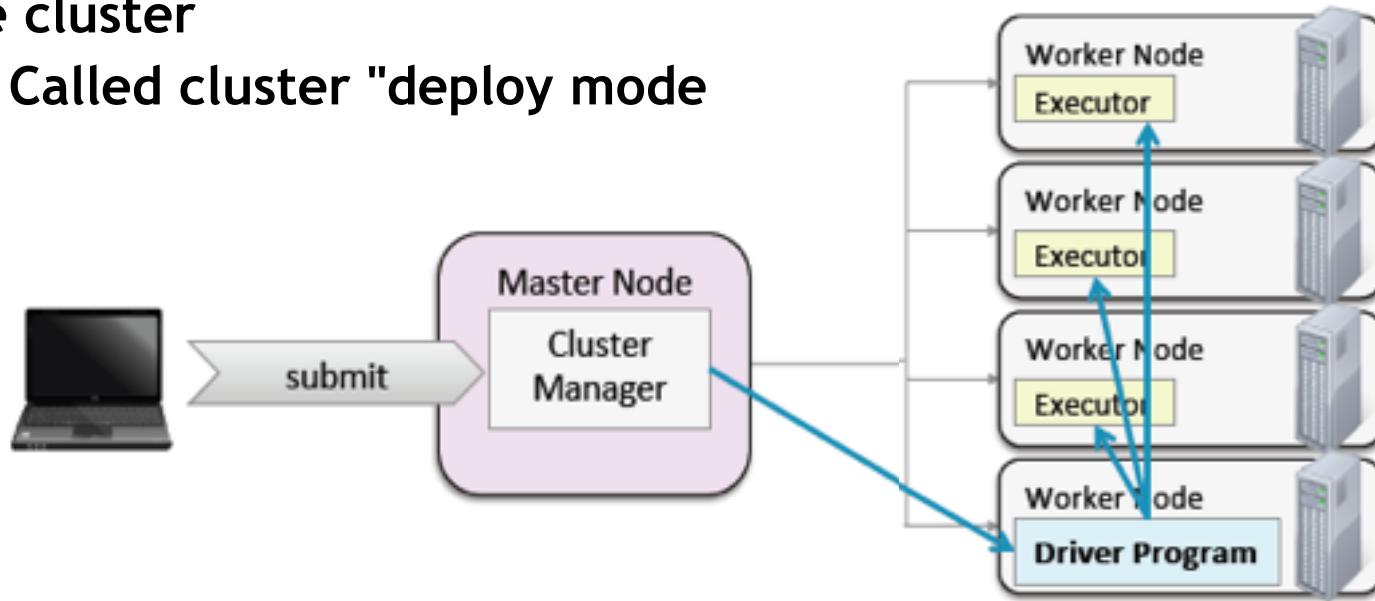
Log files

Supported Cluster Resource Managers

- **Spark Standalone**
 - Included with Spark
 - Easy to install and run
 - Limited configurability and scalability
 - Useful for testing, development, or small systems
- **Hadoop YARN**
 - Included in HDP
 - Most common for production sites
 - Allows sharing cluster resources with other applications (MapReduce, etc.)

Client Mode and Cluster Mode

- By default, the driver program runs outside the cluster
 - Called “client” deploy mode
 - Most common
 - Required for interactive use (e.g., the Spark Shell)
- It is also possible to run the driver program on a worker node in the cluster
 - Called cluster "deploy mode"



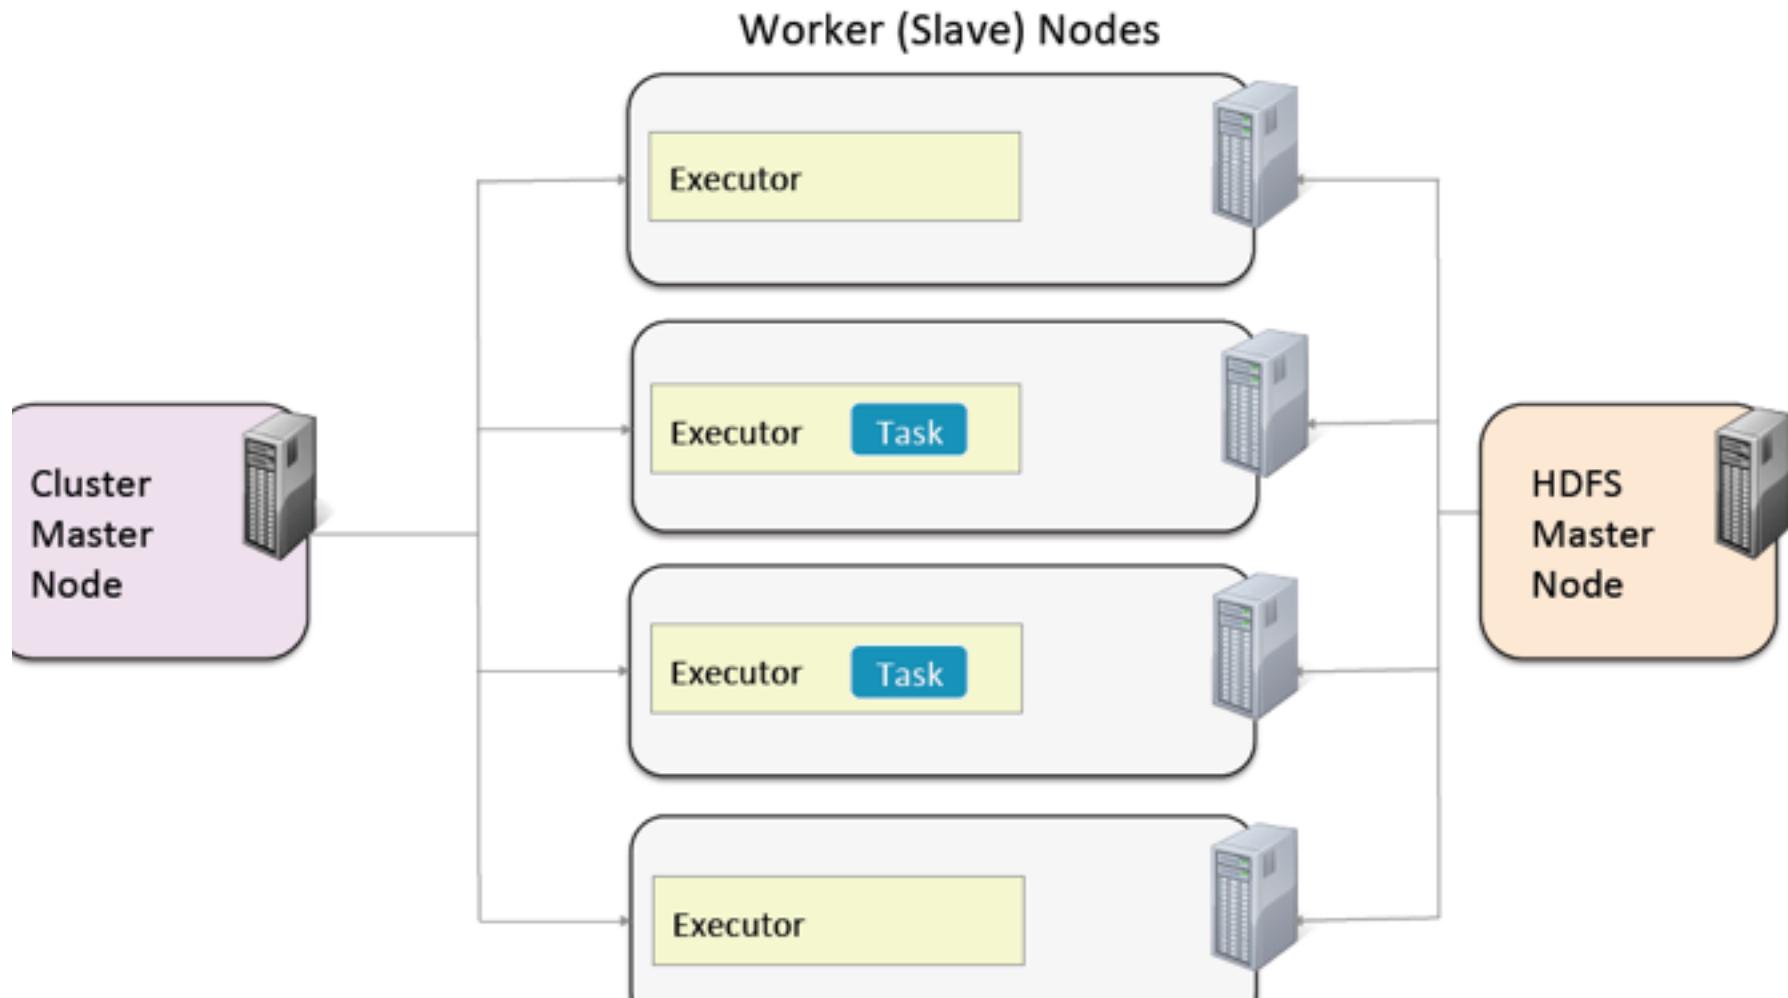
Hands On Exercise: Running Spark on a Cluster

HANDS-ON EXERCISE: RUNNING SPARK ON A CLUSTER

Chapter 6

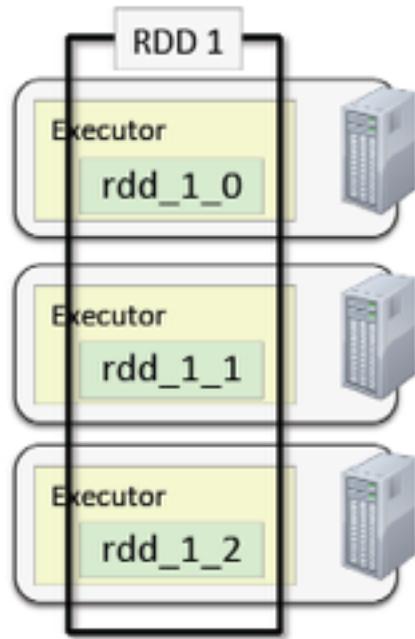
PARALLEL PROGRAMMING WITH SPARK

Spark Cluster Review



RDDs on a Cluster

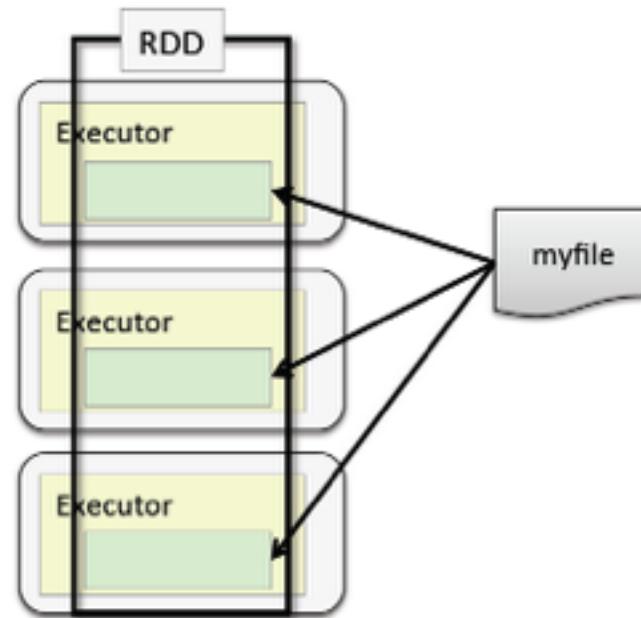
- Resilient Distributed Datasets
 - Data is partitioned across worker nodes
- Partitioning is done automatically by Spark
 - Optionally, you can control how many partitions are created



File Partitioning: Single Files

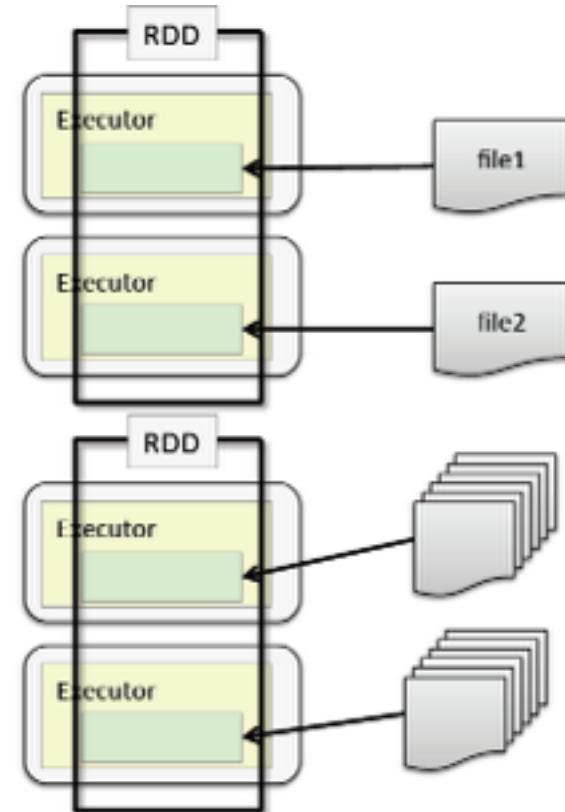
- Partitions from single files
 - Partitions based on size
 - You can optionally specify a minimum number of partitions
`textFile(file, minPartitions)`
 - Default is 2
 - More partitions = more parallelization

```
sc.textFile("myfile", 3)
```



File Partitioning: Multiple Files

- `sc.textFile("mydir/*")`
 - Each file becomes (at least) one partition
 - File-based operations can be done per-partition
- `sc.wholeTextFiles("mydir")`
 - For many small files
 - Creates a key-value PairRDD
 - key = file name
 - value = file contents



Operating on Partitions

- Most RDD operations work on each element of an RDD
- A few work on each partition
 - `foreachPartition` - call a function for each partition
 - `mapPartitions` - create a new RDD by executing a function on each partition in the current RDD
 - `mapPartitionsWithIndex` - same as `mapPartitions` but includes index of the RDD
- Functions for partition operations take iterators

Example: Count JPGs Requests per File

```
def countJpgs(index, partIter):  
    jpgcount = 0  
    for line in partIter:  
        if "jpg" in line: jpgcount += 1  
    yield (index, jpgcount)  
  
jpgcounts = sc.textFile("weblogs/*") \  
.mapPartitionsWithIndex(countJpgs)
```

```
def countJpgs(index: Int, partIter:  
Iterator[String]): Iterator[(Int, Int)] = {  
    var jpgcount = 0  
    for (line <- partIter)  
        if (line.contains("jpg")) jpgcount += 1  
    Iterator((index, jpgcount))  
}  
jpgcounts = sc.textFile("weblogs/*") .  
mapPartitionsWithIndex(countJpgs)
```

Note: Works with small files that each fit in a single partition

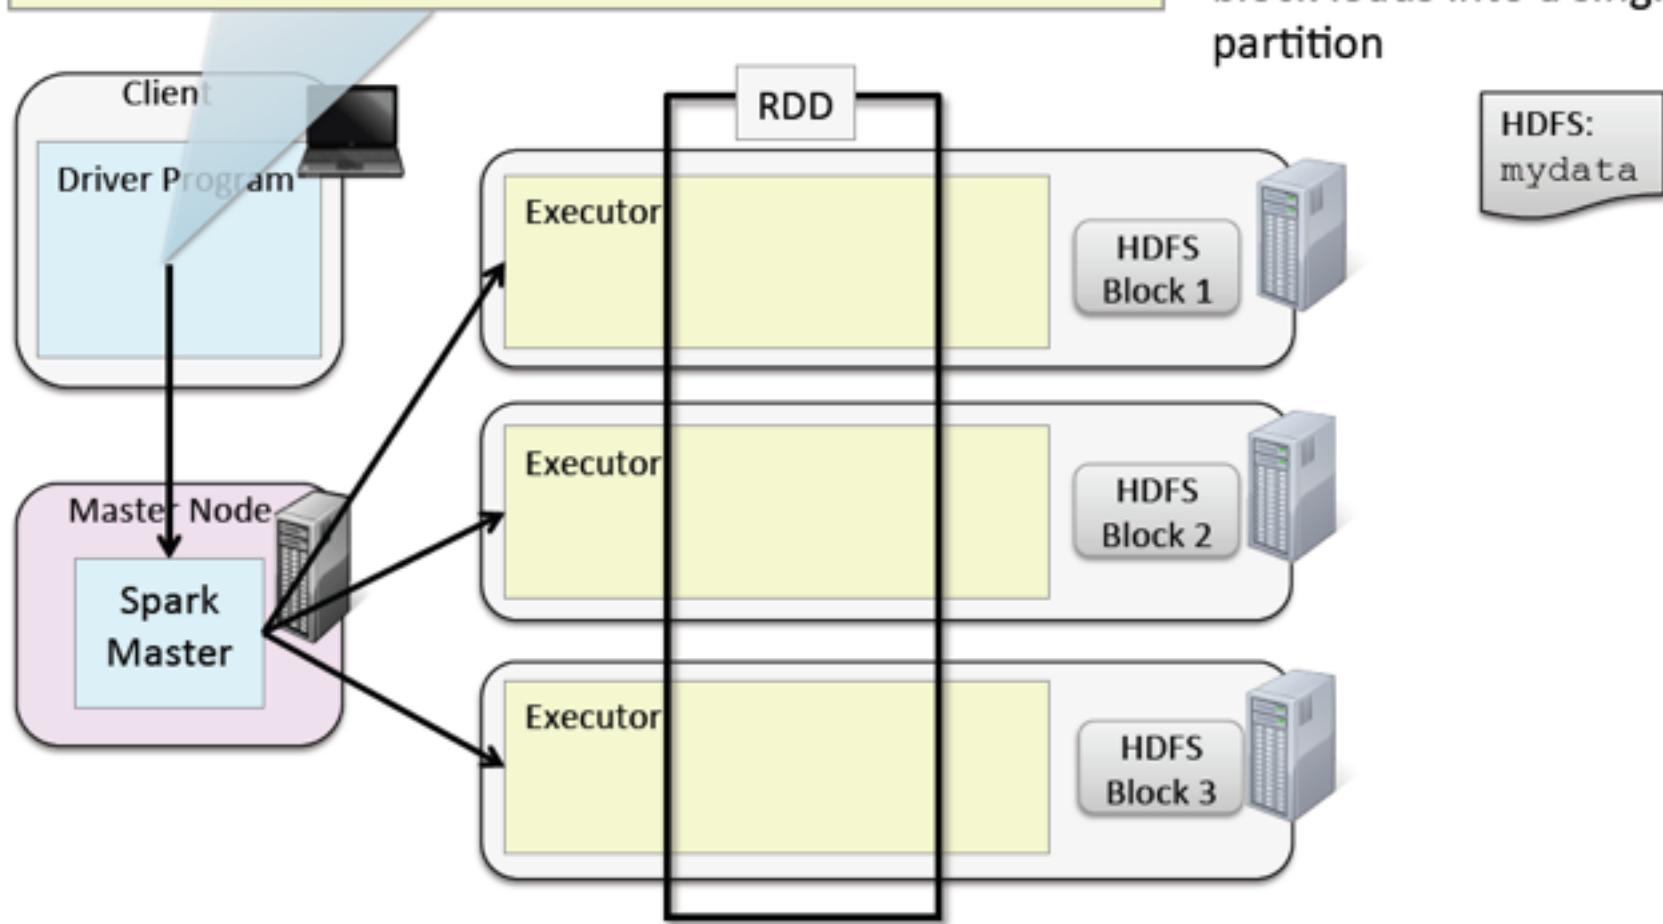
jpgcounts

(0, 237)
(1, 132)
(2, 188)
(3, 193)
...

HDFS and Data Locality

```
sc.textFile("hdfs://...mydata...").collect()
```

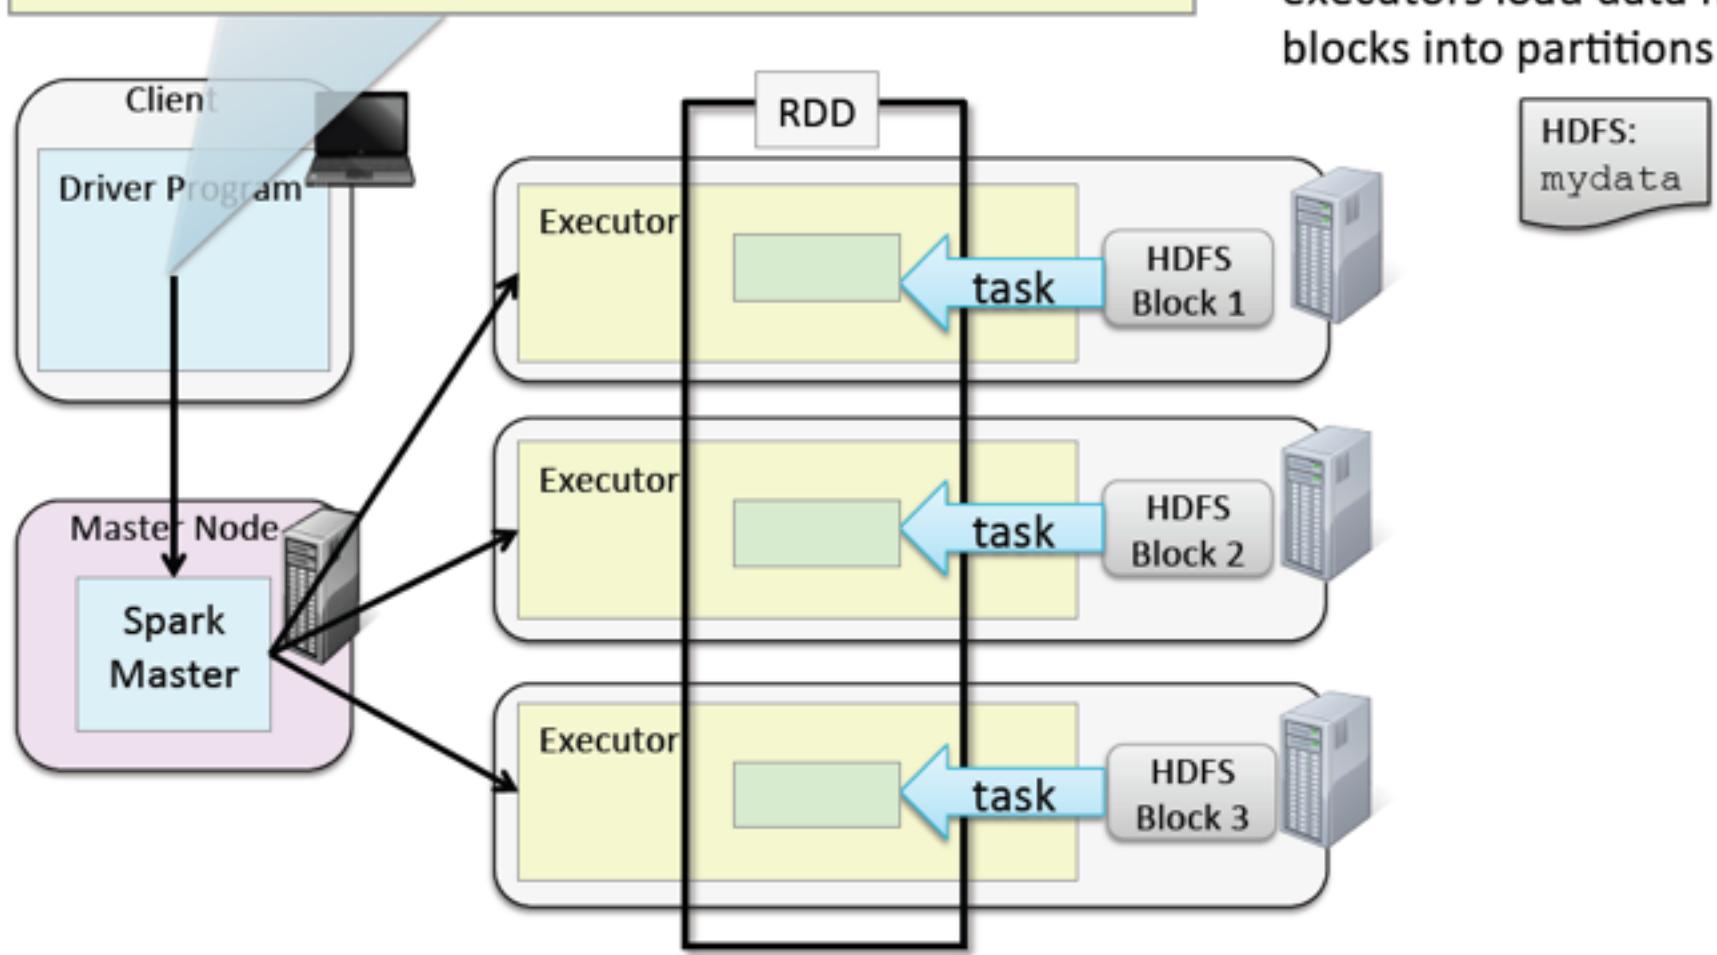
By default, Spark partitions HDFS files by block. Each block loads into a single partition



HDFS and Data Locality

```
sc.textFile("hdfs://...mydata...").collect()
```

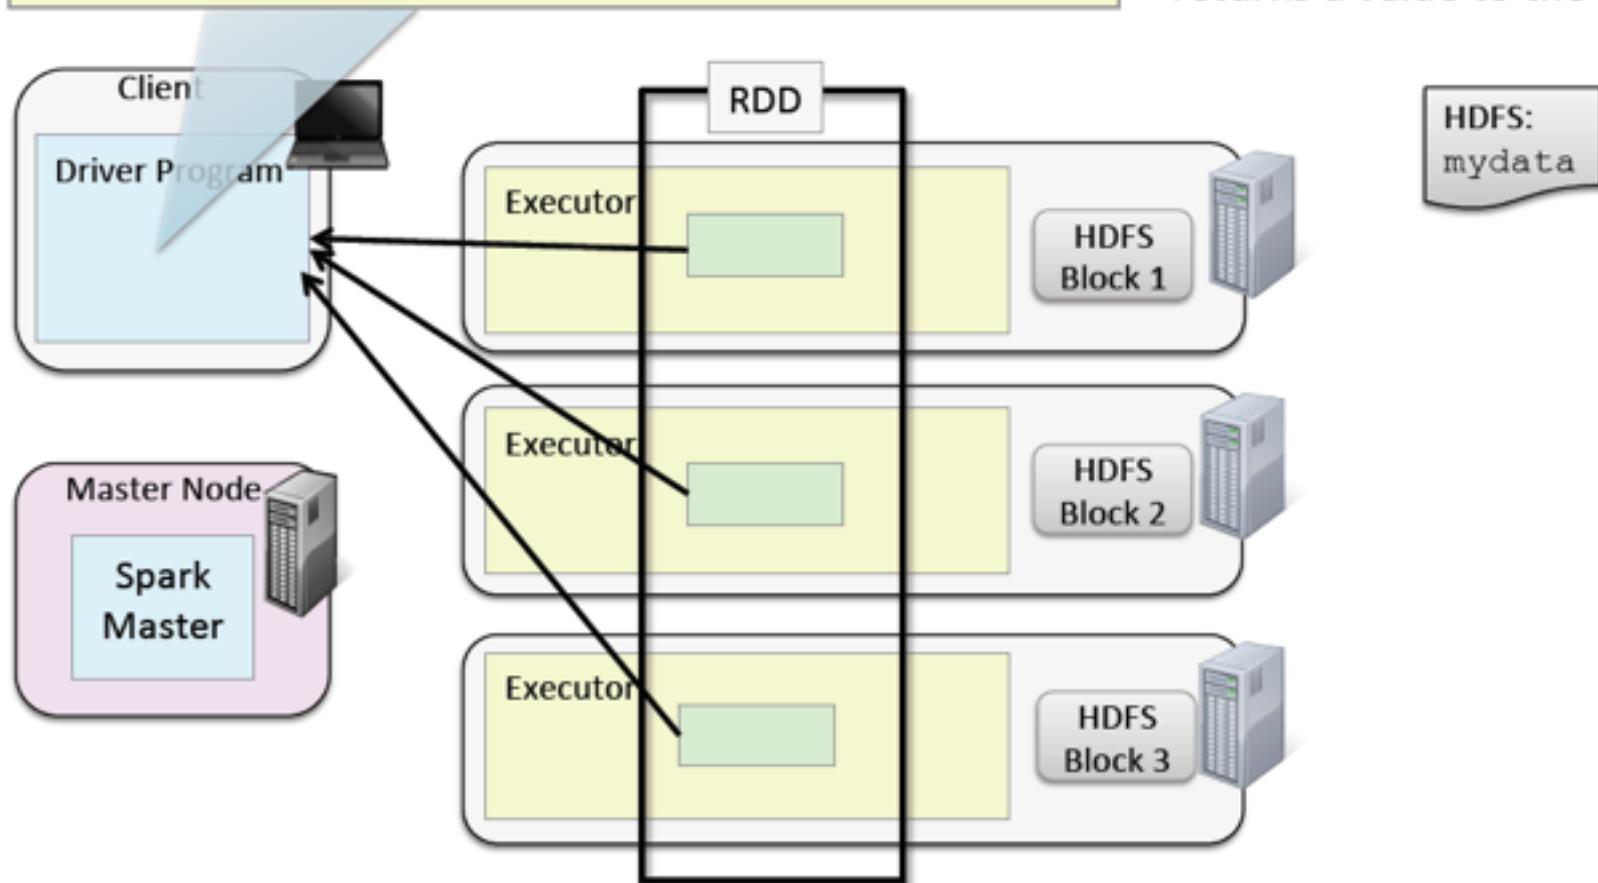
An action triggers execution: tasks on executors load data from blocks into partitions



HDFS and Data Locality

```
sc.textFile("hdfs://...mydata...").collect()
```

Data is distributed across executors until an action returns a value to the driver



Hands On Exercise: Working With Partitions

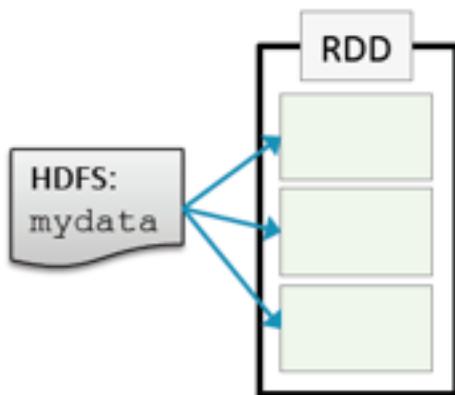
HANDS-ON EXERCISE: WORKING WITH PARTITIONS

Parallel Operations on Partitions

- RDD operations are executed in parallel on each partition
 - When possible, tasks execute on the worker nodes where the data is in memory
- Some operations preserve partitioning
 - e.g., map, flatMap, filter
- Some operations repartition
 - e.g., reduce, sort, group

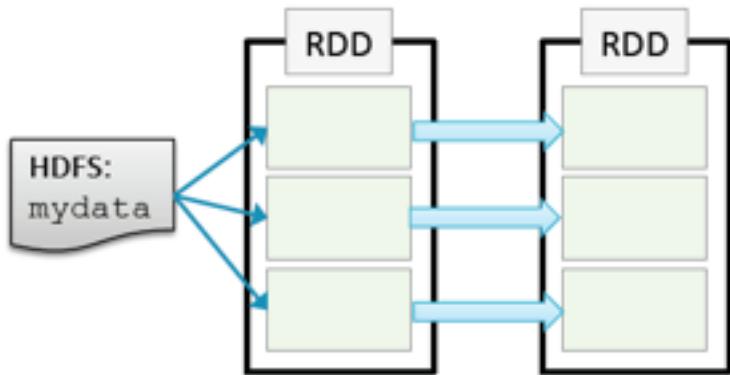
Example: Average Word Length by Letter

- > `avglens = sc.textFile(file)`



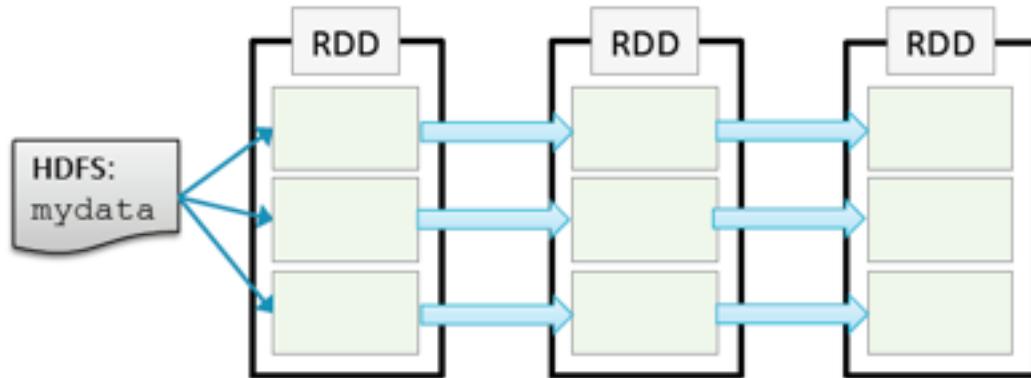
Example: Average Word Length by Letter

```
> avglens = sc.textFile(file) \  
 .flatMap(lambda line: line.split())
```



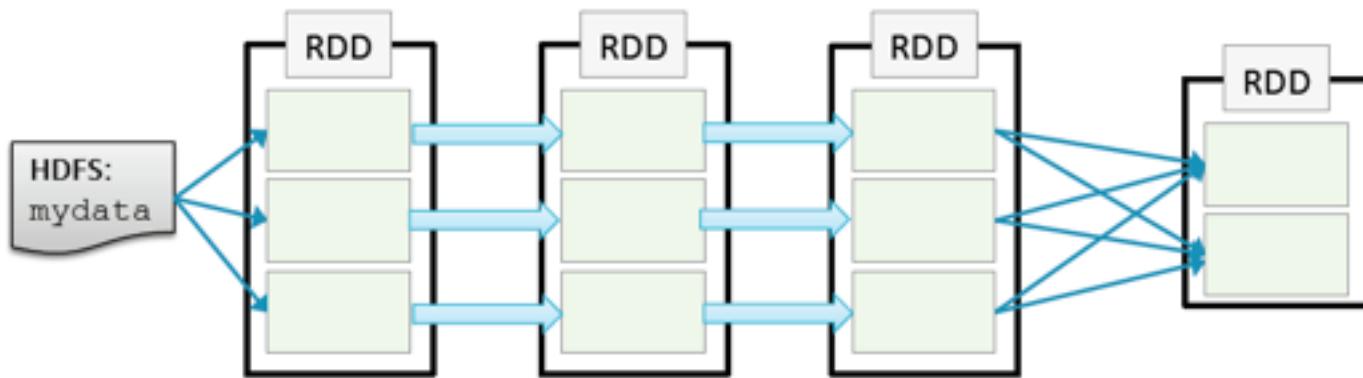
Example: Average Word Length by Letter

```
> avglens = sc.textFile(file) \
  .flatMap(lambda line: line.split()) \
  .map(lambda word: (word[0], len(word)))
```



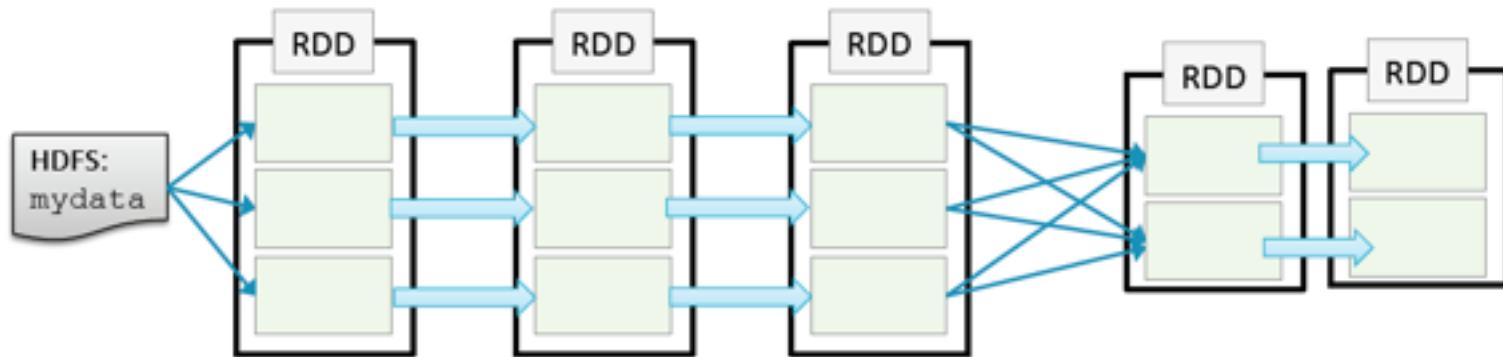
Example: Average Word Length by Letter

```
> avglens = sc.textFile(file) \
  .flatMap(lambda line: line.split()) \
  .map(lambda word: (word[0],len(word))) \
  .groupByKey()
```



Example: Average Word Length by Letter

```
> avglens = sc.textFile(file) \
    .flatMap(lambda line: line.split()) \
    .map(lambda word: (word[0],len(word))) \
    .groupByKey() \
    .map(lambda (k, values): \
        (k, sum(values)/len(values)))
```

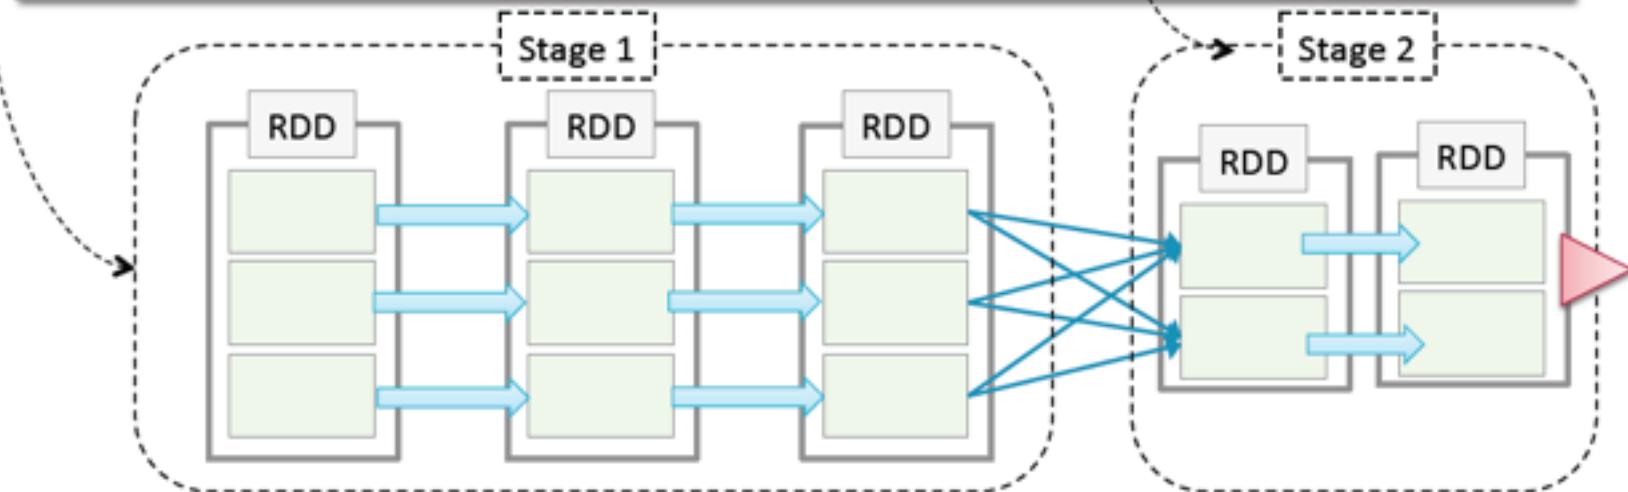


Stages

- Operations that can run on the same partition are executed in stages
- Tasks within a stage are pipelined together
- Developers should be aware of stages to improve performance

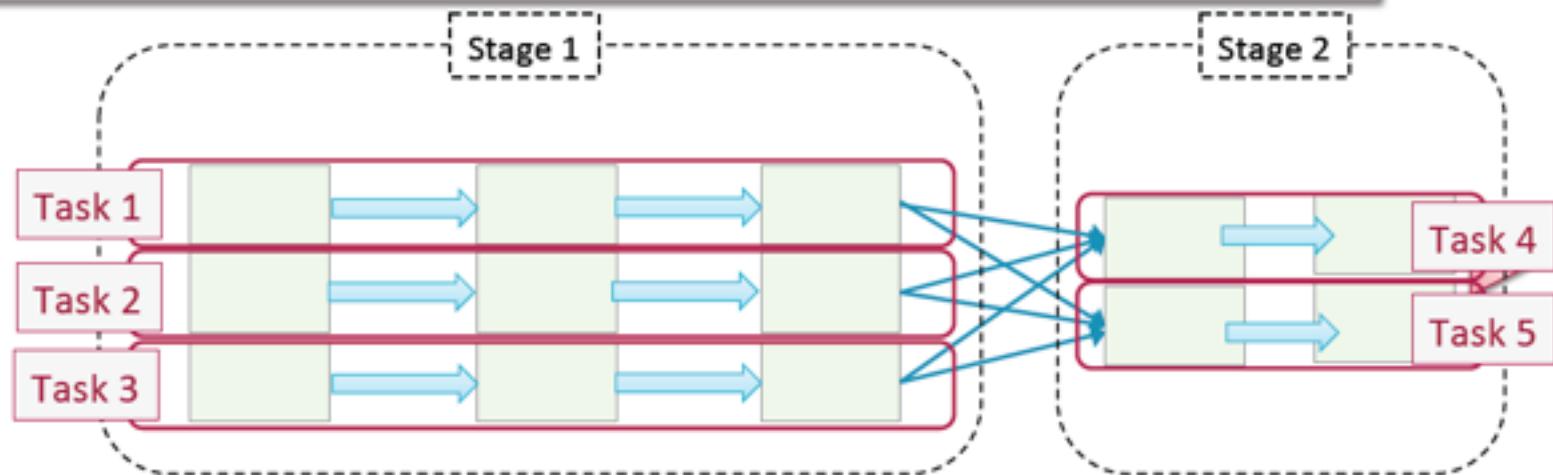
Spark Execution: Stages

```
> avglen = sc.textFile(file) \
  .flatMap(lambda line: line.split()) \
  .map(lambda word: (word[0], len(word))) \
  .groupByKey() \
  .map(lambda (k, values): \
        (k, sum(values)/len(values)))
> avglen.count()
```



Spark Execution: Stages

- > avglen = sc.textFile(file) \
.flatMap(lambda line: line.split()) \
.map(lambda word: (word[0], len(word))) \
.groupByKey() \
.map(lambda (k, values): \
 (k, sum(values)/len(values)))
> avglen.count()



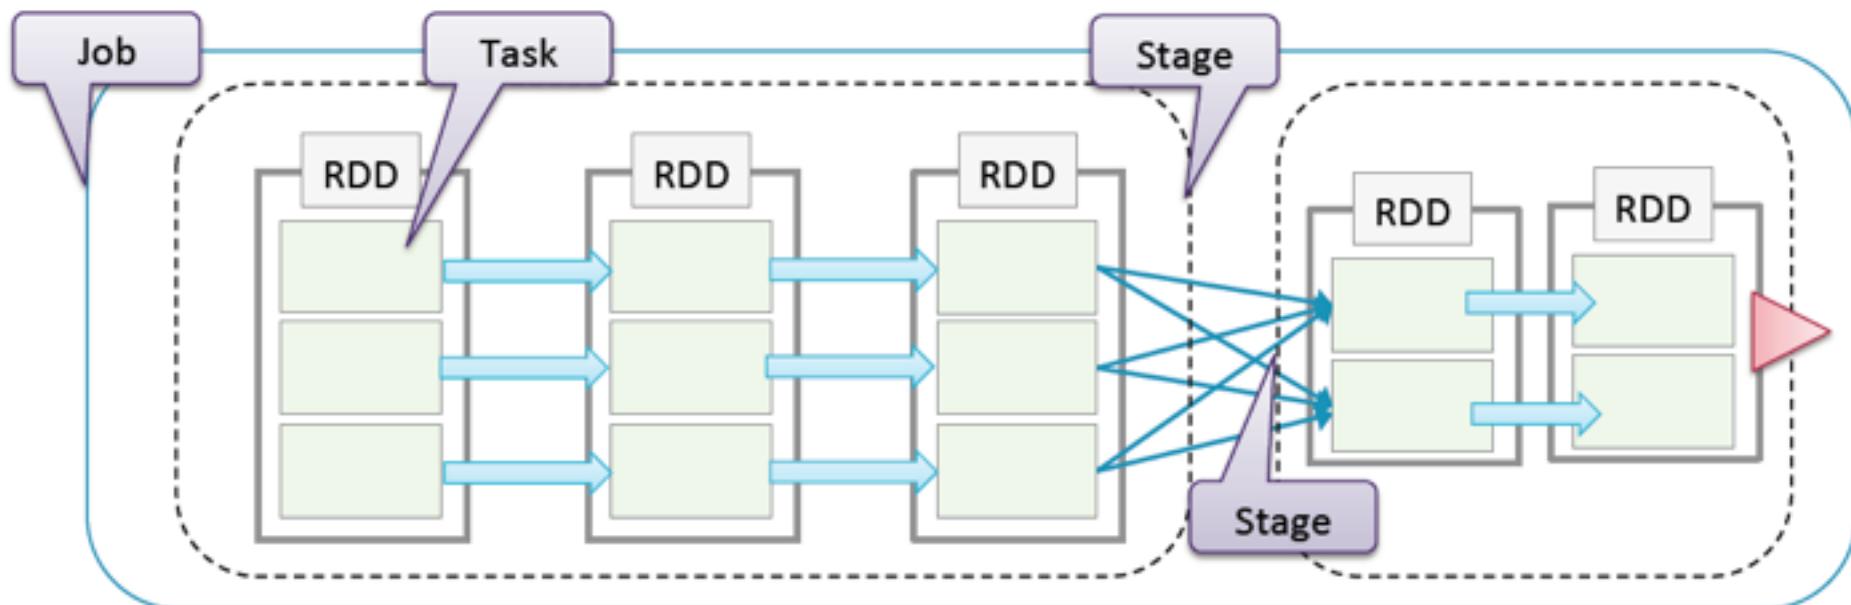
Controlling the Level of Parallelism

- RDD operations that repartition data (e.g., `reduceByKey`) take an optional additional parameter for the number of partitions/tasks
 - By default the number of cores in the cluster, or local threads
 - Configure with the `spark.default.parallelism` property

```
words.reduceByKey(lambda v1, v2: v1 + v2, 15)
```

Summary of Spark Terminology

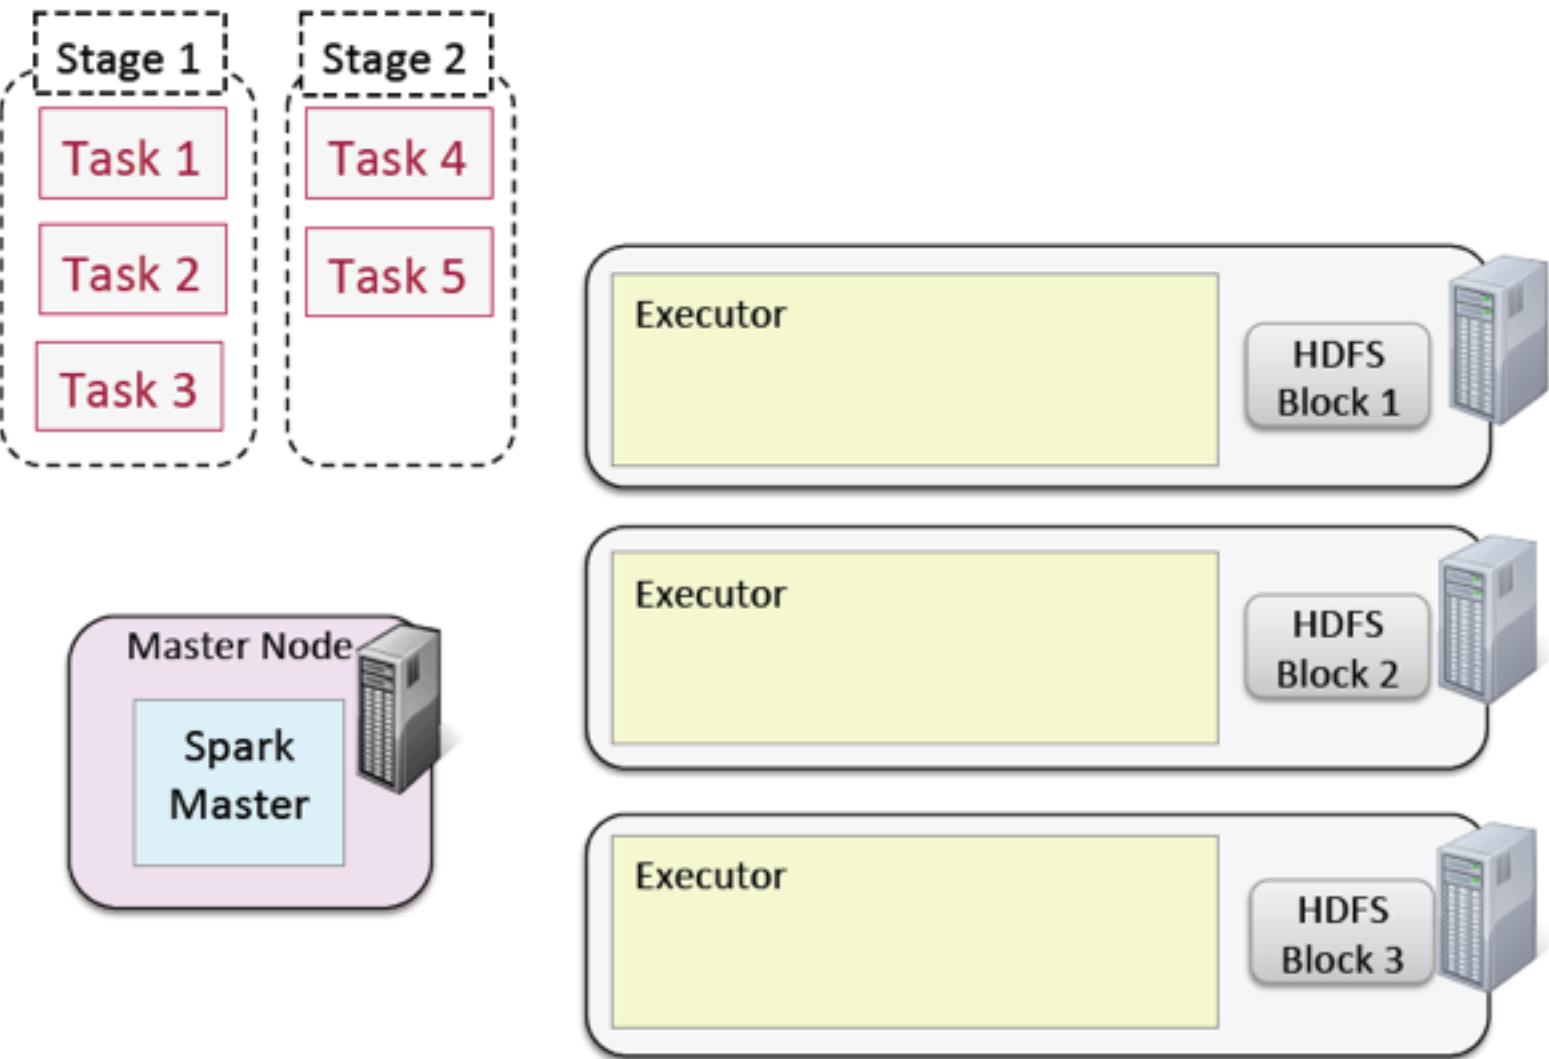
- Job - a set of tasks executed as a result of an action
- Stage - a set of tasks in a job that can be executed in parallel
- Task - an individual unit of work sent to one executor



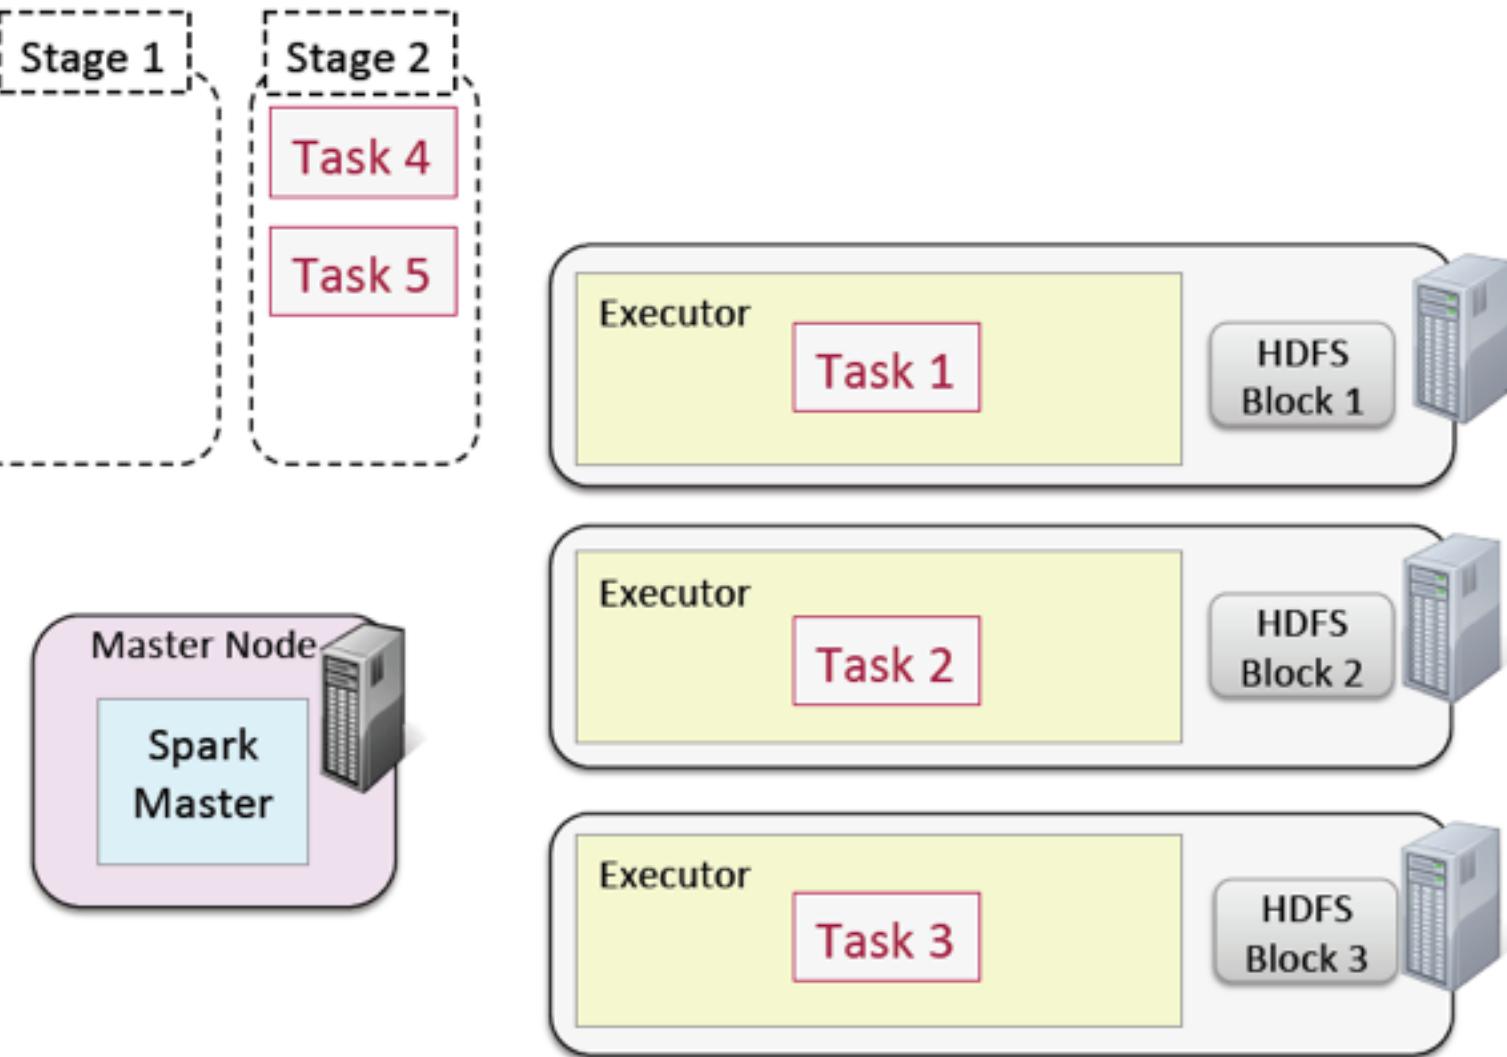
How Spark Calculates Stages

- Spark constructs a DAG (Directed Acyclic Graph) of RDD dependencies
- Narrow operations
 - Only one child depends on the RDD
 - No shuffle required between nodes
 - Can be collapsed into a single stage
 - e.g., map, filter, union
- Wide operations
 - Multiple children depend on the RDD
 - Defines a new stage
 - e.g., reduceByKey, join, groupByKey

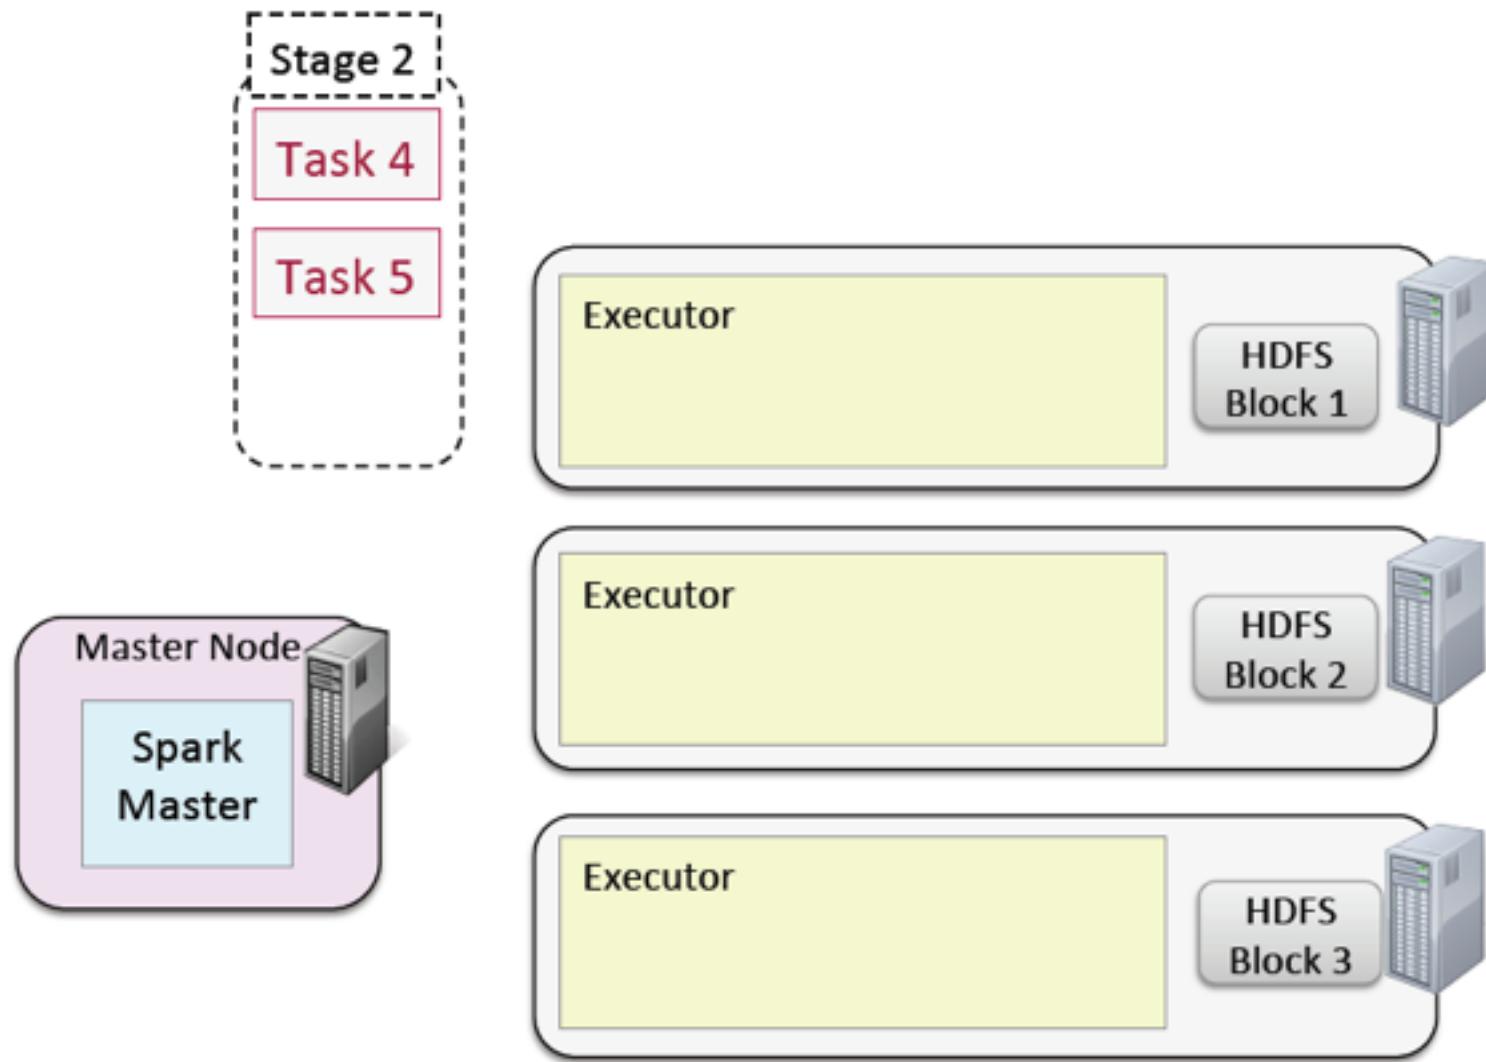
Spark Execution: Task Scheduling



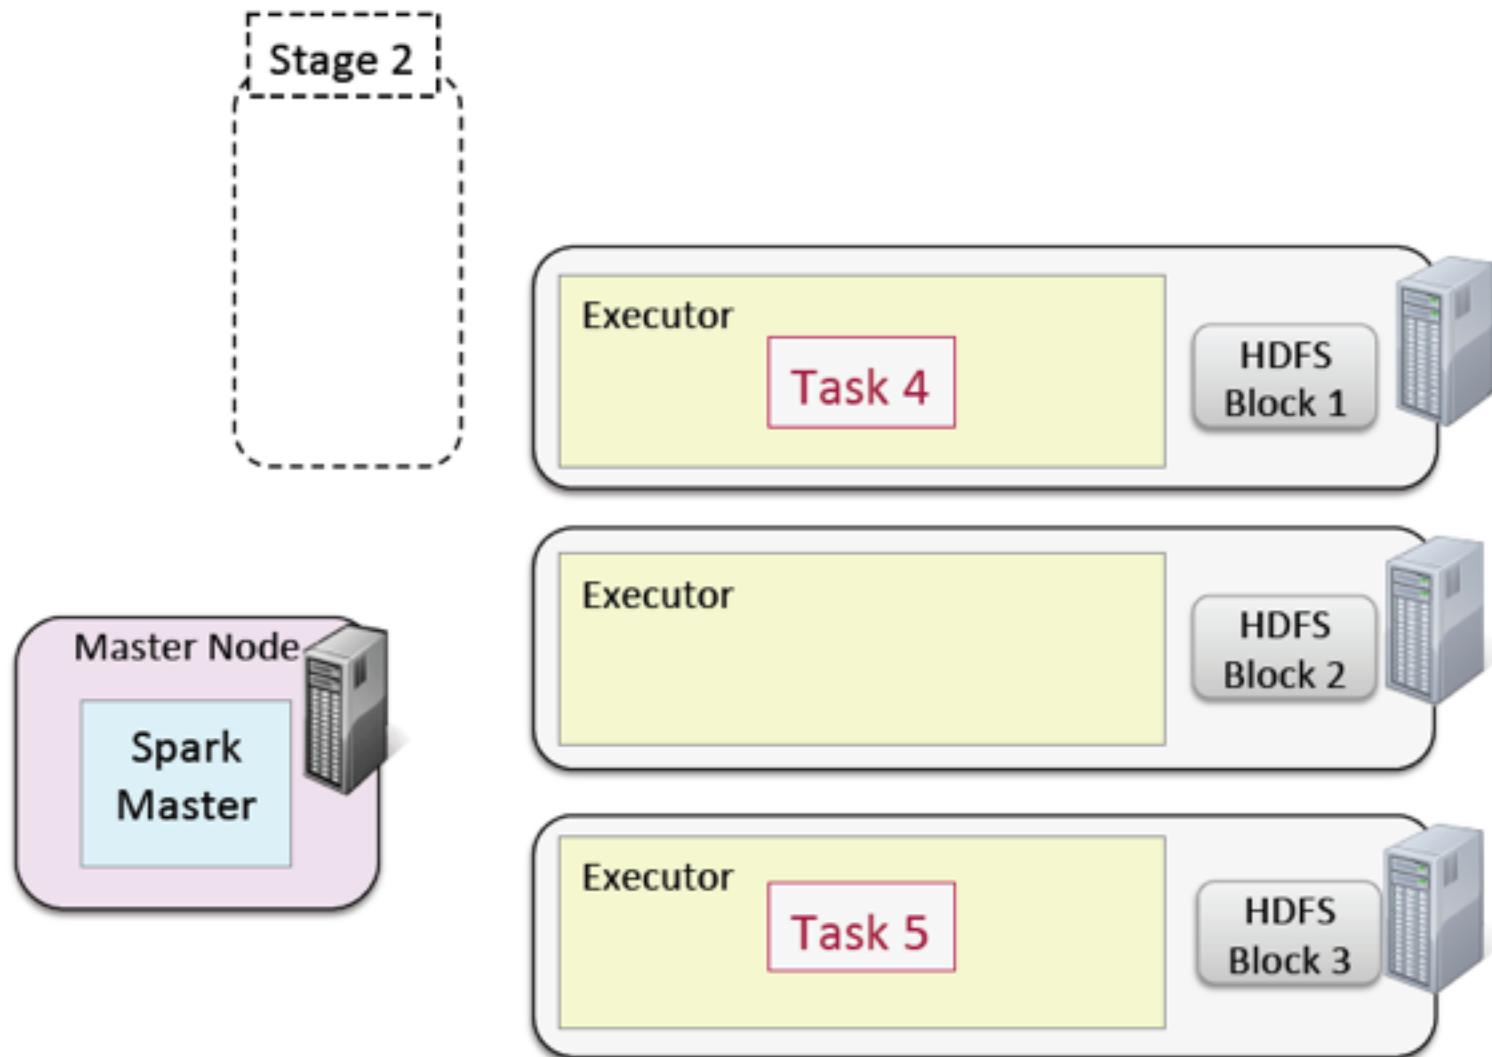
Spark Execution: Task Scheduling



Spark Execution: Task Scheduling



Spark Execution: Task Scheduling



Viewing Stages in the Spark Application UI

Spark Stages

Total Duration: 3.3 m
Scheduling Mode: FIFO
Active Stages: 0
Completed Stages: 2
Failed Stages: 0

Active Stages (0)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Shuffle Read	Shuffle Write
----------	-------------	-----------	----------	------------------------	--------------	---------------

Completed Stages (2)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Shuffle Read	Shuffle Write
0	count at <ipython-input-5-2b2806b2dd0d>:1	2014/04/30 10:41:21	224 ms	2/2		
1	groupByKey at <ipython-input-4-0d0196eb57d4>:1	2014/04/30 10:41:18	3.5 s	3/3		91.8 KB

Failed Stages (0)

Stage Id	Description	Submitted	Duration	Tasks: Succeeded/Total	Shuffle Read	Shuffle Write
----------	-------------	-----------	----------	------------------------	--------------	---------------

Stages are identified by the last operation

Number of tasks = number of partitions

Data shuffled between stages

134

Hands On Exercise: Viewing Stages and Tasks in the Spark Application UI

**HANDS-ON EXERCISE: VIEWING STAGES AND
TASKS IN THE SPARK
APPLICATION UI**

Chapter 7

CACHING AND PERSISTENCE

Example

- Each transformation operation
 - creates a new child RDD

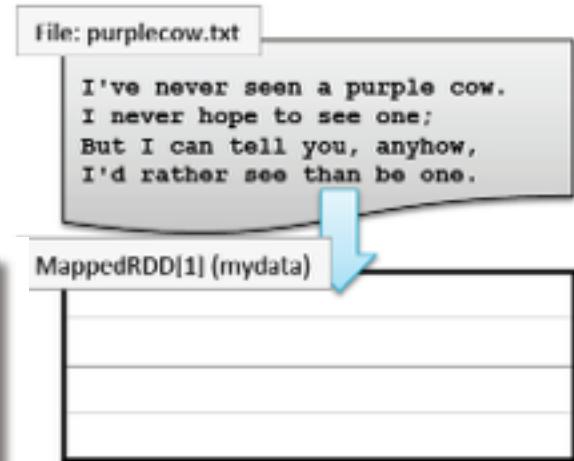
File: purplecow.txt

I've never seen a purple cow.
I never hope to see one;
But I can tell you, anyhow,
I'd rather see than be one.

Example

- Each transformation operation
 - creates a new child RDD

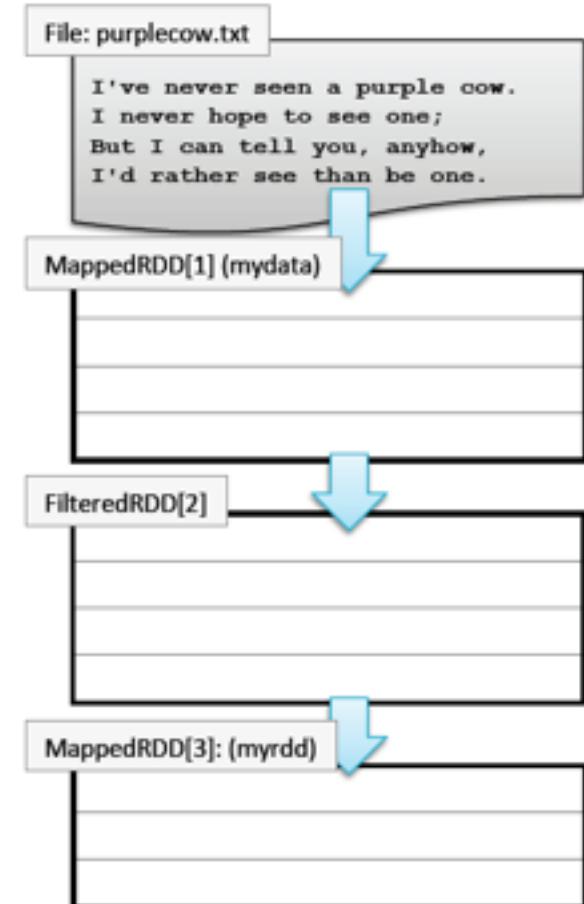
```
> mydata = sc.textFile("purplecow.txt")
```



Example

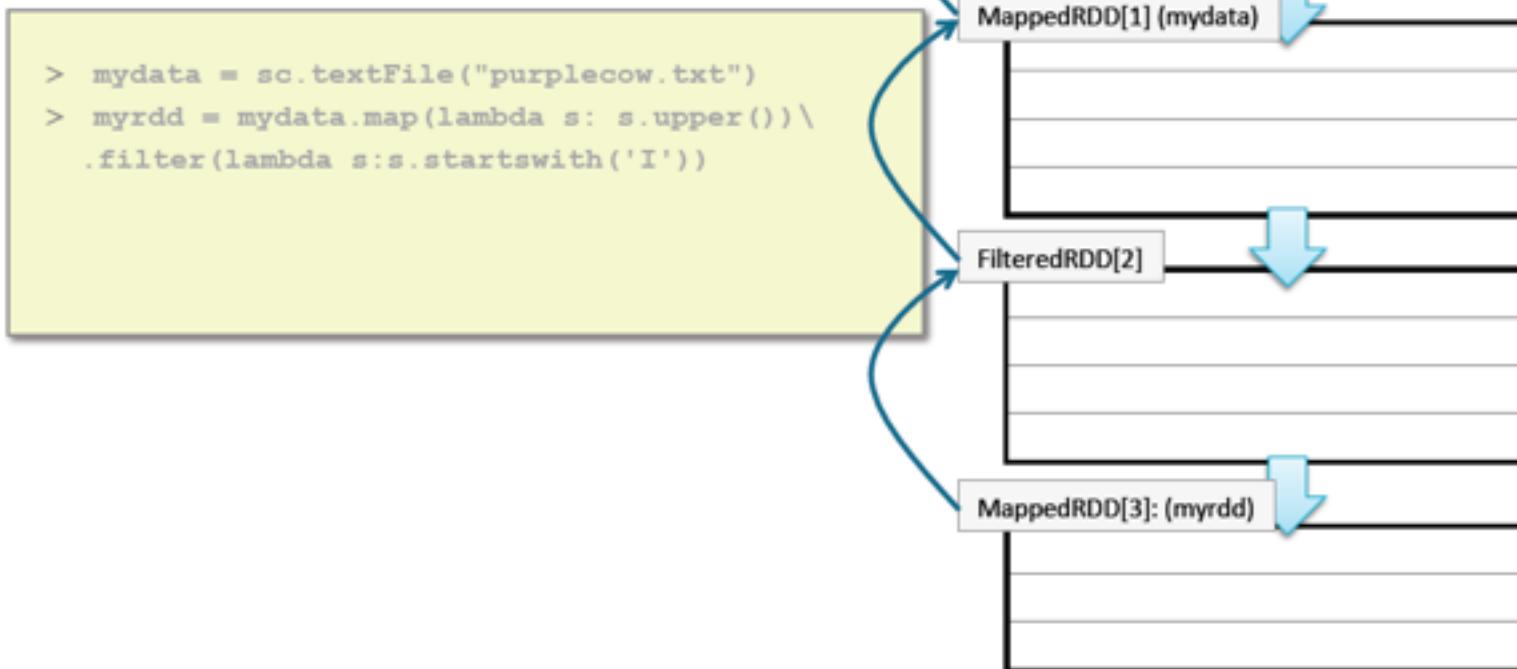
- Each transformation operation
 - creates a new child RDD

```
> mydata = sc.textFile("purplecow.txt")
> myrdd = mydata.map(lambda s: s.upper()) \
    .filter(lambda s:s.startswith('I'))
```



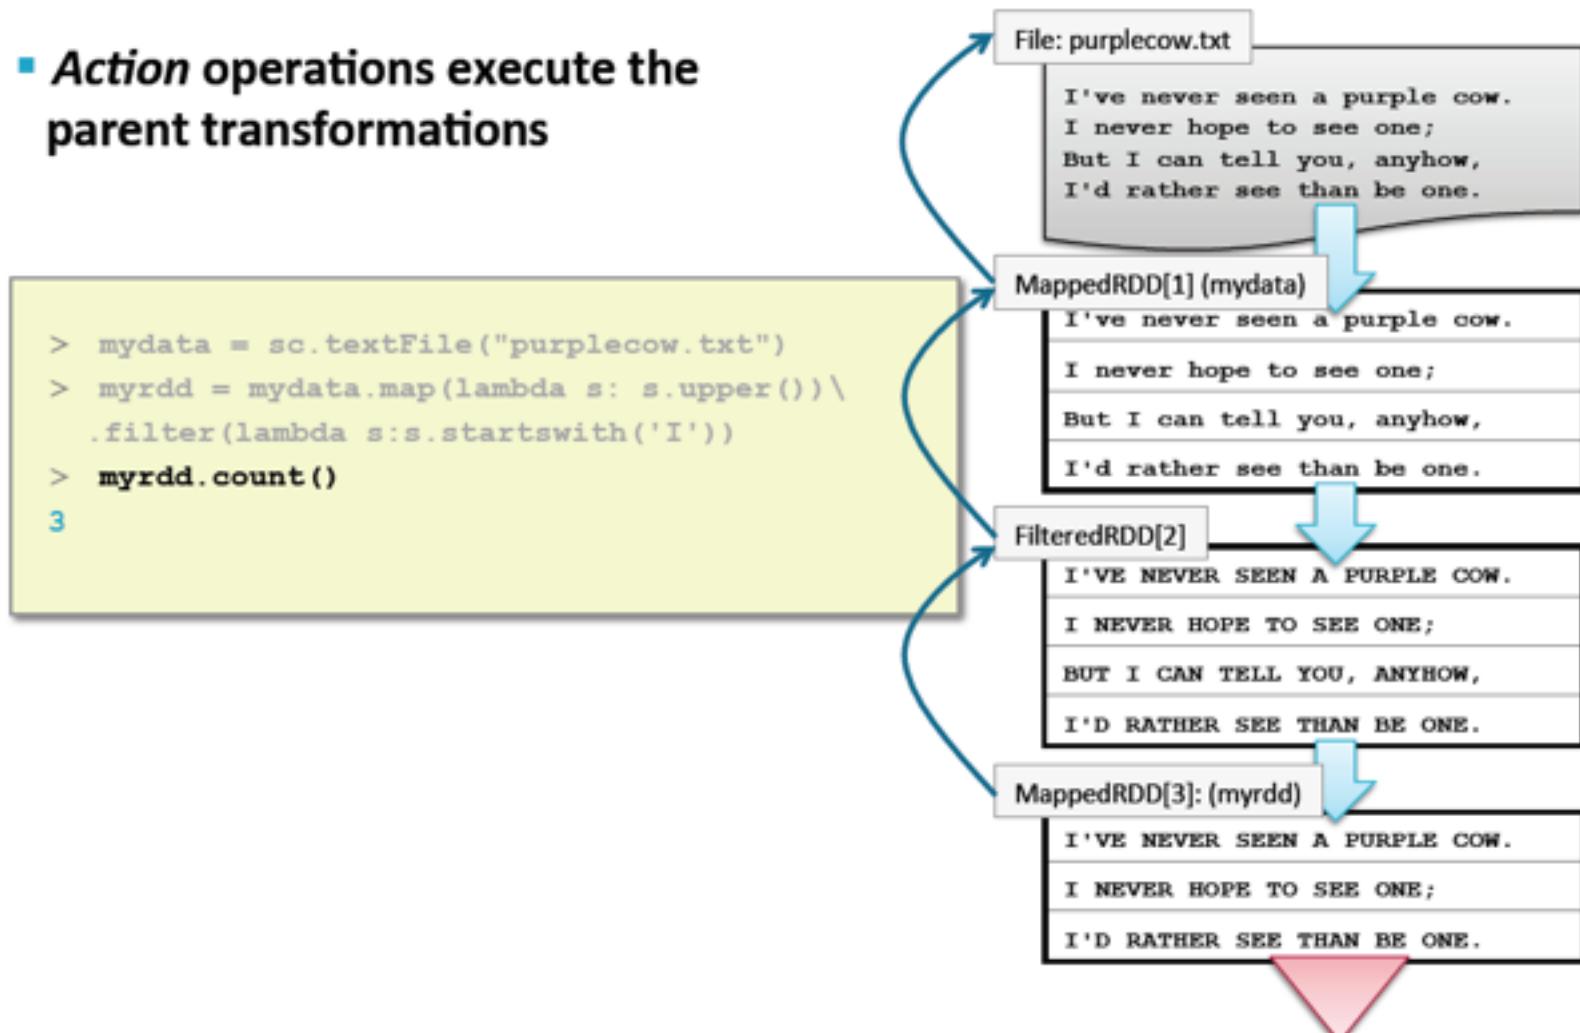
Example

- Spark keeps track of the *parent RDD* for each new RDD
- Child RDDs *depend on* their parents



Lineage Example

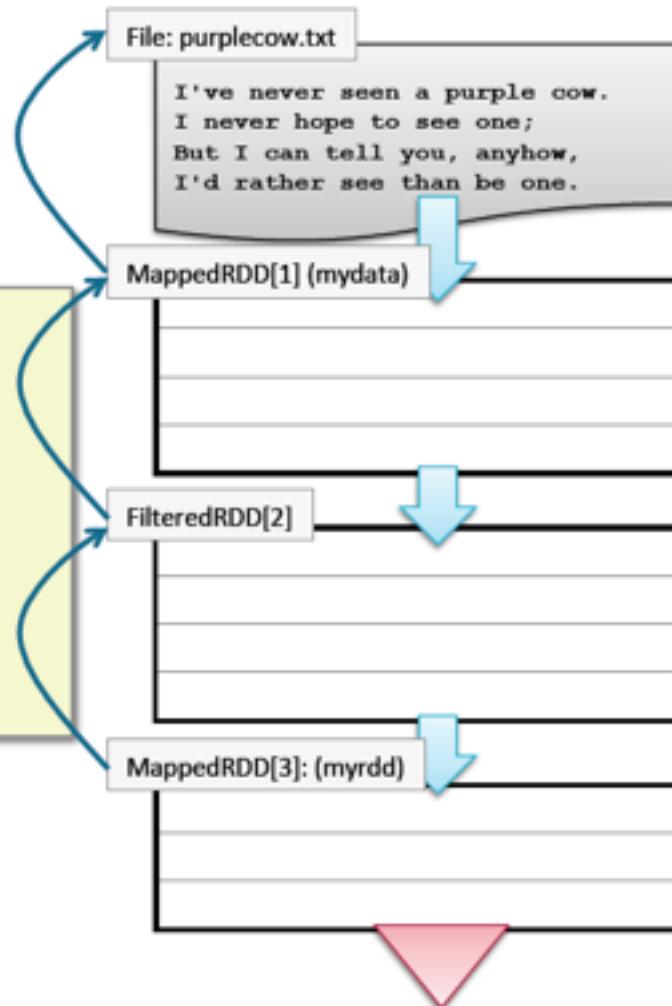
- **Action operations execute the parent transformations**



Lineage Example

- Each action re-executes the lineage transformations starting with the base
 - By default

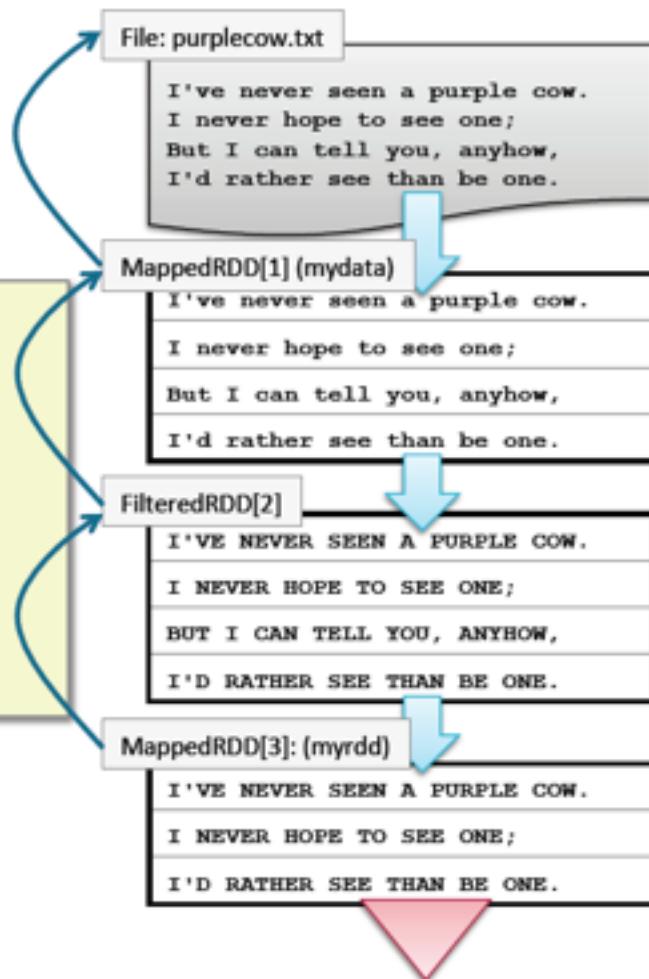
```
> mydata = sc.textFile("purplecow.txt")
> myrdd = mydata.map(lambda s: s.upper()) \
    .filter(lambda s:s.startswith('I'))
> myrdd.count()
3
> myrdd.count()
```



Lineage Example

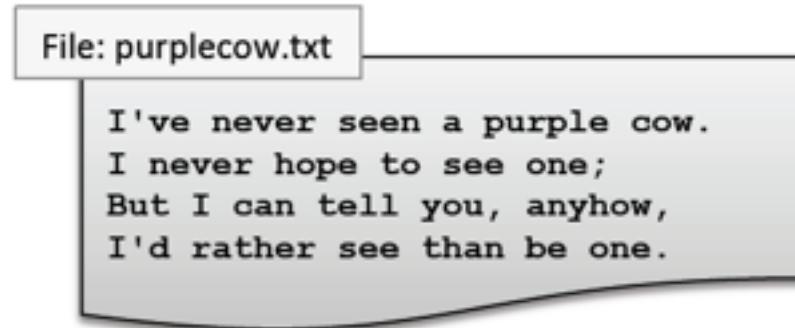
- **Each action re-executes the lineage transformations starting with the base**
 - By default

```
> mydata = sc.textFile("purplecow.txt")
> myrdd = mydata.map(lambda s: s.upper()) \
    .filter(lambda s:s.startswith('I'))
> myrdd.count()
3
> myrdd.count()
3
```



Caching

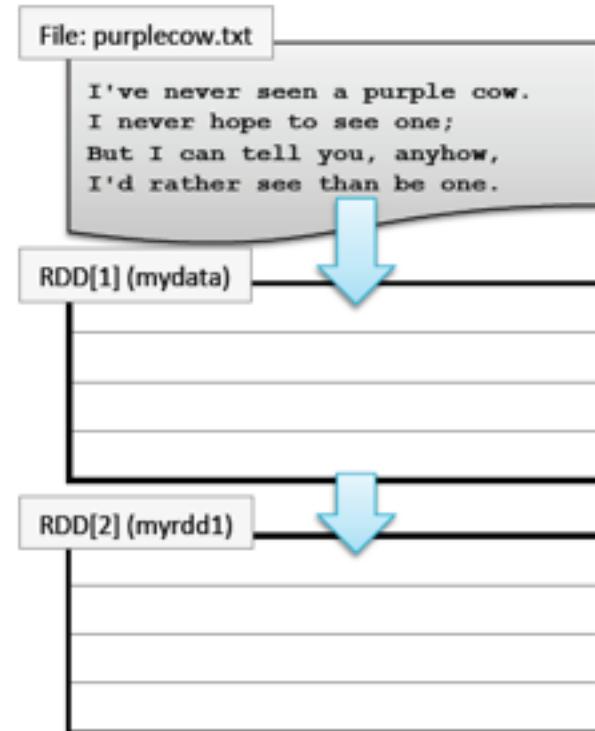
- Caching an RDD saves the data in memory



Caching

- Caching an RDD saves the data in memory

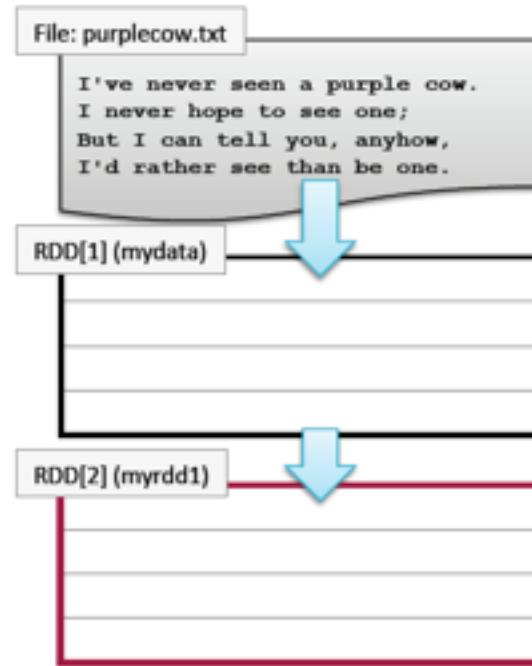
```
> mydata = sc.textFile("purplecow.txt")
> myrdd1 = mydata.map(lambda s:
    s.upper())
```



Caching

- Caching an RDD saves the data in memory

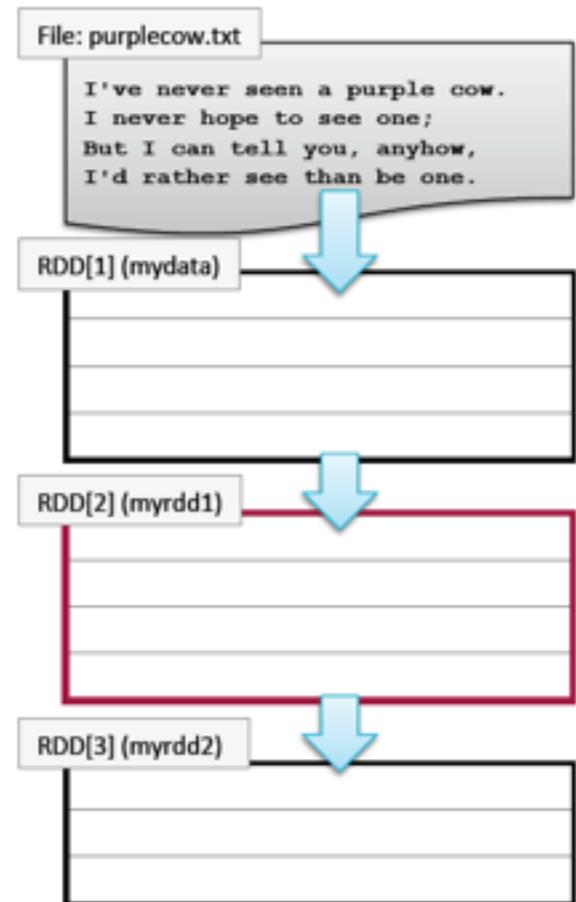
```
> mydata = sc.textFile("purplecow.txt")
> myrdd1 = mydata.map(lambda s:
>     s.upper())
> myrdd1.cache()
```



Caching

- Caching an RDD saves the data in memory

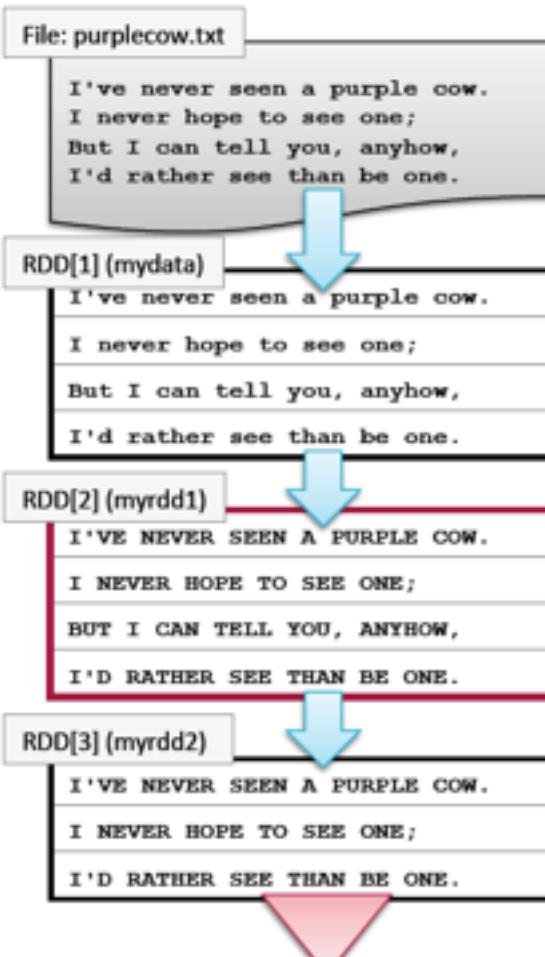
```
> mydata = sc.textFile("purplecow.txt")
> myrdd1 = mydata.map(lambda s:
  s.upper())
> myrdd1.cache()
> myrdd2 = myrdd1.filter(lambda \
  s:s.startswith('I'))
```



Caching

- Caching an RDD saves the data in memory

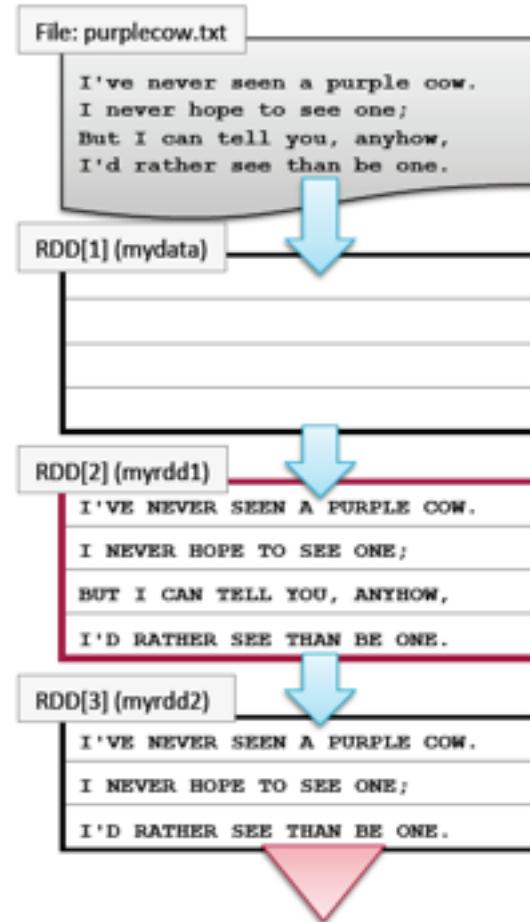
```
> mydata = sc.textFile("purplecow.txt")
> myrdd1 = mydata.map(lambda s:
>   s.upper())
> myrdd1.cache()
> myrdd2 = myrdd1.filter(lambda \
>   s:s.startswith('I'))
> myrdd2.count()
3
```



Caching

- Subsequent operations use saved data

```
> mydata = sc.textFile("purplecow.txt")
> myrdd1 = mydata.map(lambda s:
  s.upper())
> myrdd1.cache()
> myrdd2 = myrdd1.filter(lambda \
  s:s.startswith('I'))
> myrdd2.count()
3
> myrdd2.count()
3
```



Caching

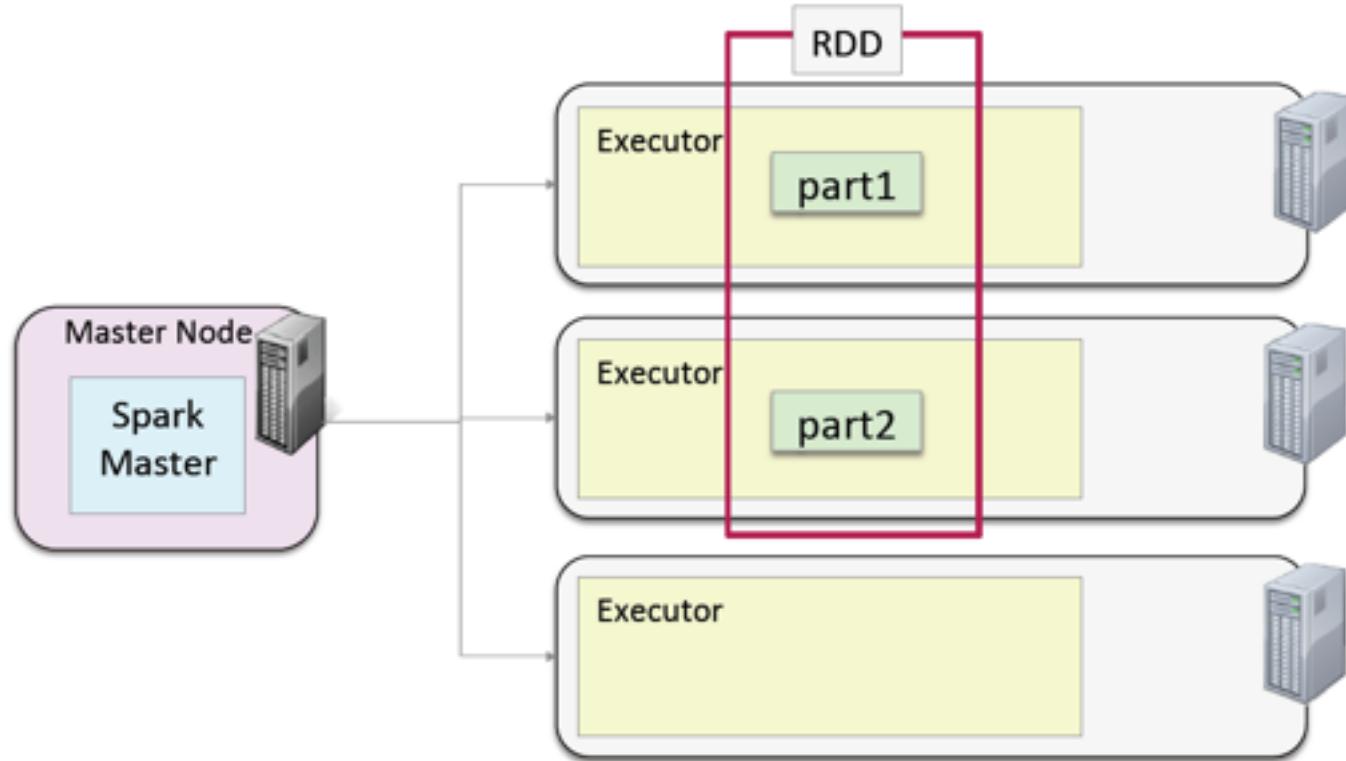
- Caching is a suggestion to Spark
 - If not enough memory is available, transformations will be re-executed when needed

Caching and Fault-Tolerance

- **RDD = Resilient Distributed Dataset**
 - Resiliency is a product of tracking lineage
 - RDDs can always be recomputed from their base if needed

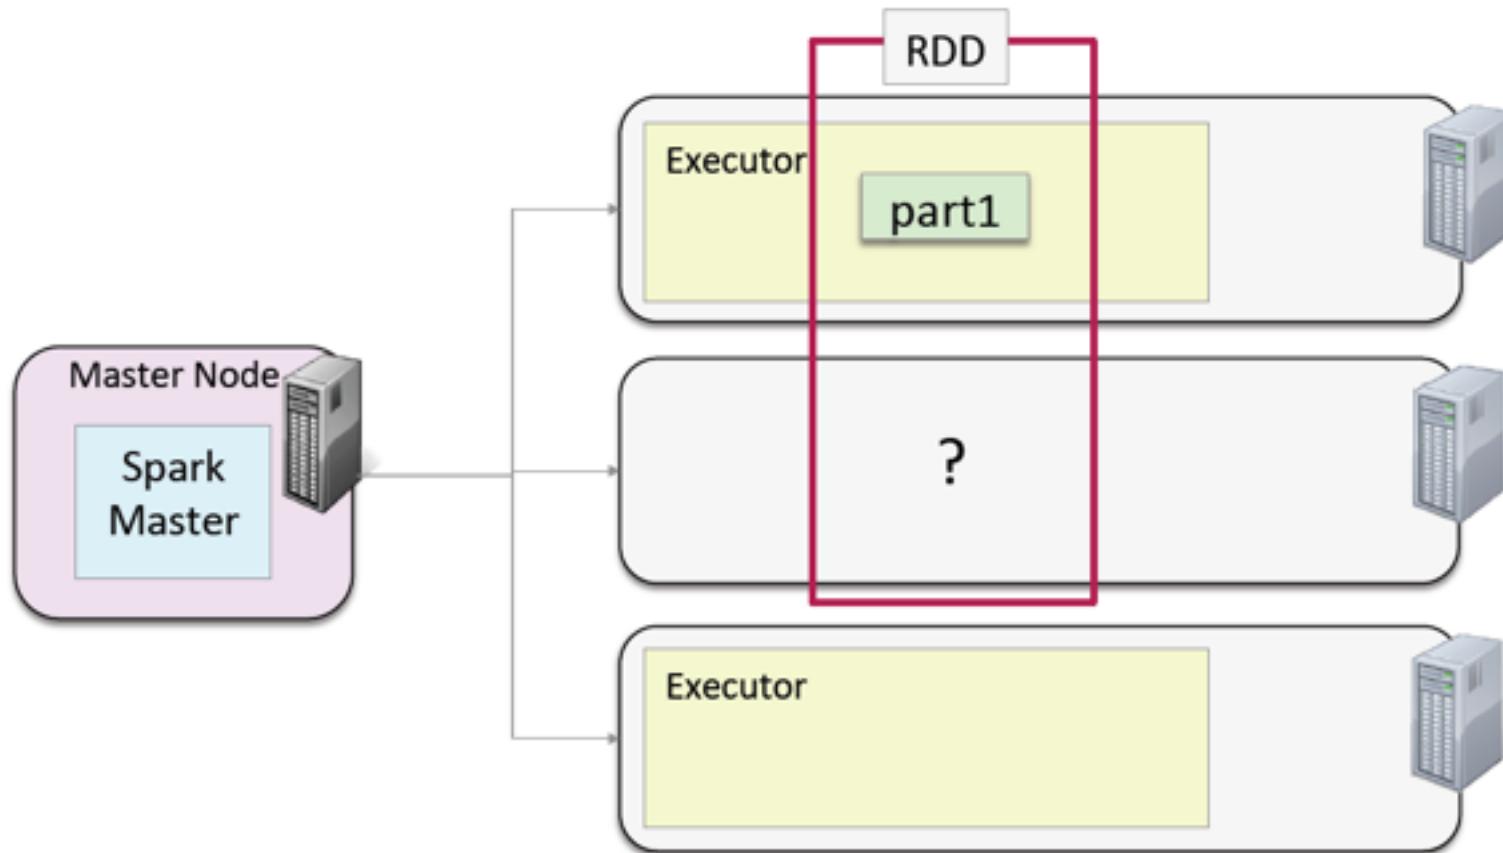
Distributed Cache

- RDD partitions are distributed across a cluster
- Cached partitions are stored in memory in Executor JVMs



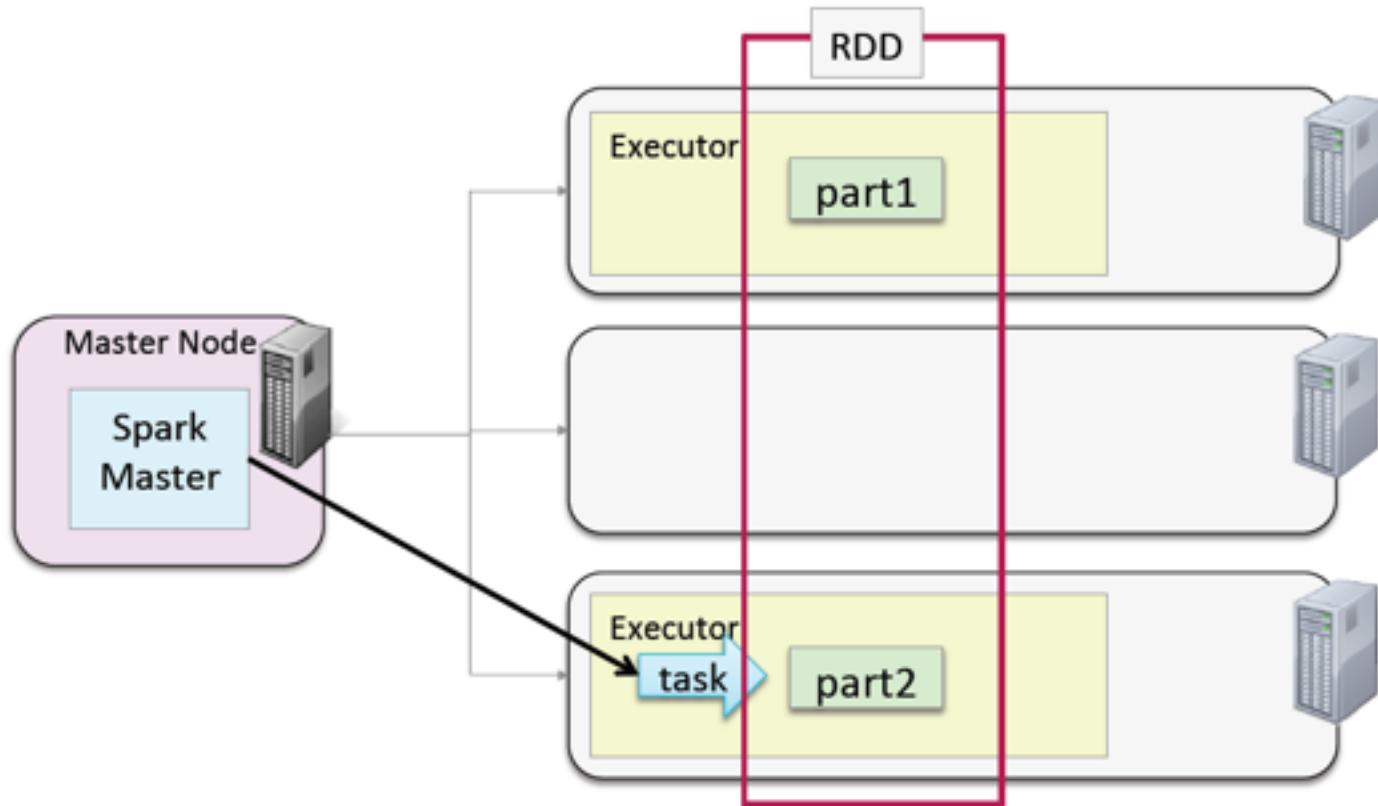
RDD Fault-Tolerance

- What happens if a cached partition becomes unavailable?



RDD Fault-Tolerance

- The SparkMaster starts a new task to recompute the partition on a different node



Persistence Levels

- The cache method stores data in memory only
- The persist method offers other Storage Levels
 - MEMORY_ONLY (default) - same as cache
 - MEMORY_AND_DISK - Store partitions on disk if they do not fit in memory
 - Called spilling
 - DISK_ONLY - Store all partitions on disk

```
> myrdd.persist(StorageLevel.DISK_ONLY)
```

Persistence Levels

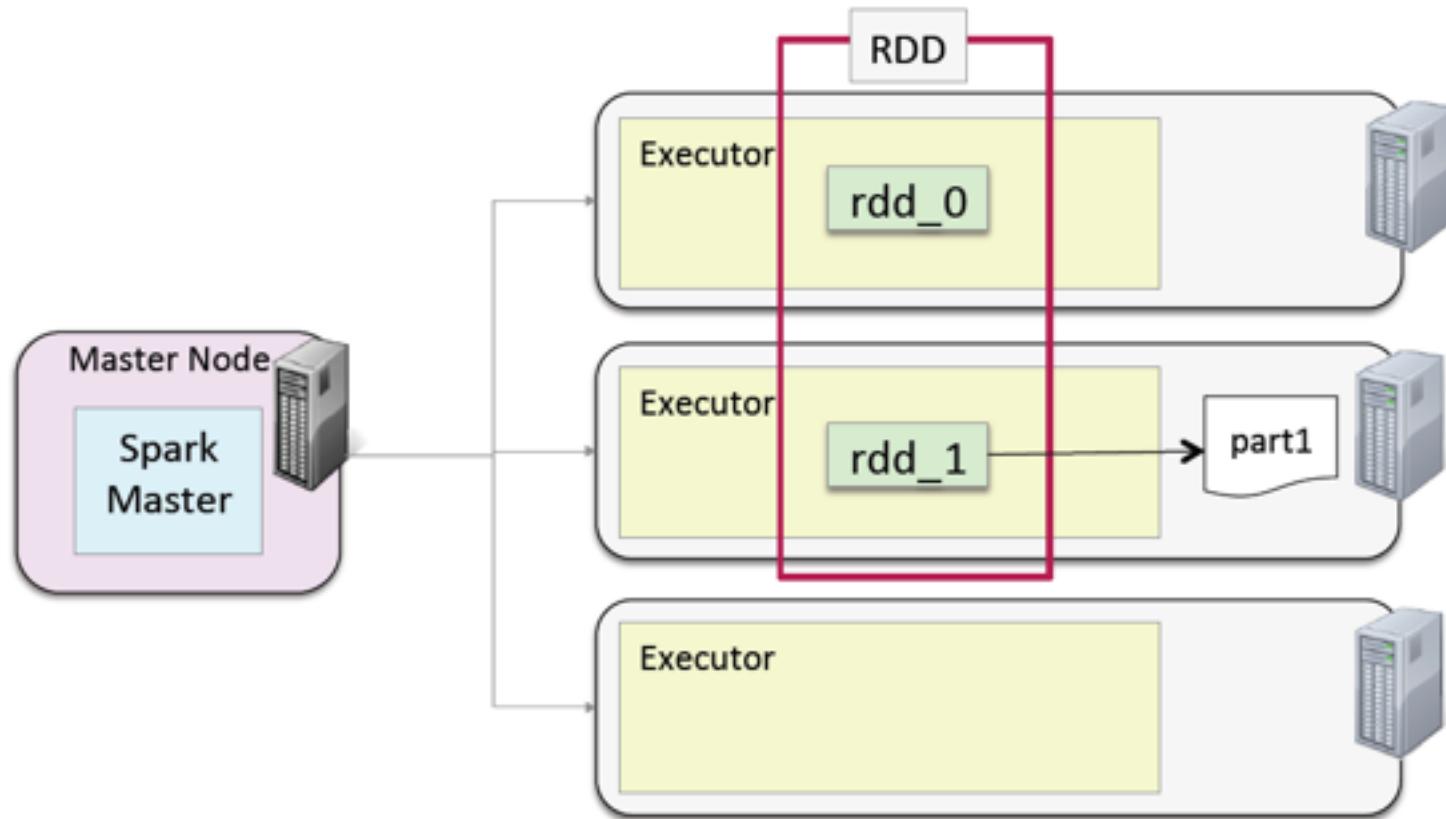
- You can choose to serialize the data in memory
 - `MEMORY_ONLY_SER` and `MEMORY_AND_DISK_SER`
 - Much more space efficient
 - Less time efficient
 - Choose a fast serialization library
- Replication - store partitions on two nodes
 - `MEMORY_ONLY_2`, `MEMORY_AND_DISK_2`, etc.

Changing Persistence Options

- To stop persisting and remove from memory and disk
 - rdd.unpersist()
- To change an RDD to a different persistence level
 - Unpersist first

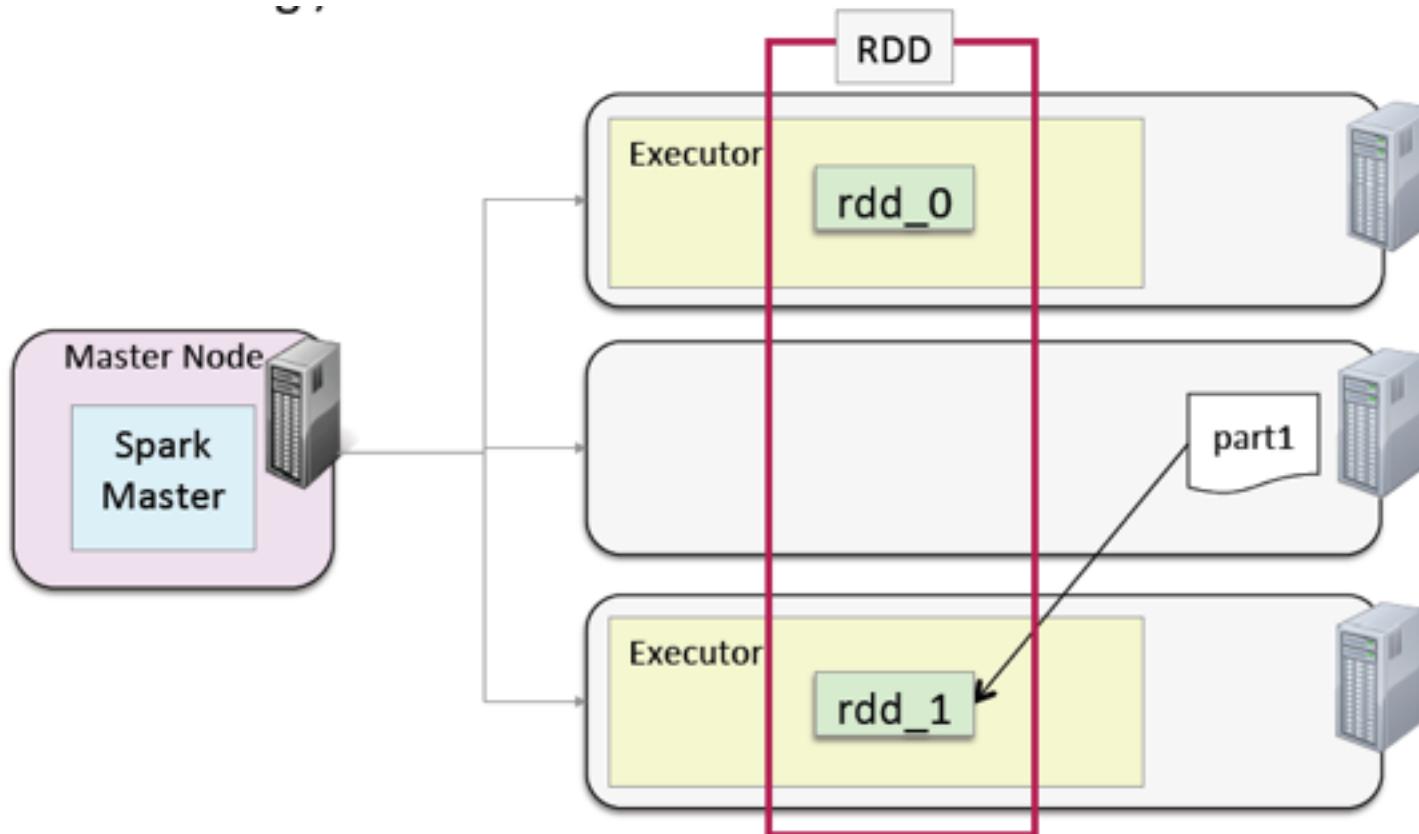
Distributed Disk Persistence

- Disk persisted partitions are stored in local files



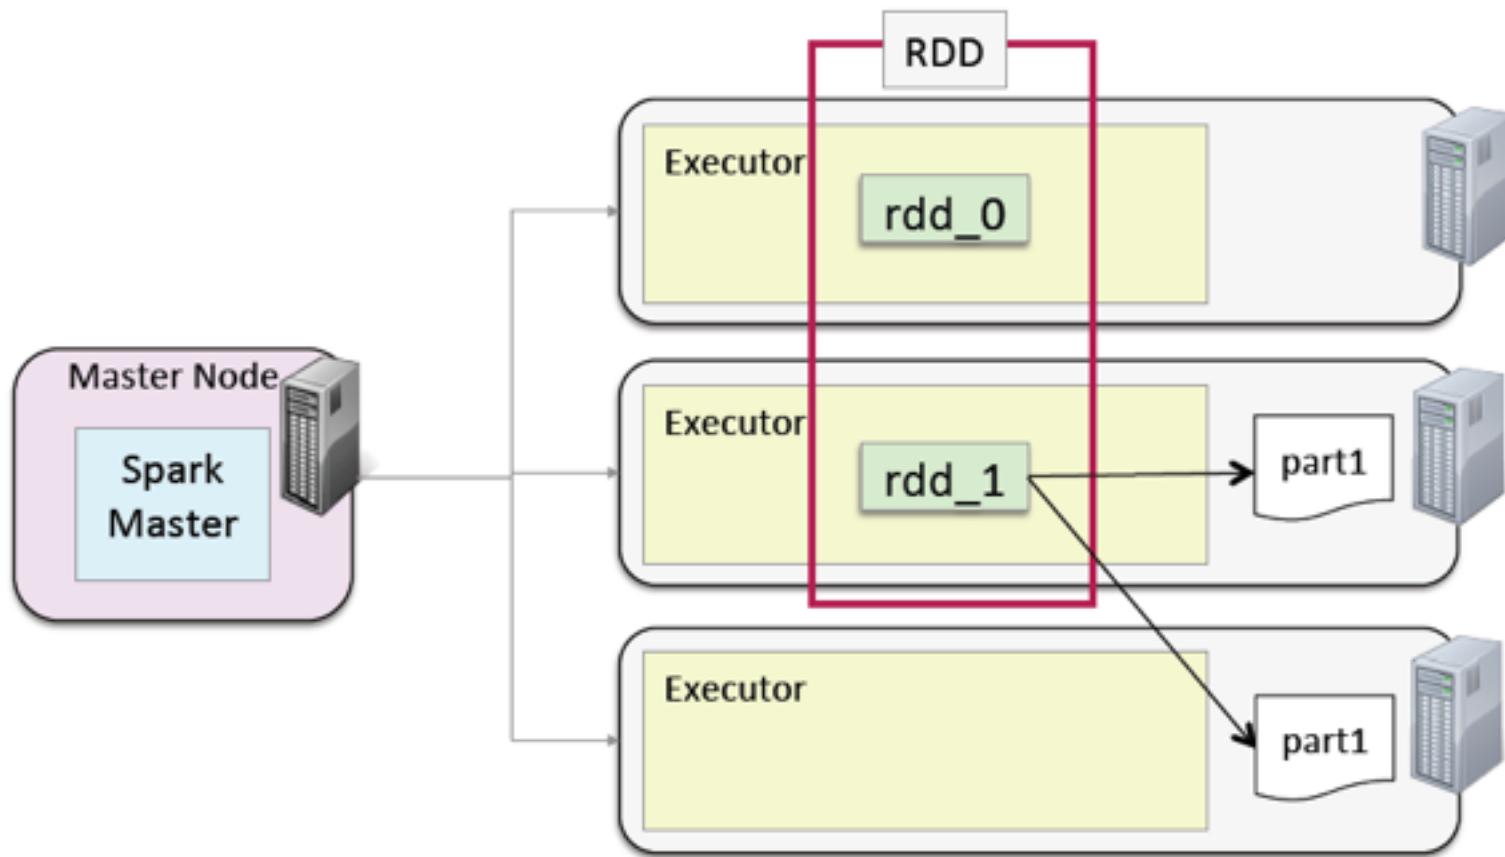
Distributed Disk Persistence

- Data on disk will be used to recreate the partition if possible
 - Will be recomputed if the data is unavailable
 - e.g., the node is down



Replication

- Persistence replication makes recomputation less likely to be necessary



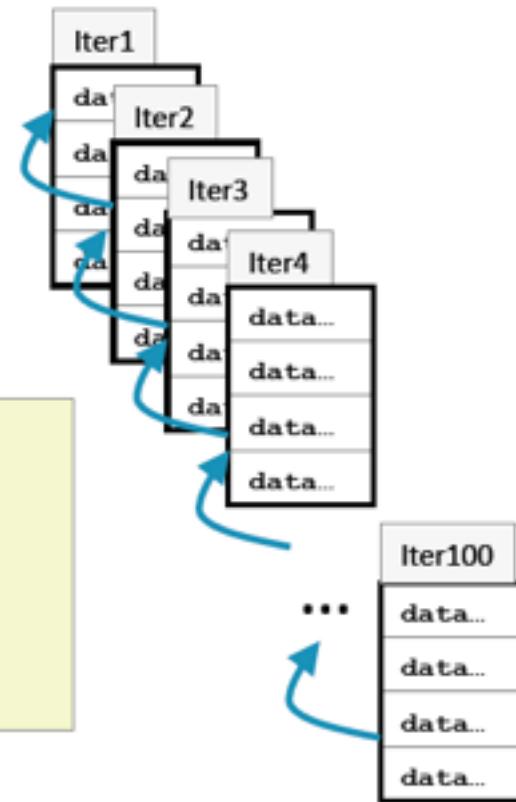
When and Where to Cache

- When should you cache a dataset?
 - When a dataset is likely to be re-used
 - e.g., iterative algorithms, machine learning
- How to choose a persistence level
 - Memory only - when possible, best performance
 - Save space by saving as serialized objects in memory if necessary
 - Disk - choose when recomputation is more expensive than disk read
 - e.g., expensive functions or filtering large datasets
 - Replication - choose when recomputation is more expensive than memory"

Checkpointing

- Maintaining RDD lineage provides resilience but can also cause problems
 - e.g., iterative algorithms, streaming
- Recovery can be very expensive
- Potential stack overflow

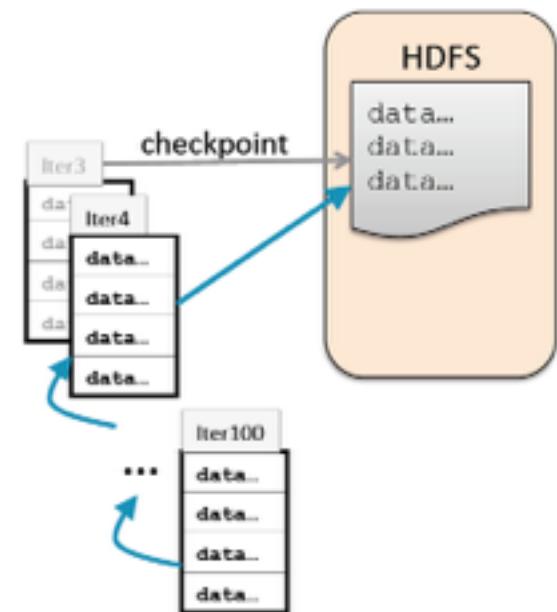
```
myrdd = ...initial-value...
while x in xrange(100):
    myrdd = myrdd.transform(...)
myrdd.saveAsTextFile()
```



Checkpointing

- Checkpointing saves the data to HDFS
- Provides fault-tolerant storage across nodes
- Lineage is not saved
- Must be checkpointed before any actions on the RDD

```
sc.setCheckpointDir(directory)
myrdd = ...initial-value...
while x in xrange(100):
    myrdd = myrdd.transform(...)
    if x % 3 == 0:
        myrdd.checkpoint()
        myrdd.count()
myrdd.saveAsTextFile()
```



Hands-On Exercise:
Caching RDDs
Checkpointing RDDs

HANDS-ON EXERCISES

Chapter 8

WRITING SPARK APPLICATIONS

Spark Shell vs. Spark Applications

- The Spark Shell allows interactive exploration and manipulation of data
 - REPL using Python or Scala
- Spark applications run as independent programs
 - Python, Scala, or Java
 - e.g., ETL processing, Streaming, and so on

The SparkContext

- Every Spark program needs a SparkContext
 - The interactive shell creates one for you
 - You create your own in a Spark application
 - Named sc by convention

Scala Example: WordCount

```
import org.apache.spark.SparkContext
import org.apache.spark.SparkContext._

object WordCount {
    def main(args: Array[String]) {
        if (args.length < 1) {
            System.err.println("Usage: WordCount <file>")
            System.exit(1)
        }

        val sc = new SparkContext()

        val counts = sc.textFile(args(0)).
            flatMap(line => line.split("\\W")).
            map(word => (word,1)).
            reduceByKey(_ + _)

        counts.take(5).foreach(println)
    }
}
```

Running a Spark Application

- The easiest way to run a Spark Application is using the spark-submit script

Python

```
$ spark-submit WordCount.py fileURL
```

Scala/
Java

```
$ spark-submit --class WordCount \  
MyJarFile.jar fileURL
```

Running a Spark Application

- Some key spark-submit options
 - **--help** - explain available options
 - **--master** - equivalent to **MASTER** environment variable for Spark Shell
 - **local[*]** - run locally with as many threads as cores (default)
 - **local[n]** - run locally with n threads
 - **local** - run locally with a single thread
 - **master URL**, e.g., **spark://masternode:7077**
 - **--deploy-mode** - either client or cluster
 - **--name** - application name to display in the UI (default is the Scala/Java class or Python program name)
 - **--jars** - additional JAR files (Scala and Java only)
 - **--pyfiles** - additional Python files (Python only)
 - **--driver-java-options** - parameters to pass to the driver JVM
- 170

Building and Running Scala Applications in the Hands-On Exercises

- Basic Maven projects are provided in the exercises/projects directory with two packages
 - stubs - starter Scala file, do exercises here
 - solution - final exercise solution

```
$ mvn package
```



```
$ spark-submit \
--class stubs.CountJPGs \
target/countjpgs-1.0.jar \
weblogs.*
```

Project Directory Structure

```
+countjpgs
  -pom.xml
  +src
    +main
      +scala
        +solution
          -CountJPGs.scala
        +stubs
          -CountJPGs.scala
    +target
      -countjpgs-1.0.jar
```

Hands-On Exercise: Writing and Running a Spark Application

HANDS-ON EXERCISE: WRITING AND RUNNING A SPARK APPLICATION

Spark Application Configuration

- Spark provides numerous properties for configuring your application
- Some example properties
 - spark.master
 - spark.app.name
 - spark.local.dir - where to store local files such as shuffle output (default/tmp)
 - spark.ui.port - port to run the Spark Application UI (default 4040)
 - spark.executor.memory - how much memory to allocate to each Executor (default 512m)
- Most are more interesting to system administrators than developers
- Spark Applications can be configured
 - At run time or
 - Programmatically

Run-time Configuration Options

- **spark-submit script**
 - e.g., `spark-submit --master spark://masternode:7077`
- **Properties file**
 - Tab - or space-separated list of properties and values
 - Load with `spark-submit --properties-file filename`
 - Example:

```
spark.master      spark://masternode:7077
spark.local.dir  /tmp
spark.ui.port    4444
```

- **Site defaults properties file**
 - `$SPARK_HOME/conf/spark-defaults.conf`
 - Template file provided

Setting Configuration Properties Programmatically

- Spark configuration settings are part of the `SparkContext`
- Configure using a `SparkConf` object
- Some example functions
 - `setAppName(name)`
 - `setMaster(master)`
 - `set(property-name, value)`
- Set functions return a `SparkConf` object to support chaining

SparkConf Example (Scala)

```
import org.apache.spark.SparkContext
import org.apache.spark.SparkContext._
import org.apache.spark.SparkConf

object WordCount {
    def main(args: Array[String]) {
        if (args.length < 1) {
            System.err.println("Usage: WordCount <file>")
            System.exit(1)
        }

        val sconf = new SparkConf()
        .setAppName("Word Count")
        .set("spark.ui.port","4141")
        val sc = new SparkContext(sconf)

        val counts = sc.textFile(args(0)).
            flatMap(line => line.split("\\W")).
            map(word => (word,1)).
            reduceByKey(_ + _)
        counts.take(5).foreach(println)
    }
}
```

Viewing Spark Properties

- You can view the Spark property setting in the Spark Application UI

The screenshot shows the PySparkShell application UI with the "Environment" tab highlighted by a red box. The main content area displays two tables: "Runtime Information" and "Spark Properties".

Runtime Information

Name	Value
Java Home	/usr/java/jdk1.7.0_51/jre
Java Version	1.7.0_51 (Oracle Corporation)
Scala Home	
Scala Version	version 2.10.3

Spark Properties

Name	Value
spark.app.name	PySparkShell
spark.driver.host	master
spark.driver.port	33121
spark.fileserver.uri	http://master:34670
spark.broadcast.uri	http://master:38591
spark.master	spark://master:7077

Spark Logging

- Spark uses Apache Log4j for logging
 - Allows for controlling logging at runtime using a properties file
 - Enable or disable logging, set logging levels, select output destination
 - For more info see <http://logging.apache.org/log4j/1.2/>
- Log4j provides several logging levels
 - Fatal
 - Error
 - Warn
 - Info
 - Debug
 - Trace
 - Off

Spark Log Files

- Log file locations depend on your cluster management platform
- Spark Standalone defaults:
 - Spark daemons: /usr/hdp/current/spark/logs
 - Individual tasks: \$SPARK_HOME/work on each worker node

Spark Worker UI - Log File Access

- Log file locations depend on your cluster management platform
- Spark Standalone defaults:
 - Spark daemons: /var/log/spark
 - Individual tasks: \$SPARK_HOME/work on each worker node

ID: worker-20140121065745-ip-10-236-129-42.ec2.internal-60105
Master URL: spark://ec2-23-20-24-104.compute-1.amazonaws.com:7077
Cores: 4 (4 Used)
Memory: 13.6 GB (12.6 GB Used)

[Back to Master](#)

Running Executors 1

ExecutorID	Cores	Memory	Job Details	Logs
4	4	12.6 GB	ID: app-20140121220135-0003 Name: PageRank User: root	stdout stderr

Configuring "Spark" Logging "(1)"

- Logging levels can be set for the cluster, for individual applications, or even for specific components or subsystems
- Default for machine: \$SPARK_HOME/conf/log4j.properties
 - Start by copying log4j.properties.template

log4j.properties.template

```
# Set everything to be logged to the console
log4j.rootCategory=INFO, console
log4j.appender.console=org.apache.log4j.ConsoleAppender
log4j.appender.console.target=System.err
...
...
```

Configuring Spark Logging (2)

- Spark will use the first log4j.properties file it finds in the Java classpath
- Spark Shell will read log4j.properties from the current directory
 - Copy log4j.properties to the working directory and edit

...my-working-directory/log4j.properties

```
# Set everything to be logged to the console
log4j.rootCategory=DEBUG, console
log4j.appender.console=org.apache.log4j.ConsoleAppender
log4j.appender.console.target=System.err
...
...
```

Hands-On Exercise: Configuring Spark Applications

HANDS-ON EXERCISE: CONFIGURING SPARK APPLICATIONS

Chapter 09

SPARK STREAMING

What is Spark Streaming?

- Spark Streaming provides real time processing of stream data
- An extension of core Spark
- Currently supports Scala and Java

Why Spark Streaming?

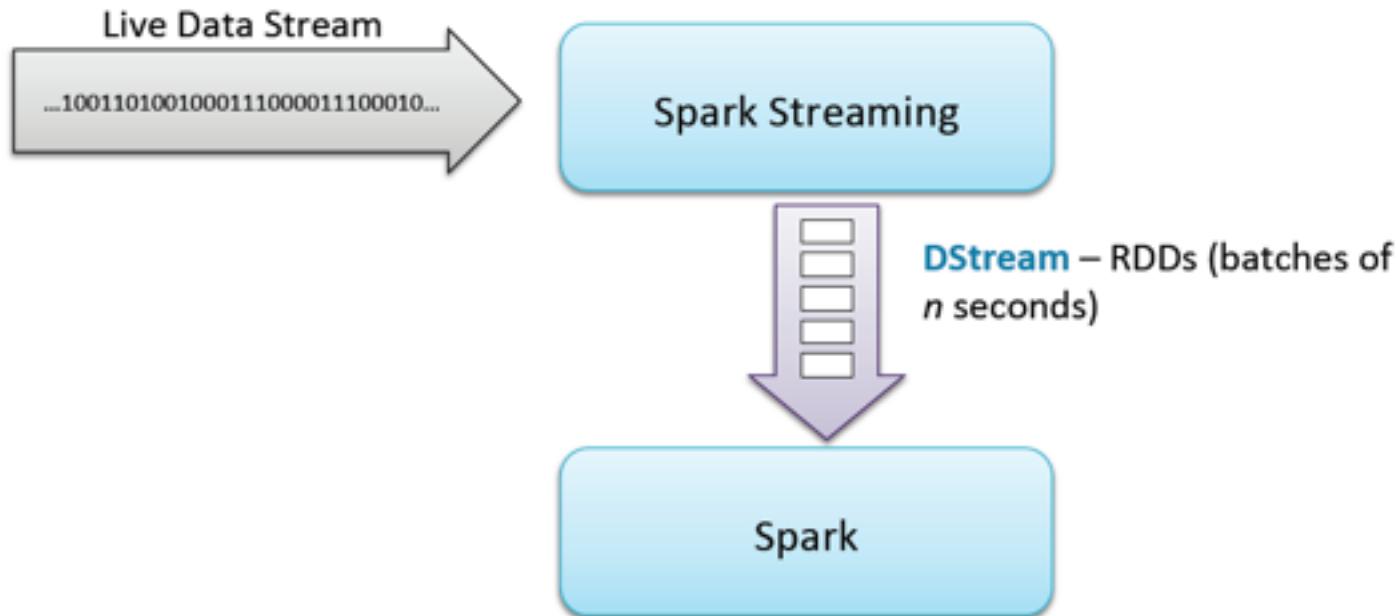
- Many big data applications need to process large data streams in real time
 - Website monitoring
 - Fraud detection
 - Ad monetization

Spark Streaming Features

- Second scale latencies
- Scalability and efficient fault tolerance
- “Once and only once” processing
- Integrates batch and real-time processing
- Easy to develop
 - Uses Spark’s high level API

Spark Streaming Overview

- Divide up data stream into batches of n seconds
- Process each batch in Spark as an RDD
- Return results of RDD operations in batches



Streaming Example: Streaming Request Count

```
object StreamingRequestCount {

    def main(args: Array[String]) {

        val ssc = new StreamingContext(new SparkConf(), Seconds(2))
        val mystream = ssc.socketTextStream(hostname, port)
        val userreqs = mystream
            .map(line => (line.split(" ") (2), 1))
            .reduceByKey((x, y) => x+y)

        userreqs.print()

        ssc.start()
        ssc.awaitTermination()
    }
}
```

Streaming Example: Configuring StreamingContext

```
object StreamingRequestCount {  
  
    def main(args: Array[String]) {  
  
        val ssc = new StreamingContext(new SparkConf(), Seconds(2))  
        val  
        val  
        use  
        ssc  
        ssc  
    }  
}
```

- A **StreamingContext** is the main entry point for Spark Streaming apps
- Equivalent to **SparkContext** in core Spark
- Configured with the same parameters as a **SparkContext** plus *batch duration* – instance of **Milliseconds**, **Seconds**, or **Minutes**
- Named **ssc** by convention

Streaming Example: Creating a DStream

```
object StreamingRequestCount {  
  
  def main(args: Array[String]) {  
  
    val ssc = new StreamingContext(new SparkConf(), Seconds(2))  
    val logs = ssc.socketTextStream(hostname, port)  
    val ...  
  


- Get a DStream (“Discretized Stream”) from a streaming data source, e.g., text from a socket

  
    use ...  
  
    ssc.start()  
    ssc.awaitTermination()  
  }  
}
```

Streaming Example: Dstream Transformations

```
object StreamingRequestCount {  
  
  def main(args: Array[String]) {  
  
    val ssc = new StreamingContext(new SparkConf(), Seconds(2))  
    val logs = ssc.socketTextStream(hostname, port)  
    val userreqs = logs  
      .map(line => (line.split(" ") (2), 1))  
      .reduceByKey((x,y) => x+y)  
  }  
}
```

- DStream operations are applied to each batch RDD in the stream
- Similar to RDD operations – **filter**, **map**, **reduce**, **joinByKey**, etc.

Streaming Example: Dstream Result Output

```
object StreamingRequestCount {  
  
  def main(args: Array[String]) {  
  
    val ssc = new StreamingContext(new SparkConf(), Seconds(2))  
    val logs = ssc.socketTextStream(hostname, port)  
    val userreqs = logs  
      .map(line => (line.split(" ") (2), 1))  
      .reduceByKey((x, y) => x+y)  
  
    userreqs.print()  
  
    s:  
    s:  
  }  
}  
}
```

- Print out the first 10 elements of each RDD

Streaming Example: Starting the Streams

```
object StreamingRequestCount {  
  
    def main(args: Array[String]) {  
  
        val ssc = new StreamingContext(new SparkConf(), Seconds(2))  
        val logs = ssc.socketTextStream(hostname, port)  
        val wordCounts = logs.flatMap(_.split(" ")).map((_, 1))  
        wordCounts.print()  
  
        ssc.start()  
        ssc.awaitTermination()  
    }  
}
```

- **start**: Starts the execution of all DStreams
- **awaitTermination**: waits for all background threads to complete before ending the main thread

Streaming Example: Streaming Request Count (Recap)

```
object StreamingRequestCount {

    def main(args: Array[String]) {

        val ssc = new StreamingContext(new SparkConf(), Seconds(2))
        val logs= ssc.socketTextStream(hostname, port)
        val userreqs = logs
            .map(line => (line.split(" ") (2),1))
            .reduceByKey((x,y) => x+y)

        userreqs.print()

        ssc.start()
        ssc.awaitTermination()
    }
}
```

Streaming Example Output

```
-----  
Time: 1401219545000 ms  
-----
```

```
(23713,2)  
(53,2)  
(24444,2)  
(127,2)  
(93,2)  
...
```

Starts 2 seconds
after `ssc.start`

Streaming Example Output

```
-----  
Time: 1401219545000 ms  
-----
```

```
(23713,2)  
(53,2)  
(24444,2)  
(127,2)  
(93,2)
```

```
...
```

```
-----  
Time: 1401219547000 ms  
-----
```

```
(42400,2)  
(24996,2)  
(97464,2)  
(161,2)  
(6011,2)
```

```
...
```

2 seconds later...

Streaming Example Output

```
-----  
Time: 1401219545000 ms  
-----
```

```
(23713,2)  
(53,2)  
(24444,2)  
(127,2)  
(93,2)
```

```
...
```

```
-----  
Time: 1401219547000 ms  
-----
```

```
(42400,2)  
(24996,2)  
(97464,2)  
(161,2)  
(6011,2)
```

```
...
```

```
-----  
Time: 1401219549000 ms  
-----
```

```
(44390,2)  
(48712,2)  
(165,2)  
(465,2)  
(120,2)
```

```
...
```



2 seconds later...

Streaming Example Output

```
-----  
Time: 1401219545000 ms  
-----
```

```
(23713,2)  
(53,2)  
(24444,2)  
(127,2)  
(93,2)
```

```
...
```

```
-----  
Time: 1401219547000 ms  
-----
```

```
(42400,2)  
(24996,2)  
(97464,2)  
(161,2)  
(6011,2)
```

```
...
```

```
-----  
Time: 1401219549000 ms  
-----
```

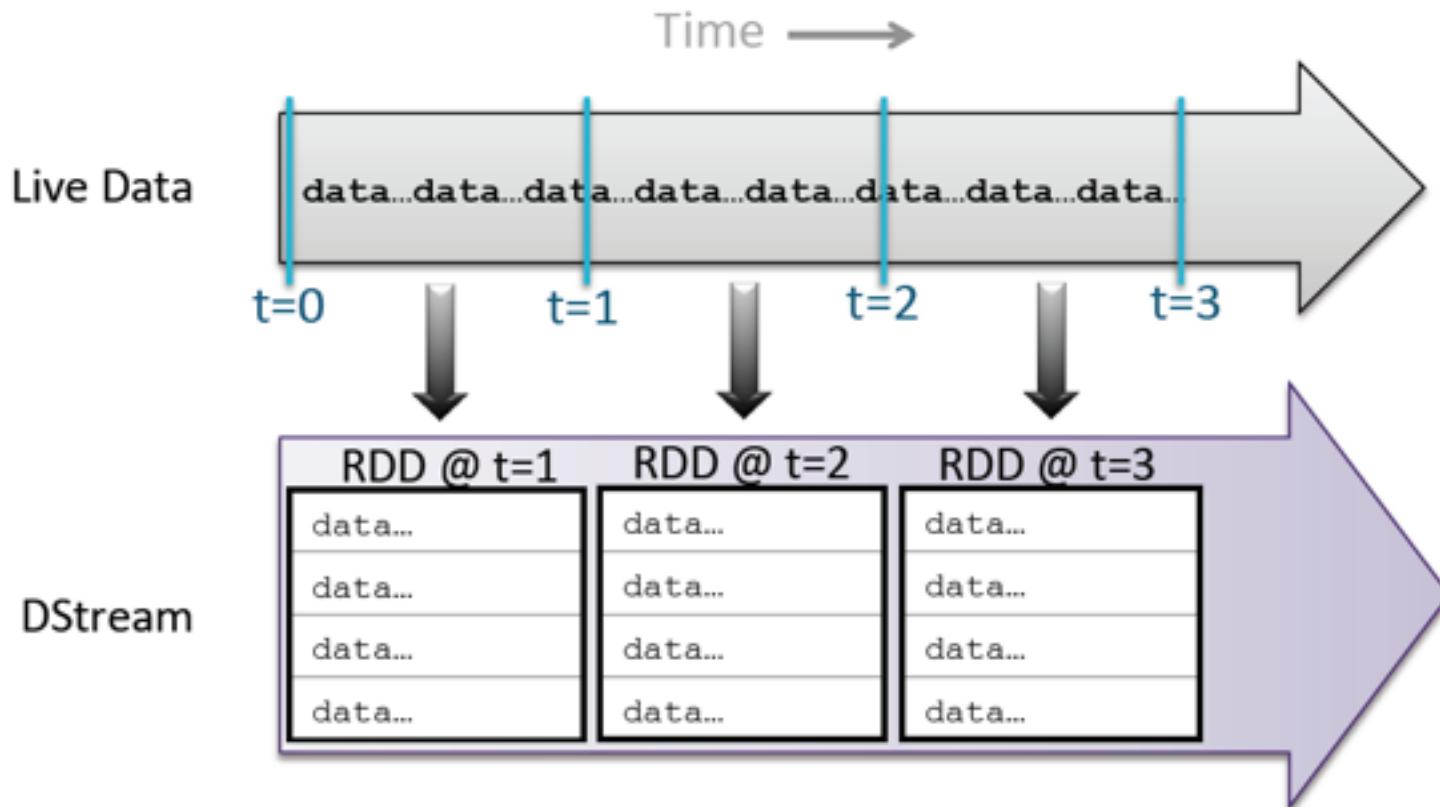
```
(44390,2)  
(48712,2)  
(165,2)  
(465,2)  
(120,2)
```

```
...
```

Continues until
termination...

DStreams

- A Dstream is a sequence of RDDs representing a data stream
 - “Discretized Stream”



Dstream Data Sources

- Dstreams are defined for a given input stream (e.g., a Unix socket)
 - Created by the StreamingContext
`ssc.socketTextStream(hostname, port)`
 - Similar to how RDDs are created by the SparkContext
- Out-of-the- box data sources
 - Network
 - Sockets
 - Other network sources, e.g., Flume, Akka Actors, Kaga, Twitter
 - Files
 - Monitors an HDFS directory for new content

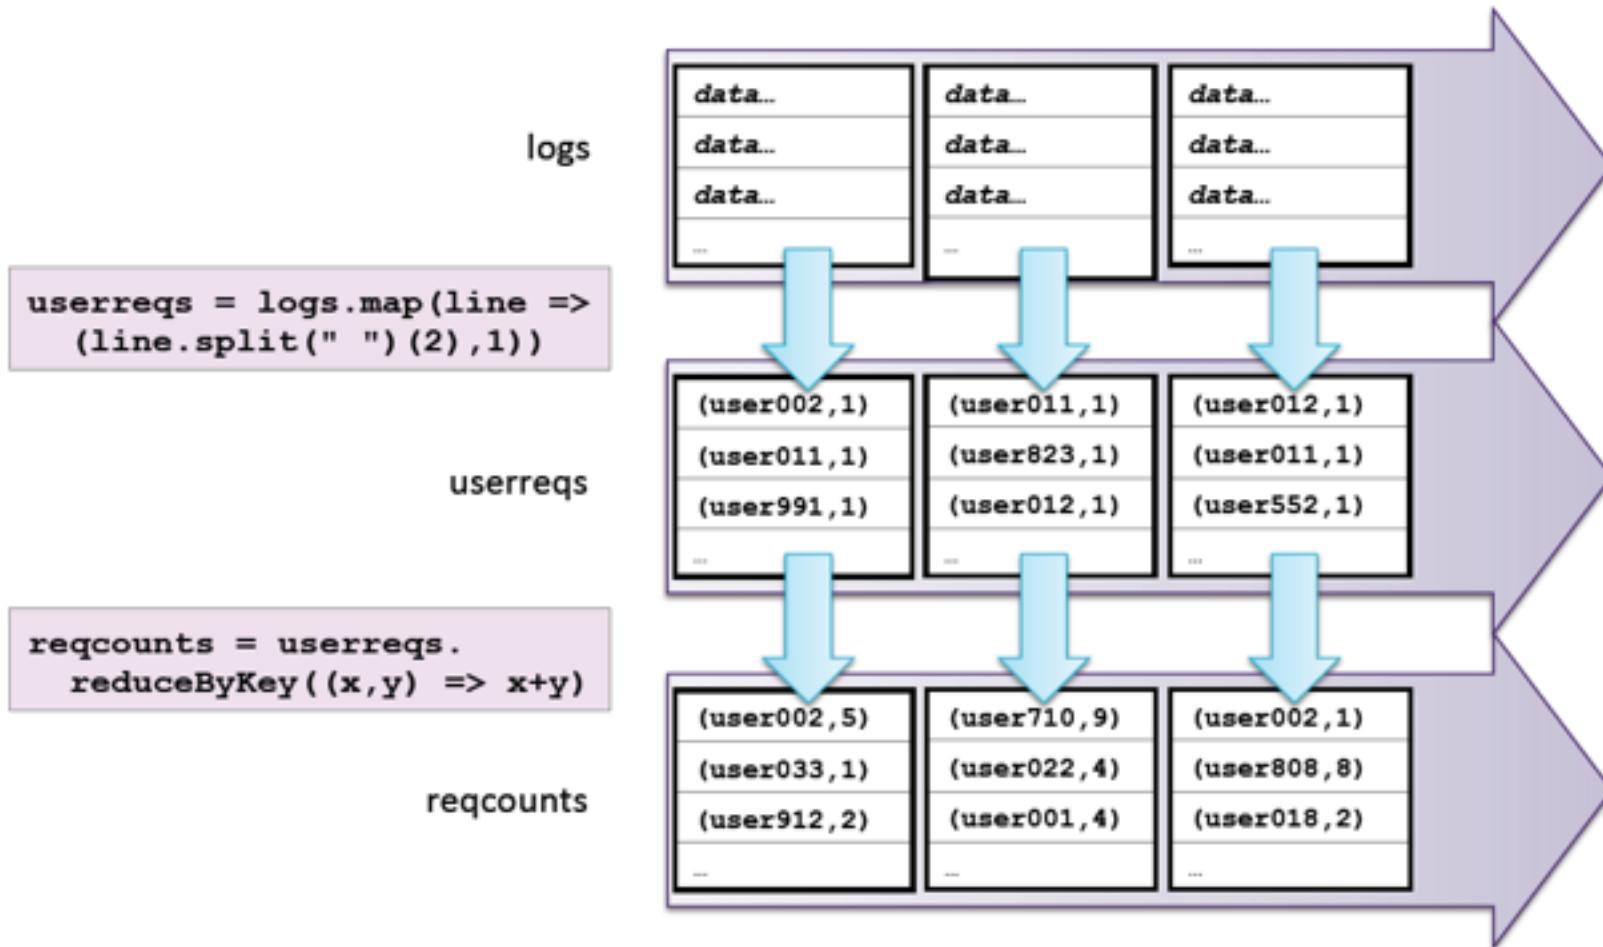
Dstream Operations

- Dstream operations are applied to every RDD in the stream
 - Executed once per duration
- Two types of Dstream operations
 - Transformations
 - Create a new Dstream from an existing one
 - Output operations
 - Write data (for example, to a file system, database, or console)
 - Similar to RDD actions

Dstream Transformations

- Many RDD transformations are also available on DStreams
 - Regular transformations such as map, flatMap, filter
 - Pair transformations such as reduceByKey, groupByKey, joinByKey
- What if you want to do something else?
 - transform(function)
 - Creates a new Dstream by executing function on RDDs in the current DStream

Dstream Transformations

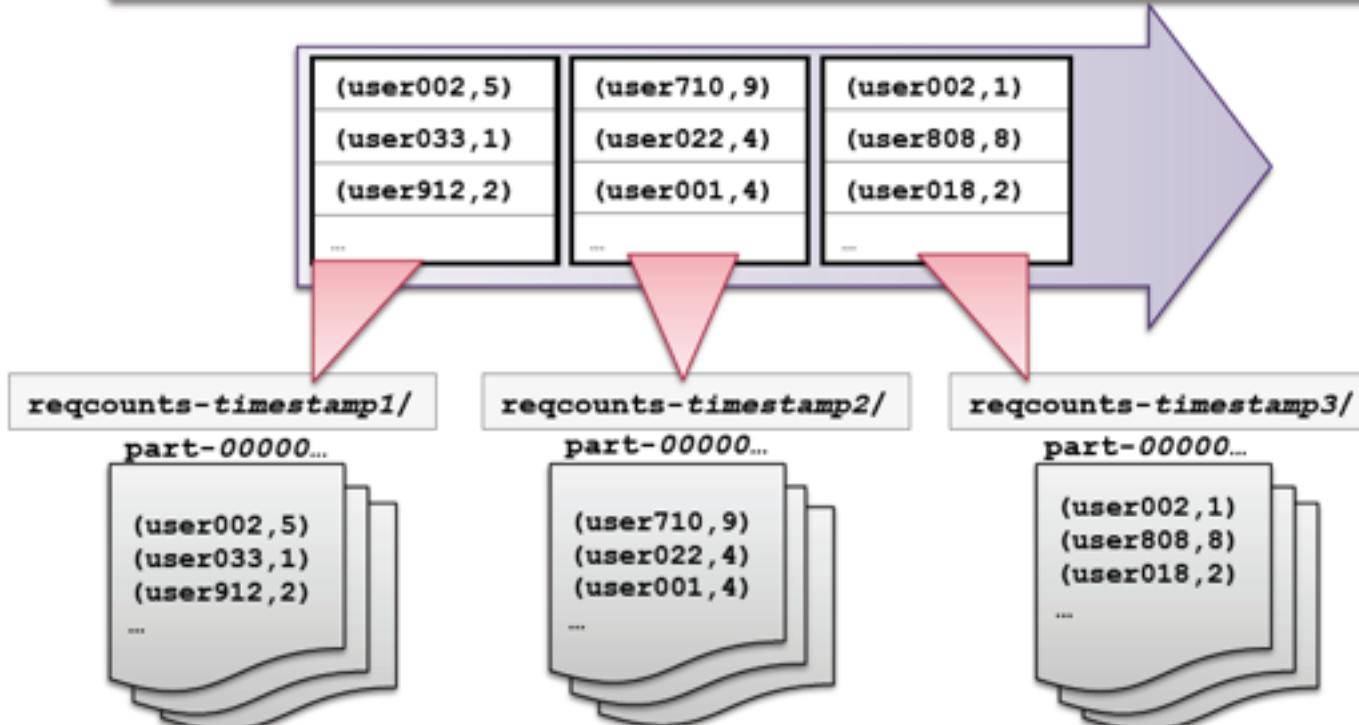


Dstream Output Operations

- **Console output**
 - `print` - prints out the first 10 elements of each RDD
- **File output**
 - `saveAsTextFiles` - save data as text
 - `saveAsObjectFiles` - save as serialized object files
- **Other types of functions**
 - `foreachRDD (function, time)` - performs a function on each RDD in the Dstream

Saving Dstream Results as Files

```
val userreqs = logs
  .map(line => (line.split(" ") (2), 1))
  .reduceByKey((x,y) => x+y)
userreqs.print()
userreqs.saveAsTextFiles (".../outdir/reqcounts")
```



Find Top Users

```
...  
val userreqs = logs  
  .map(line => (line.split(" ") (2), 1))  
  .reduceByKey((x,y) => x+y)  
userreqs.saveAsTextFiles(path)  
  
val sortedreqs = userreqs  
  .map(pair => pair.swap)  
  .transform(rdd => rdd.sortByKey(false))  
  
sortedreqs.foreach(  
  println("Top users: ")  
  rdd.take(5).foreach(  
    pair => printf("User: %s (%s)\n", pair._2, pair._1))  
}  
}  
  
ssc.start()  
ssc.awaitTermination()  
...
```

Transform each RDD: swap userID/count, sort by count

Find Top Users

```
...
val userreqs = logs
  .map(line => (line.split(" ") (2) ,1))
  .reduceByKey((x,y) => x+y)
userreqs.saveAsTextFiles(path)

val sortedreqs = userreqs
  .map(pair => (pair._2, pair._1))
  .transform(Print out the top 5 users as userID (count))
  .sortByKey(false)

sortedreqs.foreachRDD((rdd,time) => {
  println("Top users @ " + time)
  rdd.take(5).foreach(
    pair => printf("User: %s (%s)\n",pair._2, pair._1))
})
ssc.start()
ssc.awaitTermination()
...
```

Find Top Users - Output

-

```
Top users @ 1401219545000 ms
User: 16261 (8)
User: 22232 (7)
User: 66652 (4)
User: 21205 (2)
User: 24358 (2)
```

t = 0 (2 seconds
after `ssc.start()`)

Find Top Users - Output

-

```
Top users @ 1401219545000 ms
User: 16261 (8)
User: 22232 (7)
User: 66652 (4)
User: 21205 (2)
User: 24358 (2)
Top users @ 1401219547000 ms
User: 53667 (4)
User: 35600 (4)
User: 62 (2)
User: 165 (2)
User: 40 (2)
```

$t=1$
(2 seconds later)

Find Top Users - Output

- ```
Top users @ 1401219545000 ms
User: 16261 (8)
User: 22232 (7)
User: 66652 (4)
User: 21205 (2)
User: 24358 (2)
Top users @ 1401219547000 ms
User: 53667 (4)
User: 35600 (4)
User: 62 (2)
User: 165 (2)
User: 40 (2)
Top users @ 1401219549000 ms
User: 31 (12)
User: 6734 (10)
User: 14986 (10)
User: 72760 (2)
User: 65335 (2)
Top users @ 1401219551000 ms
...
```

$t = 2$   
(2 seconds later)

Continues until  
termination...

# Using Spark Streaming with Spark Shell

- Spark Streaming is designed for batch applications, not interactive use
- Spark Shell can be used for limited testing
  - Adding operations after the Streaming Context has been started is unsupported

```
$ spark-shell --master local[2]
```

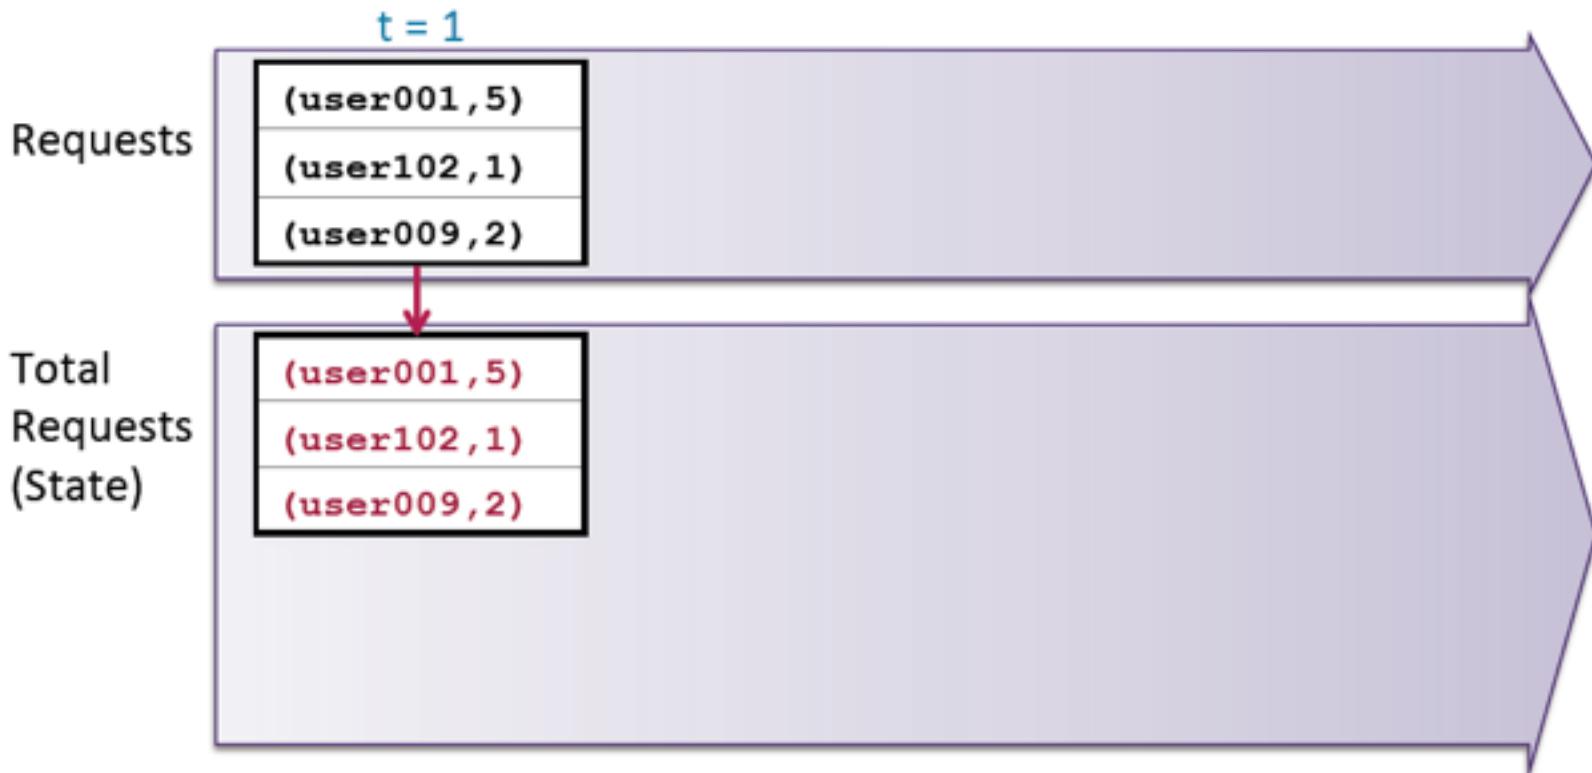
- Stopping and restarting the Streaming Context is unsupported

Hands-On Exercise: Exploring Spark Streaming

# HANDS-ON EXERCISE: EXPLORING SPARK STREAMING

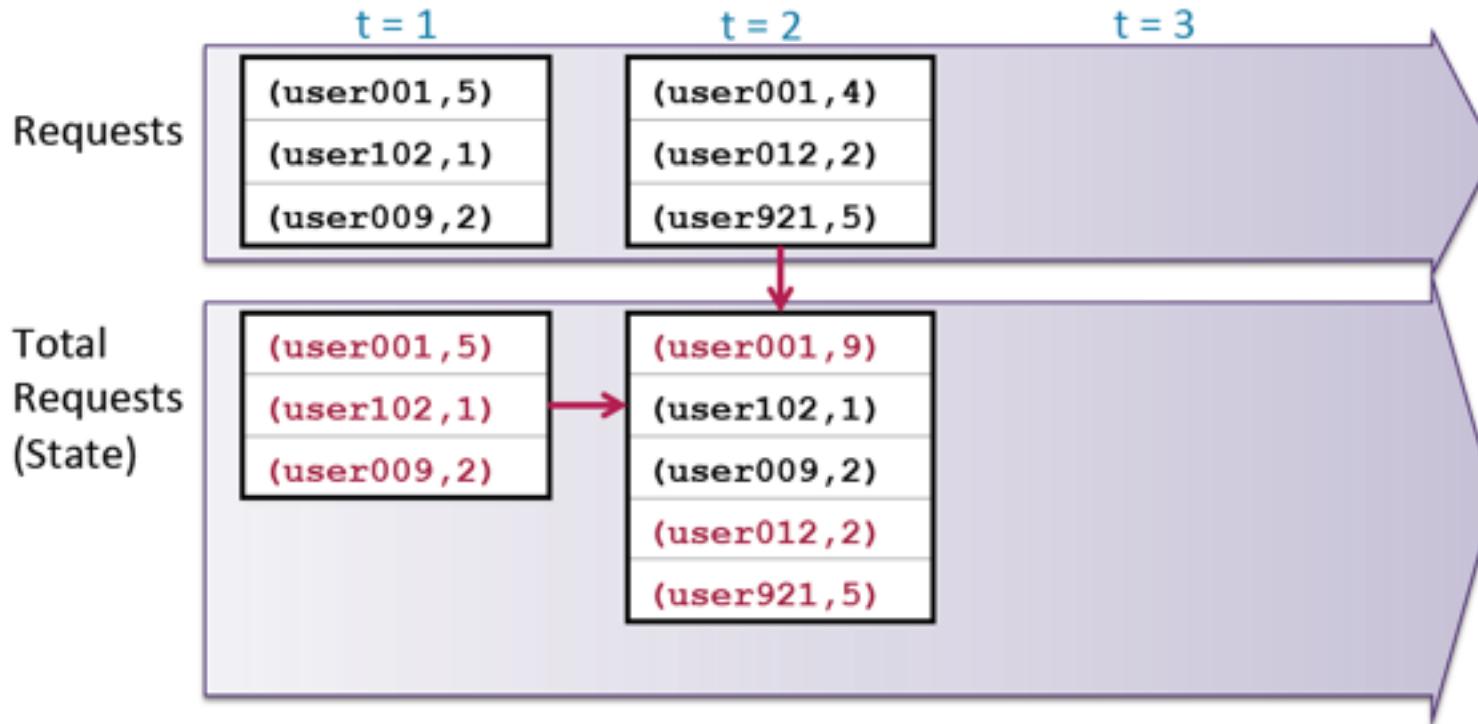
# State Dstreams

- Use the `updateStateByKey` function to create a state DStream
- Example: Total request count by User ID



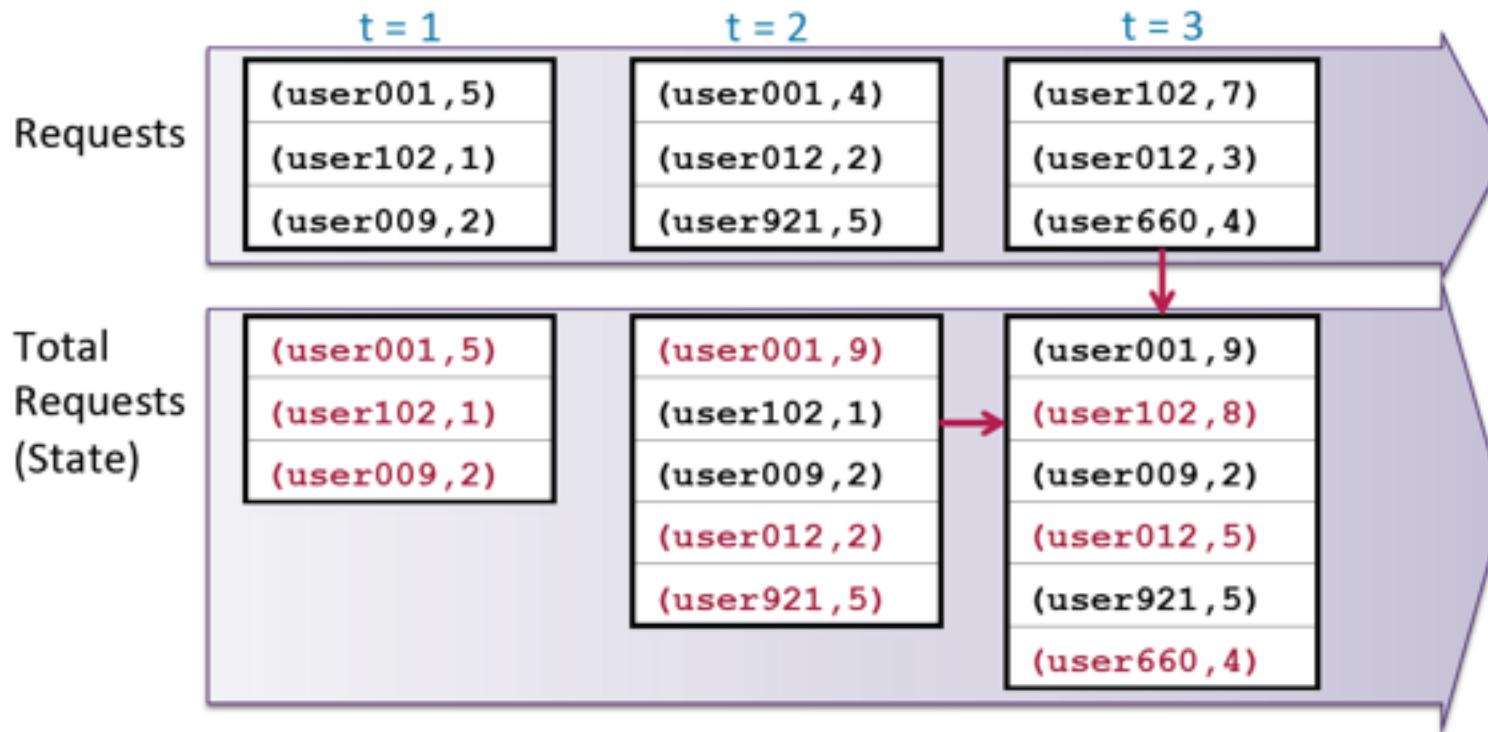
# State Dstreams

- Use the `updateStateByKey` function to create a state DStream
- Example: Total request count by User ID



# State Dstreams

- Use the `updateStateByKey` function to create a state DStream
- Example: Total request count by User ID



# Total User Request Count

```
...
val userreqs = logs
 .map(line => (line.split(" ") (2) ,1))
 .reduceByKey((x,y) => x+y)
...

ssc.checkpoint("checkpoints")
val totalUserreqs = userreqs.updateStateByKey(updateCount)
totalUserreqs.print()
```

```
ssc.start()
ssc.awaitTermination()
...
```

Set checkpoint directory to enable checkpointing.  
Required to prevent infinite lineages.

# Total User Request Count

```
...
val userreqs = logs
 .map(line => (line.split(" ") (2), 1))
 .reduceByKey((x,y) => x+y)
...
ssc.checkpoint("checkpoints")
val totalUserreqs = userreqs.updateStateByKey(updateCount)
totalUserreqs.print()
```

next slide...

ssc.start()  
ssc.awaitTermination()

Compute a state DStream based on the previous states updated with the values from the current batch of request counts

# Total User Request Count - Update Function

```
def updateCount = (newCounts: Seq[Int], state: Option[Int]) => {
 val newCount = newCounts.foldLeft(0) (_ + _)
 val previousCount = state.getOrElse(0)
 Some(newCount + previousCount)
}
```

New Values

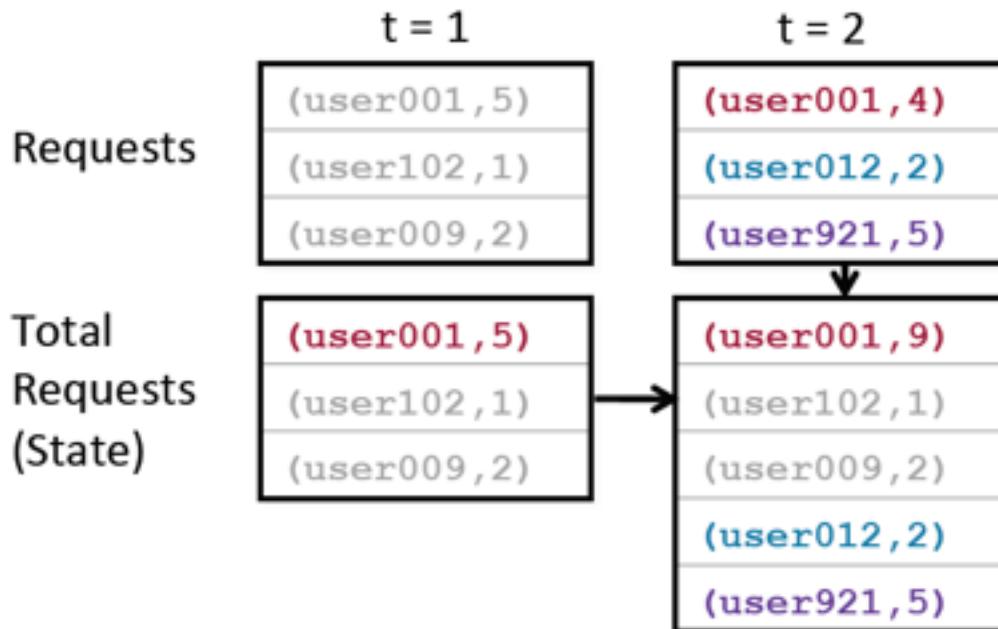
Current State (or **None**)

New State

Given an existing state for a key (user), and new values (counts), return a new state (sum of current state and new counts)

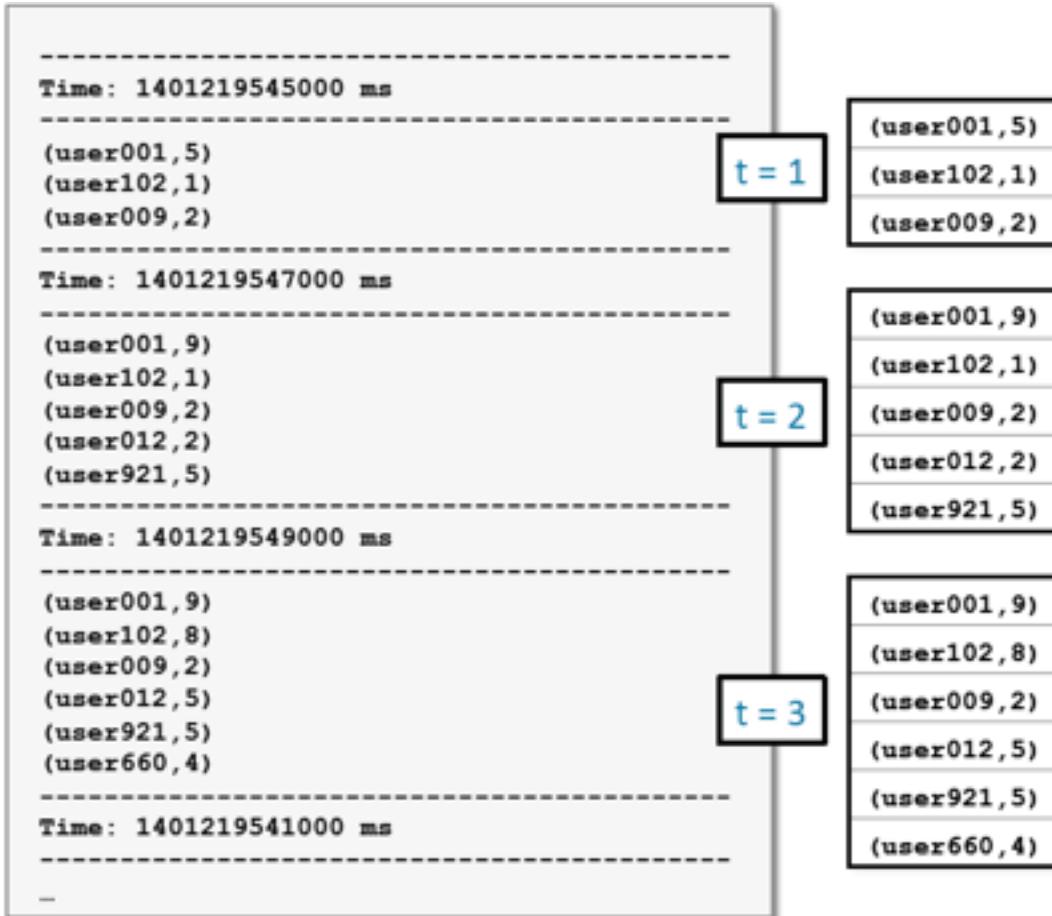
# Total User Request Count - Update Function

- Example at t=2
  - user001: updateCount([4],Some[5])  9
  - user012: updateCount([2],None))  2
  - user921: updateCount([5],None))  5



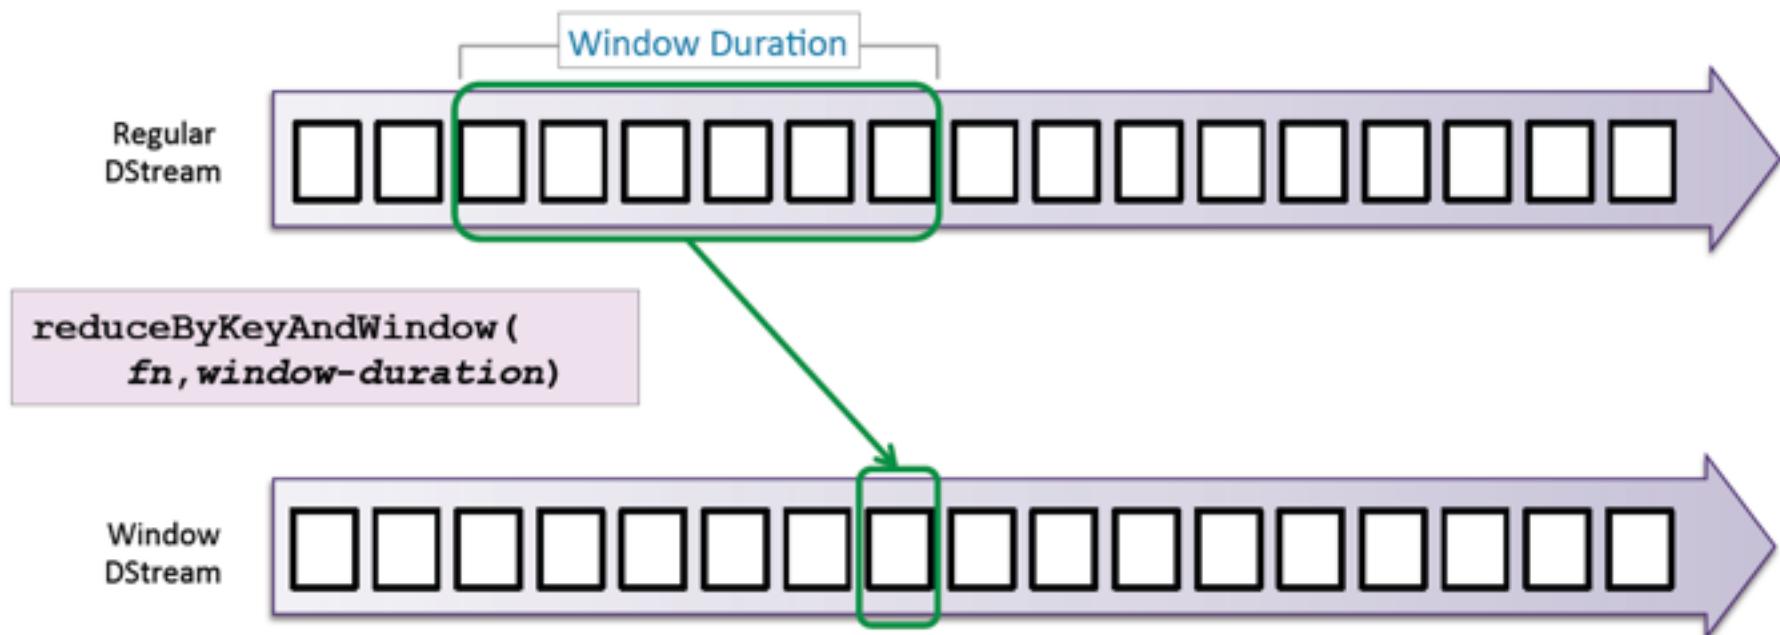
# Maintaining State - Output

- 



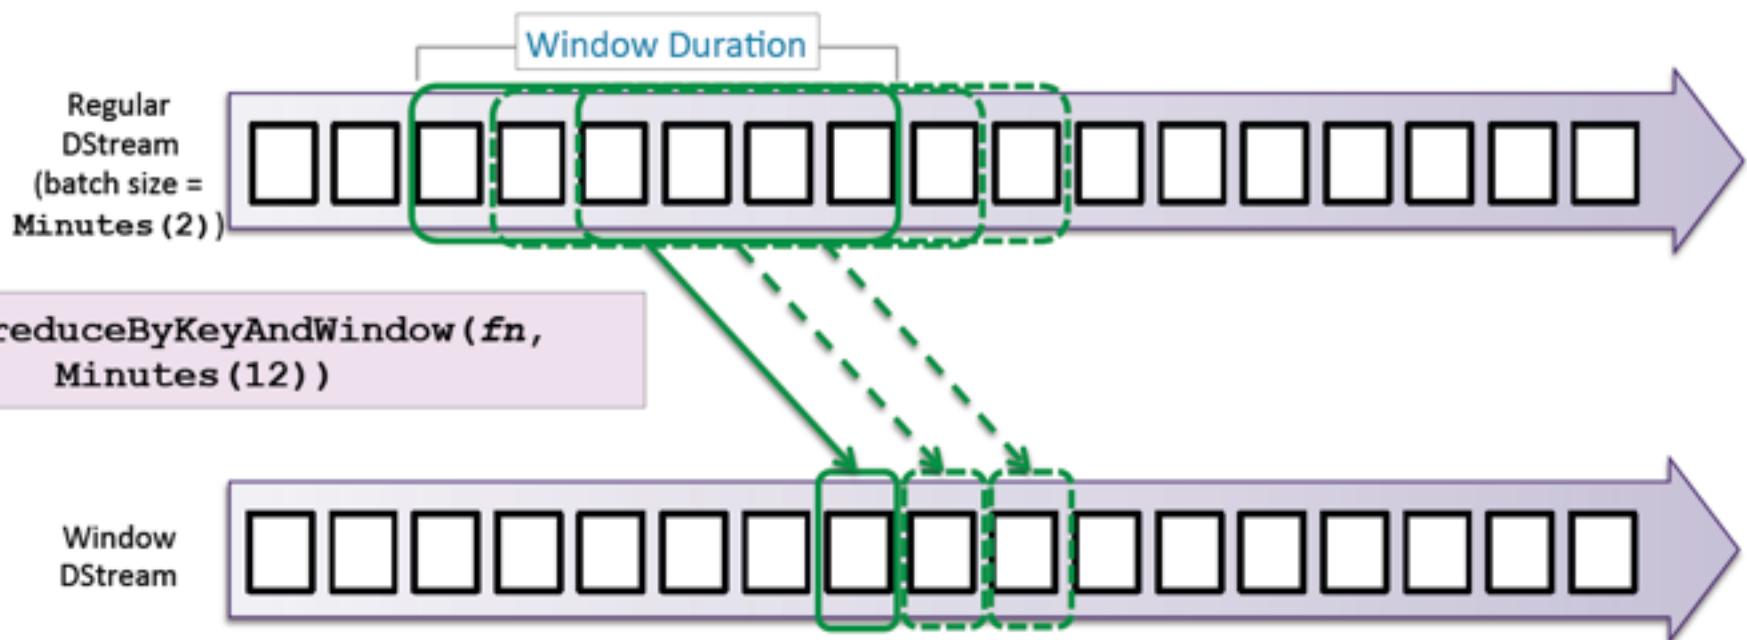
# Sliding Window Operations

- Regular Dstream operations execute for each RDD based on SSC duration
- “Window” operations span RDDs over a given duration
  - e.g., `reduceByKeyAndWindow`, `countByWindow`



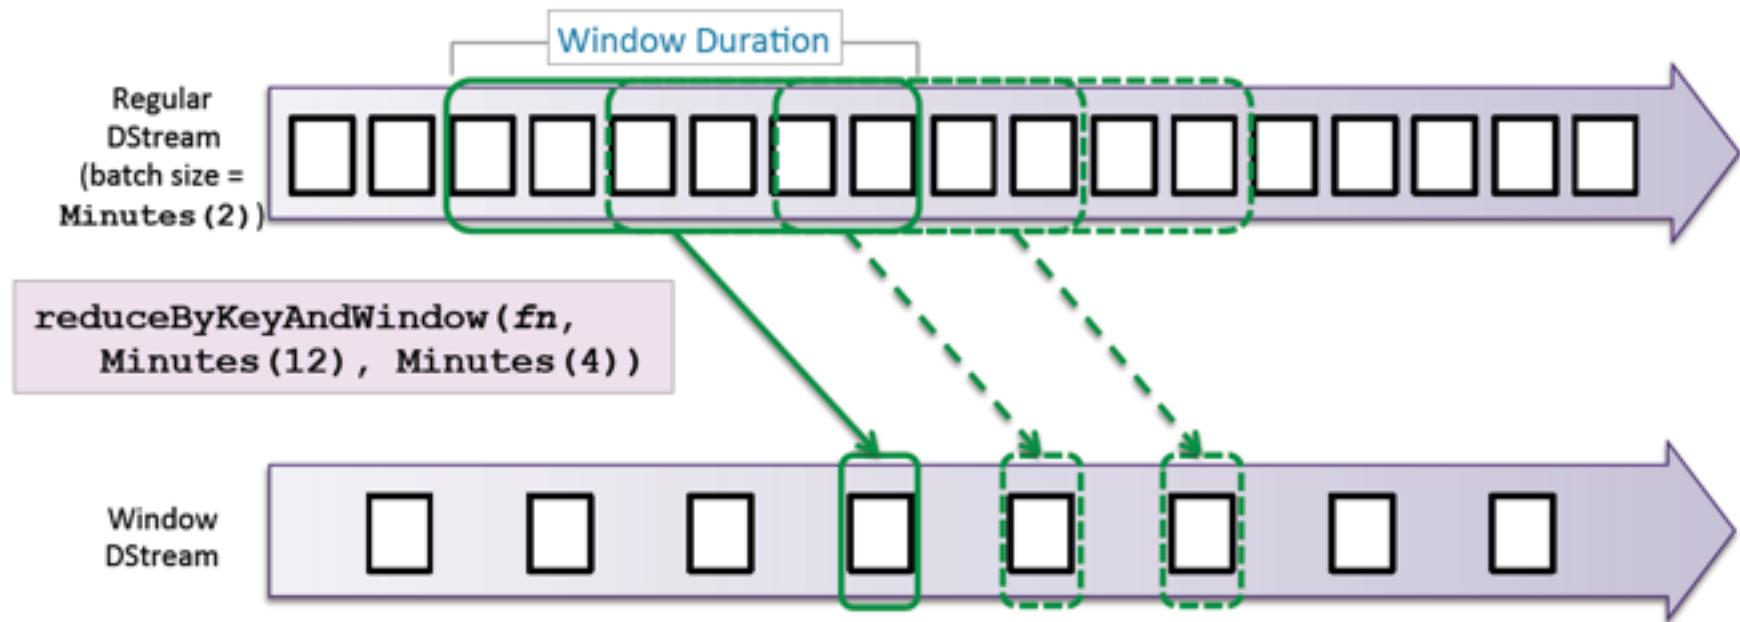
# Sliding Window Operations

- By default window operations will execute with an “interval” the same as the SCC duration
  - i.e., for 2 minute batch duration, window will “slide” every 2 minutes



# Sliding Window Operations

- You can specify a different slide duration (must be a multiple of the SSC duration)



# Count and Sort User Requests by Window

```
...
val ssc = new StreamingContext(new StreamConf(), Seconds(2))
val logs = ssc.socketTextStream(hostname, port)

val reqcountsByWindow = logs.
 map(line => (line.split(' ') (2), 1)).
 reduceByKeyAndWindow((x: Int, y: Int) => x+y,
 Minutes(5), Seconds(30))

val topreqsByWindow = reqcountsByWindow.
 map(pair =>
 transform(r
topreqsByWind

ssc.start()
ssc.awaitTermination()
...
```

Every 30 seconds, count requests by user over the last 5 minutes

# Count and Sort User Requests by Window

```
...
val ssc = new StreamingContext(new StreamConf(), Seconds(2))
val logs = ssc.socketTextStream(hostname, port)
...
val reqcountsByWindow = logs.
 map(line =>
 reduceByKeyAndWindow(_ + 1)
 Minutes(5))

 Sort and print the top users for every RDD (every 30
 seconds)

val topreqsByWindow=reqcountsByWindow.
 map(pair => pair.swap).
 transform(rdd => rdd.sortByKey(false))
topreqsByWindow.map(pair => pair.swap).print()

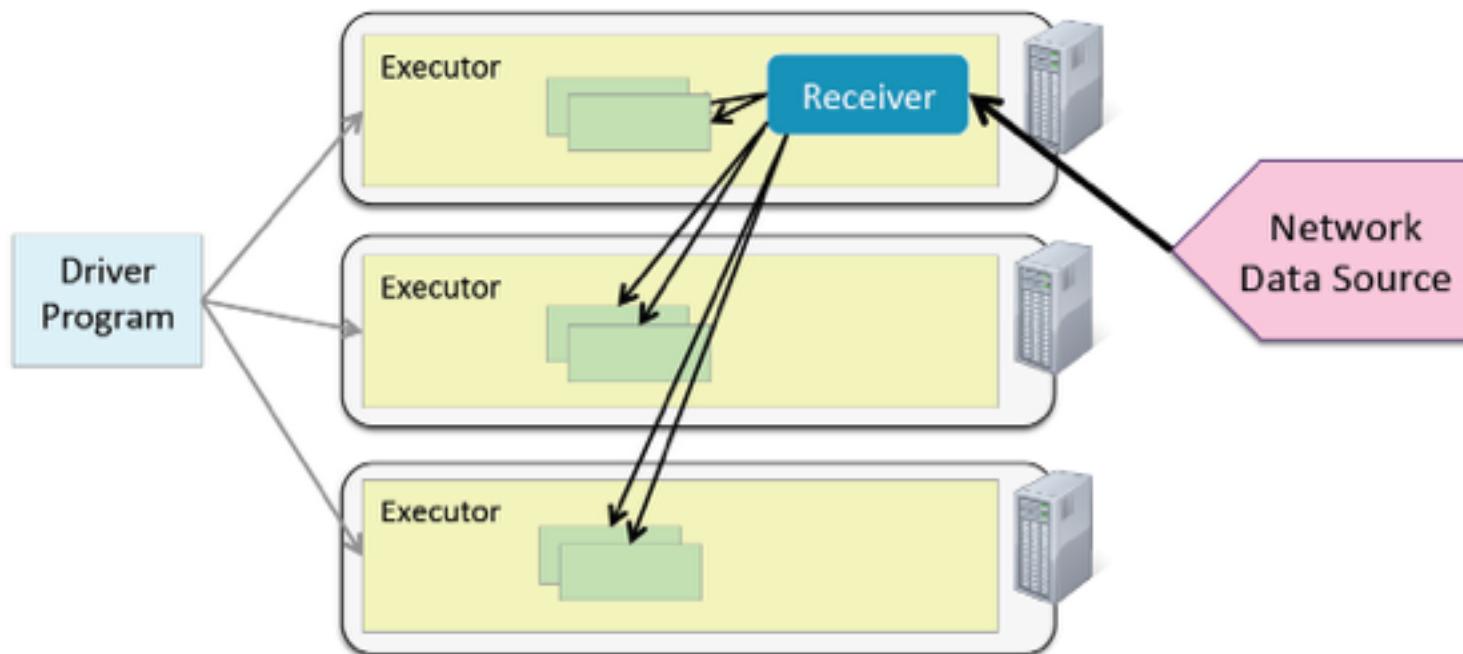
ssc.start()
ssc.awaitTermination()
...
```

# Special Considerations for Streaming Applications

- Spark Streaming applications are by definition long-running
  - Require some different approaches than typical Spark applications
- Metadata accumulates over time
  - Use checkpointing to trim RDD lineage data
    - Required to use windowed and state operations
    - Enabled by setting the checkpoint directory: `ssc.checkpoint(directory)`
- Monitoring
  - The StreamingListener API lets you collect statistics

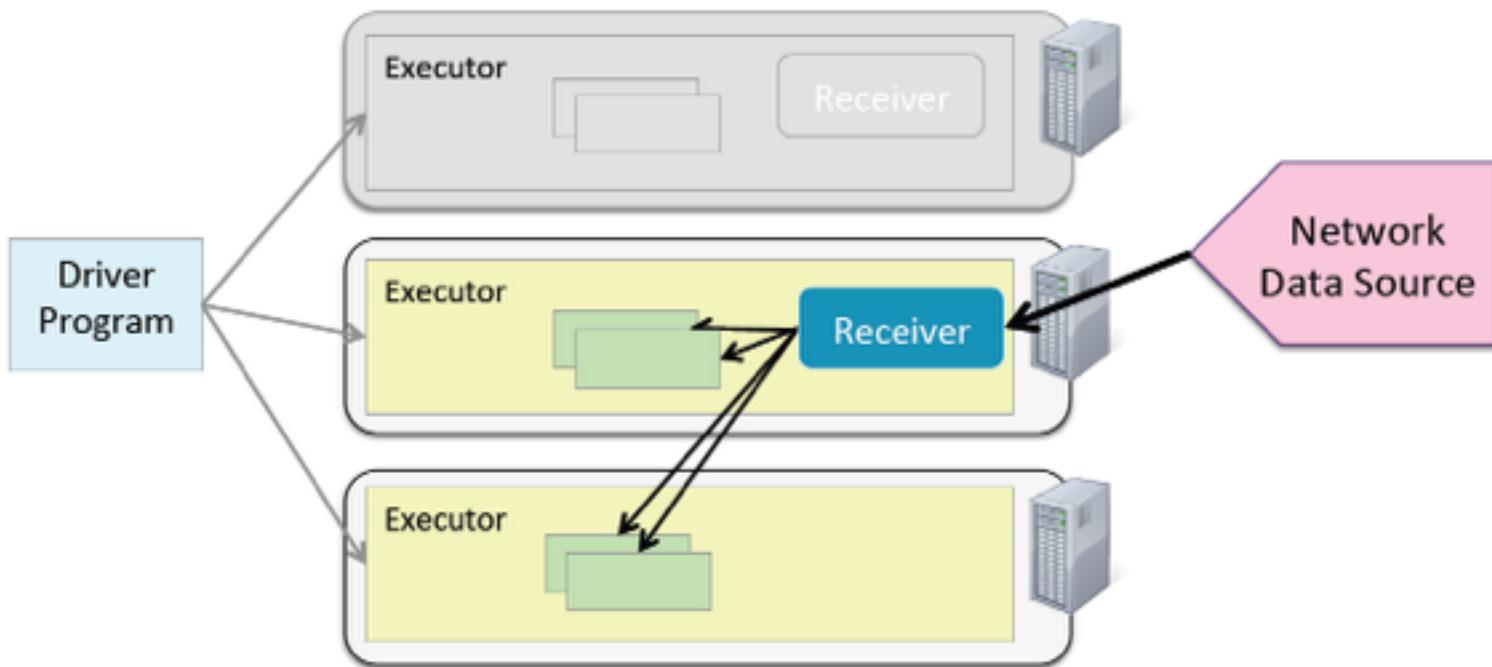
# Spark Fault Tolerance

- Network data is received on a worker node
  - Receiver distributes data (RDDs) to the cluster
- Spark Streaming persists windowed RDDs by default (replication = 2)



# Spark Fault Tolerance

- If the receiver fails, Spark will restart it on a different Executor
  - Potential for brief loss of incoming data



# Building and Running Spark Streaming Applications

- Building Spark Streaming Applications
  - Link with the main Spark Streaming library (included with Spark)
  - Link with additional Spark Streaming libraries if necessary, e.g., Kafka, Flume, Twitter
- Running Spark Streaming Applications
  - Use at least two threads if running locally

# The Spark Streaming Application UI

- The Streaming tab in the Spark App UI provides basic metrics about the application

The screenshot shows the Spark Application UI interface. At the top, there is a navigation bar with tabs: Stages, Storage, Environment, Executors, **Streaming**, and a URL placeholder. The 'Streaming' tab is highlighted with a red box. Below the tabs, the page title is 'StreamingRequestCount application UI'. The main content area is titled 'Streaming' and displays the following information:

Started at: Thu Jun 05 13:00:22 PDT 2014  
Time since start: 42 seconds 302 ms  
Network receivers: 1  
Batch interval: 2 seconds  
Processed batches: 21  
Waiting batches: 0

Statistics over last 21 processed batches

Receiver Statistics

| Receiver         | Status | Location  | Records in last batch [2014/06/05 13:01:04] | Minimum rate [records/sec] | Median rate [records/sec] | Maximum rate [records/sec] | Last Error |
|------------------|--------|-----------|---------------------------------------------|----------------------------|---------------------------|----------------------------|------------|
| SocketReceiver-0 | ACTIVE | localhost | 39                                          | 0                          | 20                        | 20                         | -          |

Batch Processing Statistics

| Metric           | Last batch | Minimum | 25th percentile | Median | 75th percentile | Maximum         |
|------------------|------------|---------|-----------------|--------|-----------------|-----------------|
| Processing Time  | 336 ms     | 215 ms  | 281 ms          | 394 ms | 525 ms          | 1 second 239 ms |
| Scheduling Delay | 3 ms       | 1 ms    | 2 ms            | 3 ms   | 5 ms            | 11 ms           |
| Total Delay      | 339 ms     | 223 ms  | 285 ms          | 396 ms | 526 ms          | 1 second 250 ms |

Hands-On Exercise: Writing a Spark Streaming Application

## HANDS-ON EXERCISE: WRITING A SPARK STREAMING APPLICATION

Chapter 10

# COMMON PATTERNS IN SPARK PROGRAMMING

# Common Spark Use Cases

- Spark is especially useful when working with any combination of:
  - Large amounts of data
    - Distributed storage
  - Intensive computations
    - Distributed computing
  - Iterative algorithms
    - In-memory processing and pipelining

# Common Spark Use Cases

- Examples
  - Risk analysis
    - “How likely is this borrower to pay back a loan?”
  - Recommendations
    - “Which products will this customer enjoy?”
  - Predictions
    - “How can we prevent service outages instead of simply reacting to them?”
  - Classification
    - “How can we tell which mail is spam and which is legitimate?”

# "PageRank"

- PageRank gives web pages a ranking score based on links from other pages
  - Higher scores given for more links, and links from other high ranking pages
- Why do we care?
  - PageRank is a classic example of big data analysis (like WordCount)
  - Lots of data- needs an algorithm that is distributable and scalable
  - Iterative - the more iterations, the better than answer

# Graph Analytics

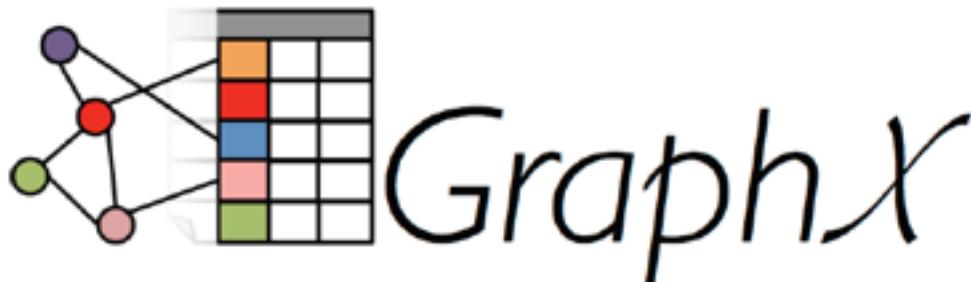
- Many data analytics problems work with “data parallel” algorithms
  - Records can be processed independently of each other
  - Very well suited to parallelizing
- Some problems focus on the relationships between the individual data items. For example:
  - Social networks
  - Web page hyperlinks
  - Roadmaps
- These relationships can be represented by graphs
  - Requires “graph parallel” algorithms

# Graph Analysis Challenges at Scale

- Graph Creation
  - Extracting relationship information from a data source
    - For example, extracting links from web pages
- Graph Representation
  - e.g., adjacency lists in a table
- Graph Analysis
  - Inherently iterative, hard to parallelize
  - This is the focus of specialized libraries like Pregel, GraphLab
- Post analysis processing
  - e.g., incorporating product recommendations into a retail site

# Graph Analysis in Spark

- Spark is very well suited to graph parallel algorithms
- GraphX
  - UC Berkeley AMPLab project on top of Spark
  - Unifies optimized graph computation with Spark's fast data parallelism and interactive abilities



# Machine Learning

- Most programs tell computers exactly what to do
  - Database transactions and queries
  - Controllers
    - Phone systems, manufacturing processes, transport, weaponry, etc.
  - Media delivery
  - Simple search
  - Social systems
    - Chat, blogs, email, etc.
- An alternative technique is to have computers learn what to do
- Machine Learning refers to programs that leverage collected data to drive future program behavior
- This represents another major opportunity to gain value from data

# The ‘Three Cs’

- Machine Learning is an active area of research and new applications
- There are three well-established categories of techniques for exploiting data:
  - Collaborative filtering (recommendations)
  - Clustering
  - Classification

# Collaborative Filtering

- Collaborative Filtering is a technique for recommendations
- Example application: given people who each like certain books, learn to suggest what someone may like in the future based on what they already like
- Helps users navigate data by expanding to topics that have affinity with their established interests
- Collaborative Filtering algorithms are agnostic to the different types of data items involved
  - Useful in many different domains

# Clustering

- Clustering algorithms discover structure in collections of data
  - Where no formal structure previously existed
- They discover what clusters, or groupings, naturally occur in data
- Examples:
  - Finding related news articles
  - Computer vision (groups of pixels that cohere into objects)

# Classification

- The previous two techniques are considered ‘unsupervised’ learning
  - The algorithm discovers groups or recommendations itself
- Classification is a form of ‘supervised’ learning
- A classification system takes a set of data records with known labels
  - Learns how to label new records based on that information
- Example:
  - Given a set of e-mails identified as spam/not spam, label new e-mails as spam/not spam
  - Given tumors identified as benign or malignant, classify new tumors

# Machine Learning Challenges

- Highly computation intensive and iterative
- Many traditional numerical processing systems do not scale to very large datasets
  - e.g., MatLab

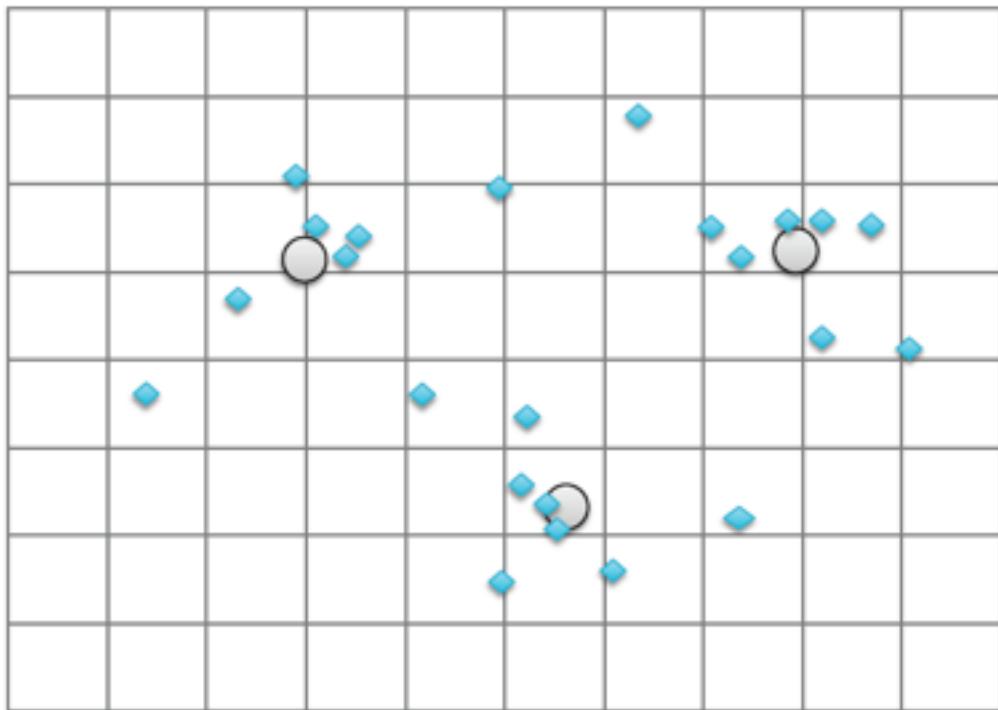
# MLlib: Machine Learning on Spark

- MLlib is part of Apache Spark
- Includes many common ML functions
  - ALS (alternating least squares)
  - k-means
  - Logistic Regression
  - Linear Regression
  - Gradient Descent
- Still a ‘work in progress’

# K-means Clustering

- K-means Clustering
  - A common iterative algorithm used in graph analysis and machine learning
  - You will implement a simplified version in the Hands-On Exercises

# Approximate k-means Clustering



1. Choose K random points as starting centers
2. Find all points closest to each center
3. Find the center (mean) of each cluster
4. If the centers changed by more than  $c$ , iterate again  
...
5. Close enough!

## Hands-On Exercise: Iterative Processing in Spark

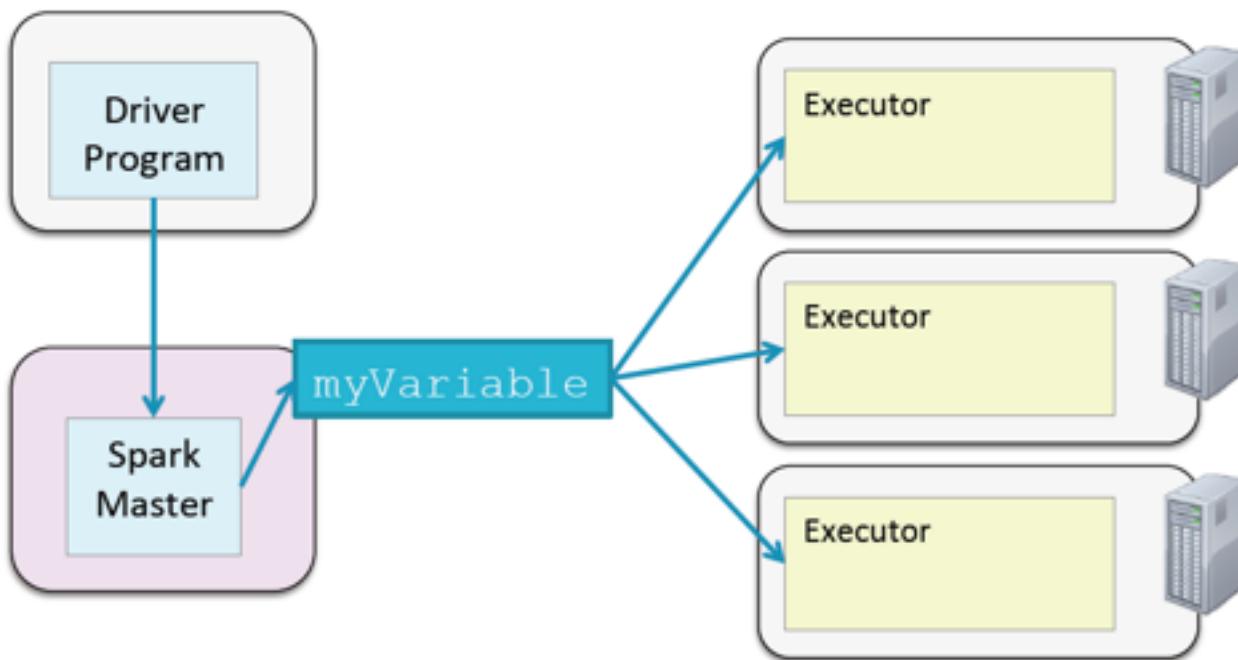
# HANDS-ON EXERCISE

Chapter 11

# IMPROVING SPARK PERFORMANCE

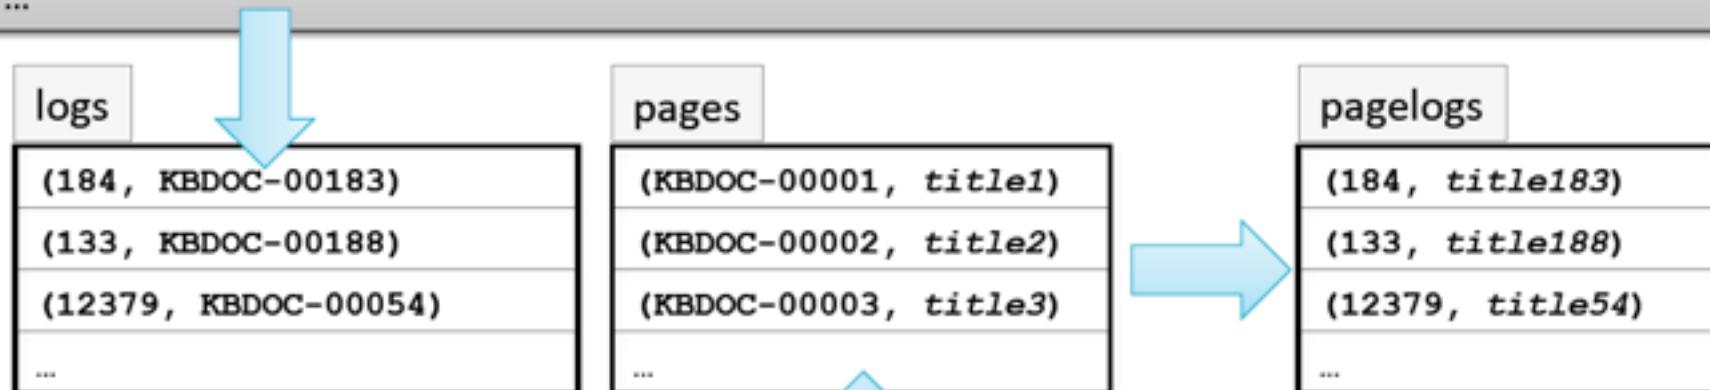
# Broadcast Variables

- Broadcast variables are set by the driver and retrieved by the workers
- They are read only once set
- The first read of a Broadcast variable retrieves and stores its value on the node



# Match User IDs with Requested Page Titles

```
227.35.151.122 - 184 [16/Sep/2013:00:03:51 +0100] "GET /KBDOC-00183.html HTTP/1.0" 200 ...
146.218.191.254 - 133 [16/Sep/2013:00:03:48 +0100] "GET /KBDOC-00188.html HTTP/1.0" 200 ...
176.96.251.224 - 12379 [16/Sep/2013:00:02:29 +0100] "GET /KBDOC-00054.html HTTP/1.0" 16011...
...
```



KBDOC-00001:*MeeToo 4.1 - Back up files*

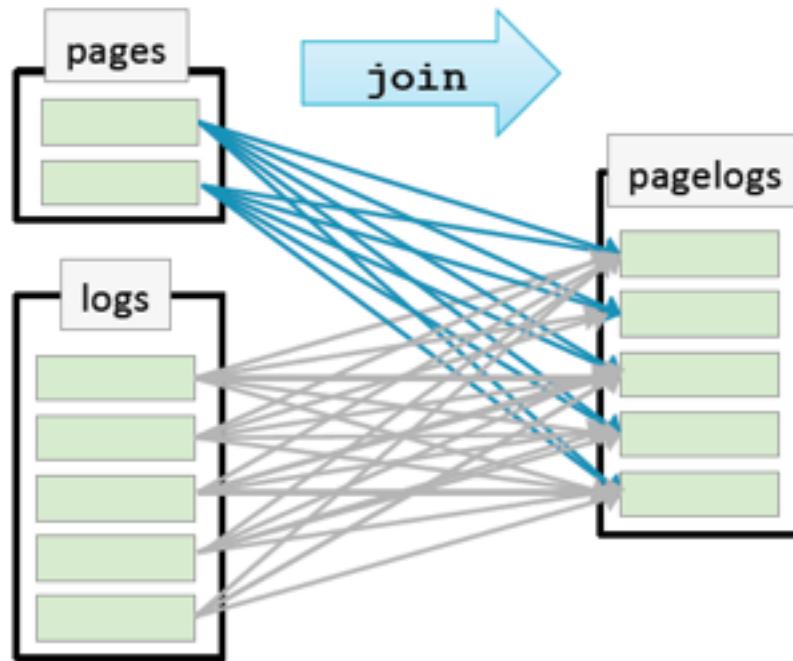
KBDOC-00002:Sorrento F24L - Change the phone ringtone and notification sound

KBDOC-00003:Sorrento F41L – overheating

...

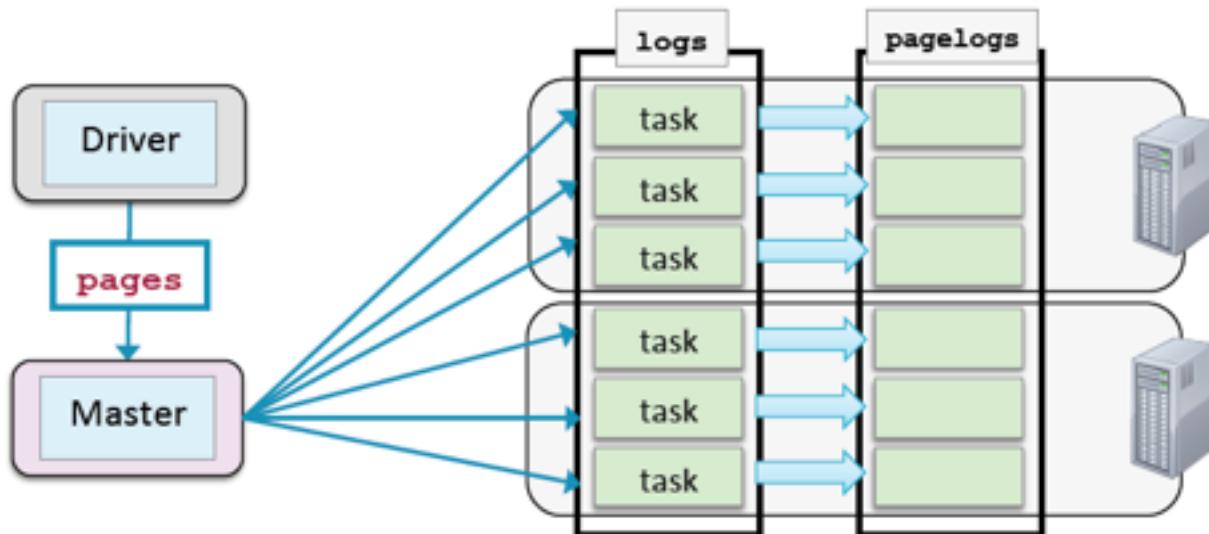
# Join a Web Server Log with Page Titles

```
logs = sc.textFile(logfile).map(fn)
pages = sc.textFile(pagefile).map(fn)
pagelogs = logs.join(pages)
```



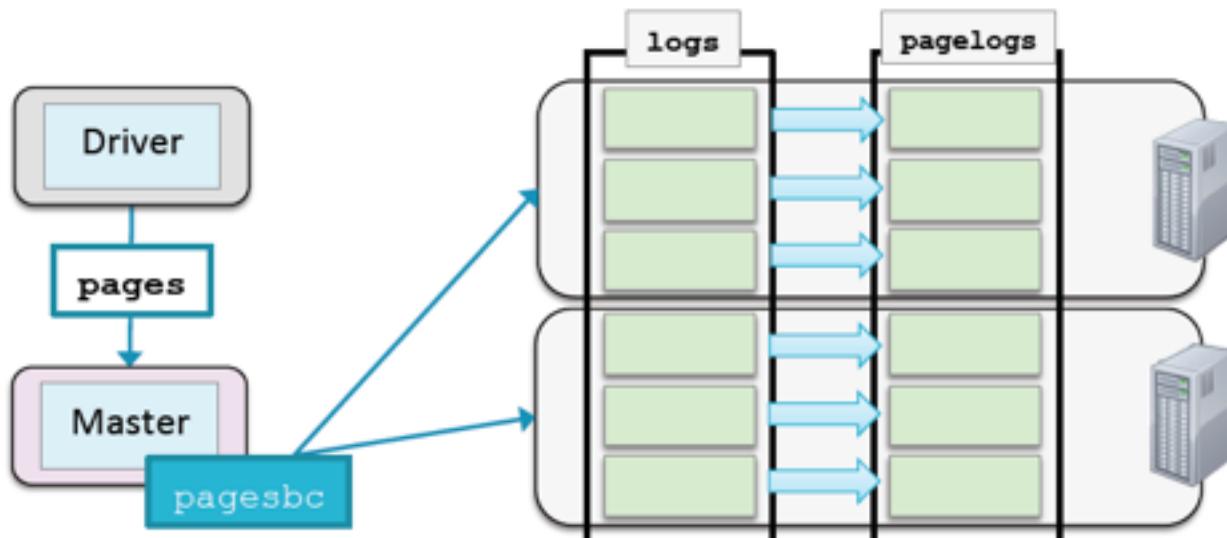
# Pass a Small Table as a Parameter

```
logs = sc.textFile(logfile).map(fn)
pages = dict(map(fn,open(pagefile)))
pagelogs = logs.map(lambda (userid,pageid):
 (userid,pages[pageid]))
```



# Broadcast a Small Table

```
logs = sc.textFile(logfile).map(...)
pages = dict(map(fn,open(pagefile)))
pagesbc = sc.broadcast(pages)
pagelogs = logs.map(lambda (userid, pageid):
 (userid,pagesbc.value[pageid])))
```



# Broadcast Variables

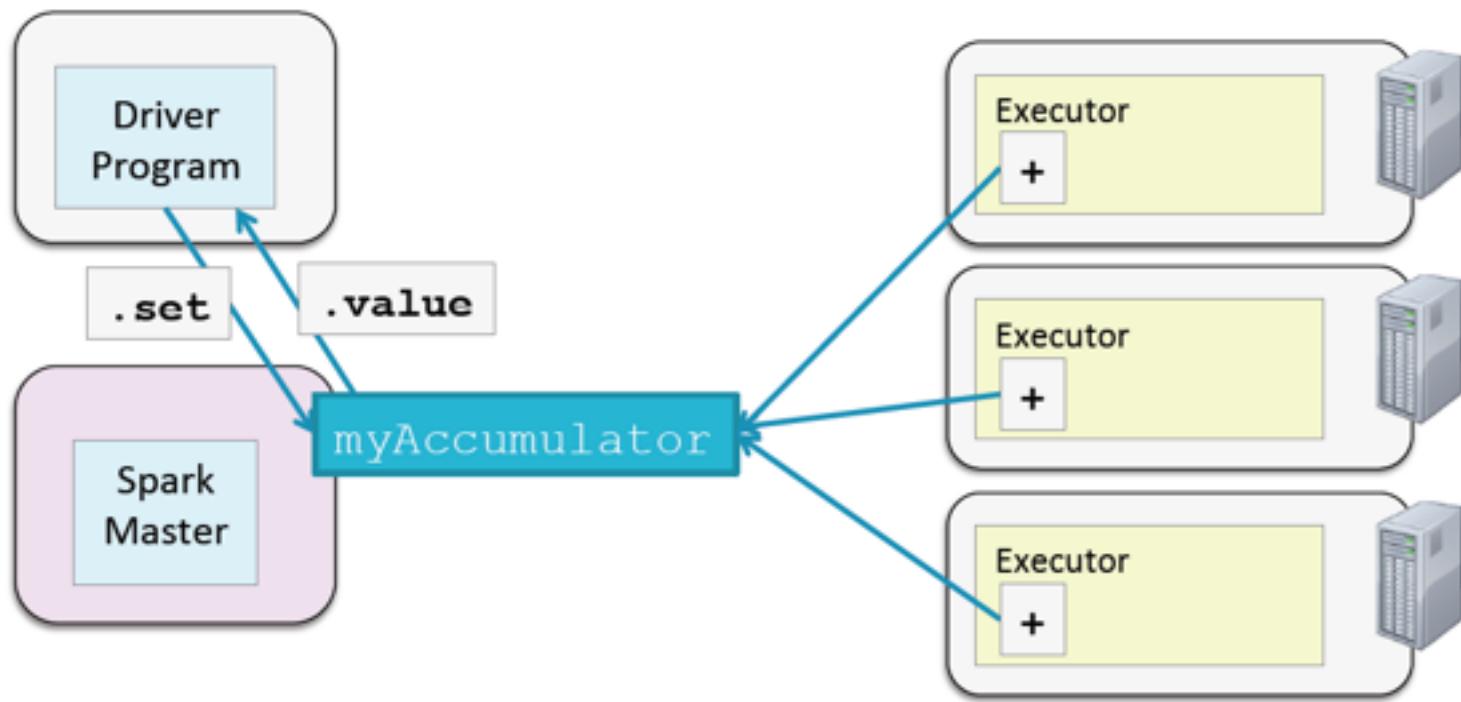
- Why use Broadcast variables?
  - Use to minimize transfer of data over the network, which is usually the biggest bottleneck
  - Spark Broadcast variables are distributed to worker nodes using a very efficient peer-to-peer algorithm

In this Hands-On Exercise, you will filter web server logs for requests

## HANDS-ON EXERCISE: USING BROADCAST VARIABLES

# Accumulators

- Accumulators are shared variables
  - Worker nodes can add to the value
  - Only the driver application can access the value



# Accumulator Example: Average Word Length

```
def addTotals(word,words,letters):
 words += 1
 letters += len(word)

totalWords = sc.accumulator(0)
totalLetters = sc.accumulator(0.0)

words = sc.textFile(myfile) \
 .flatMap(lambda line: line.split())

words.foreach(lambda word: \
 addTotals(word,totalWords,totalLetters))

print "Average word length: ", \
 totalLetters.value/totalWords.value
```

# More About Accumulators

- Accumulators will only be incremented once per task
  - If tasks must be rerun due to failure, Spark will correctly add only for the task which succeeds
- Only the driver can access the value
  - Updates are only sent to the master, not to all workers
    - Code will throw an exception if you use `.value` on worker nodes
- Supports the increment (`+=`) operator
- Can use integers or doubles
  - `sc. accumulator(0)`
  - `sc. accumulator(0.0)`
- Can customize to support any data type
  - Extend the Accumulator Param class

Hands-On Exercise: Using Accumulators

## HANDS-ON EXERCISE: USING ACCUMULATORS

# Performance Issue: Serialization

- **Serialization affects**
  - Network bandwidth
  - - Memory (save memory by serializing)
- **Default method of serialization in Spark is basic Java serialization**
  - - Simple but slow

# Using Kryo Serialization

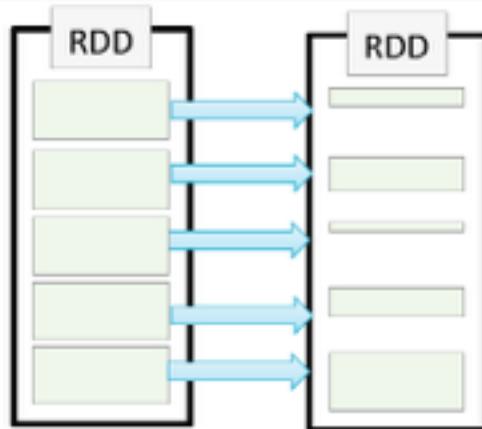
- Use Kryo serialization for Scala and Java
  - To enable, set `spark.serializer = spark.KryoSerializer`
- To enable Kryo for your custom classes
  - Create a `KryoRegistrar` class and set `spark.kryo.registrator=MyRegistrar`
  - Register your classes with Kryo

```
class MyRegistrar extends spark.KryoRegistrar {
 def registerClasses(kryo: Kryo) {
 kryo.register(classOf[MyClass1])
 kryo.register(classOf[MyClass2])
 ...
 }
}
```

# Performance Issue: Small Partitions

- Problem: filter() can result in partitions with small amounts of data
  - Results in many small tasks

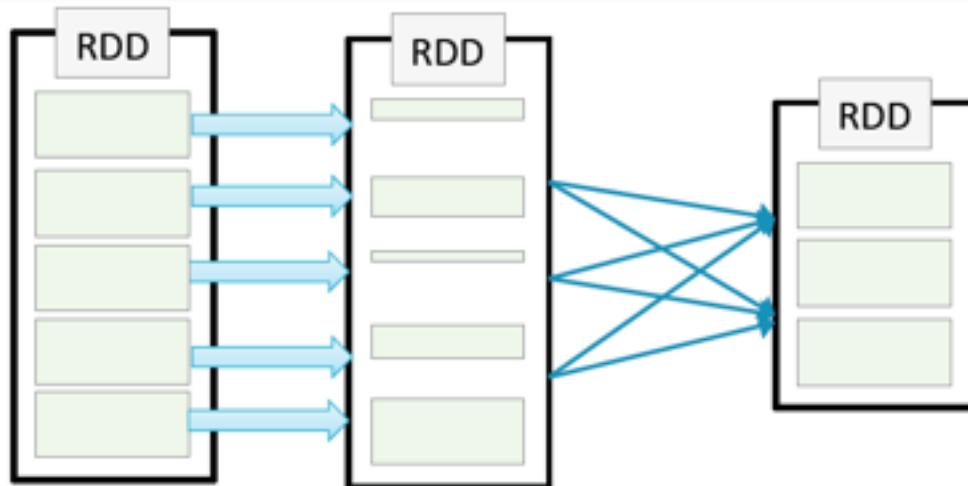
```
sc.textFile(file) \
 .filter(lambda s: s.startswith('I')) \
 .map(lambda s: \
 (s.split()[0],(s.split()[1],s.split()[2])))
```



# Solution: Repartition/Coalesce

- Solution: repartition(n)
  - This is the same as coalesce(n, shuffle=false)

```
sc.textFile(file) \
 .filter(lambda s: s.startswith('I')) \
 .repartition(3) \
 .map(lambda s: \
 (s.split()[0], (s.split()[1],s.split()[2])))
```



# Performance Issue: Passing Too Much Data in Functions

- Problem: Passing large amounts of data to parallel functions results in poor performance

```
hashmap = some_massive_hash_map()
...
myrdd.map(lambda x: hashmap(x)).countByValue()
```

# Performance Issues: Passing Too Much Data in Functions

- Solution:
  - If the data is relatively small, use a Broadcast variable

```
hashmap = some_massive_hash_map()
bhashmap = sc.broadcast(hashmap)
...
myrdd.map(lambda x: bhashmap(x)).countByValue()
```

- If the data is very large, parallelize into an RDD

```
hashmap = some_massive_hash_map()
hashmaprdd = sc.parallelize(hashmap)
...
myrdd.join(bhashmaprdd).countByValue()
```

# Diagnosing Performance Issues

- The Spark Application UI provides useful metrics to find performance problems

Completed Stages (3)

| Stage Id | Description                    | Submitted           | Duration | Tasks: Succeeded/Total | Shuffle Read | Shuffle Write |
|----------|--------------------------------|---------------------|----------|------------------------|--------------|---------------|
| 2        | saveAsTextFile at <console>:30 | 2014/06/06 06:50:41 | 4 s      | 61/61                  |              |               |
| 0        | top at <console>:30            |                     |          |                        |              |               |
| 1        | reduceByKey at <console>:29    |                     |          |                        |              |               |

**Spark** Stages Storage Environment Executors Spark shell application UI

**Stage Details** (Curved arrow pointing from the Stage Details section to the Stage ID column in the table)

**Details for Stage 4**

Total task time across all tasks: 0.4 s

**Summary Metrics for 61 Completed Tasks**

| Metric                           | Min  | 25th percentile | Median | 75th percentile | Max   |
|----------------------------------|------|-----------------|--------|-----------------|-------|
| Result serialization time        | 0 ms | 0 ms            | 0 ms   | 0 ms            | 4 ms  |
| Duration                         | 0 ms | 1 ms            | 3 ms   | 10 ms           | 34 ms |
| Time spent fetching task results | 0 ms | 0 ms            | 0 ms   | 0 ms            | 0 ms  |
| Scheduler delay                  | 8 ms | 11 ms           | 12 ms  | 14 ms           | 26 ms |

**Aggregated Metrics by Executor**

| Executor ID | Address         | Task Time | Total Tasks | Failed Tasks | Succeeded Tasks | Shuffle Read | Shuffle Write | Shuffle Spill (Memory) | Shuffle Spill (Disk) |
|-------------|-----------------|-----------|-------------|--------------|-----------------|--------------|---------------|------------------------|----------------------|
| 0           | localhost:40410 | 1 s       | 61          | 0            | 61              | 0.0 B        | 0.0 B         | 0.0 B                  | 0.0 B                |

# Diagnosing Performance Issues

- Where to look for performance issues
  - Scheduling and launching tasks
  - Task execution
  - Shuffling
  - Collecting data

# Scheduling and Launching Issues

- Scheduling and launching taking too long
  - Are you passing too much data to tasks?
    - `myrdd.map(lambda x: HugeLookupTable(x))`
  - Use a Broadcast variable or an RDD

Summary Metrics for 11 Completed Tasks

| Metric                           | Min   | 25th percentile | Median | 75th percentile | Max   |
|----------------------------------|-------|-----------------|--------|-----------------|-------|
| Result serialization time        | 0 ms  | 0 ms            | 0 ms   | 0 ms            | 4 ms  |
| Duration                         | 20 ms | 23 ms           | 30 ms  | 44 ms           | 0.3 s |
| Time spent fetching task results | 0 ms  | 0 ms            | 0 ms   | 0 ms            | 0 ms  |
| Scheduler delay                  | 11 s  | 11 s            | 12 s   | 12 s            | 15 s  |

# Task Execution Issues

- Task execution taking too long?
  - Are there tasks with a very high per-record overhead?
    - e.g., `mydata.map(dbLookup)`
    - Each lookup call opens a connection to the DB, reads, and closes
  - Try `mapPartitions`

# Task Execution Issues

- Are a few tasks taking much more time than others?
  - Repartition, partition on a different key, or write a custom partitioner

| Summary Metrics for 182 Completed Tasks |       |                 |        |                 |       |
|-----------------------------------------|-------|-----------------|--------|-----------------|-------|
| Metric                                  | Min   | 25th percentile | Median | 75th percentile | Max   |
| Result serialization time               | 0 ms  | 0 ms            | 0 ms   | 0 ms            | 1 ms  |
| Duration                                | 11 ms | 15 ms           | 17 ms  | 20 ms           | 50 ms |
| Time spent fetching                     | 0 ms  | 0 ms            | 0 ms   | 0 ms            | 0 ms  |

Task durations should be fairly even

Example: empty partitions due to filtering

| 182 Partitions |                                   |                  |              |                 |              |
|----------------|-----------------------------------|------------------|--------------|-----------------|--------------|
| Block Name     | Storage Level                     | Size in Memory ▲ | Size on Disk | Executors       | Partition ID |
| rdd_8_87       | Memory Deserialized 1x Replicated | 2.2 MB           | 0.0 B        | localhost:56859 | 0            |
| rdd_8_0        | Memory Deserialized 1x Replicated | 112.0 B          | 0.0 B        | localhost:56859 | 1            |
| rdd_8_1        | Memory Deserialized 1x Replicated | 112.0 B          | 0.0 B        | localhost:56859 | 2            |
| rdd_8_10       | Memory Deserialized 1x Replicated | 112.0 B          | 0.0 B        | localhost:56859 | 3            |
| rdd_8_100      | Memory Deserialized 1x Replicated | 112.0 B          | 0.0 B        | localhost:56859 | 4            |
| rdd_8_104      | Memory Deserialized 1x Replicated | 112.0 B          | 0.0 B        | localhost:56859 | 5            |

# Shuffle Issues

- Writing shuffle results taking too long?
  - Make sure you have enough memory for buffer cache
  - Make sure spark.local.dir is a local disk, ideally dedicated

| Completed Stages (3) |                                |                     |
|----------------------|--------------------------------|---------------------|
| Stage Id             | Description                    | Submitted           |
| 2                    | saveAsTextFile at <console>:30 | 2014/06/06 06:50:41 |
| 0                    | top at <console>:30            | 2014/06/06 06:49:56 |
| 1                    | reduceByKey at <console>:29    | 2014/06/06 06:49:47 |

Saves to disk if too big for buffer cache

| File Read | Shuffle Write |
|-----------|---------------|
|           | 168.9 KB      |

| Tasks      |         |         |                |           |                     |          |      |                 |              |               |        |
|------------|---------|---------|----------------|-----------|---------------------|----------|------|-----------------|--------------|---------------|--------|
| Task Index | Task ID | Status  | Locality Level | Executor  | Time                | Duration | Time | Result Ser Time | Write Time ▲ | Shuffle Write | Errors |
| 26         | 26      | SUCCESS | PROCESS_LOCAL  | localhost | 2014/06/06 05:11:44 | 66 ms    | 4 ms |                 | 9 ms         | 2.4 KB        |        |
| 21         | 21      | SUCCESS | PROCESS_LOCAL  | localhost | 2014/06/06 05:11:44 | 74 ms    |      |                 | 8 ms         | 2.4 KB        |        |
| 22         | 22      | SUCCESS | PROCESS_LOCAL  | localhost | 2014/06/06 05:11:44 | 0.1 s    |      |                 | 6 ms         | 2.4 KB        |        |

Look for big write times

# Collecting Data to the Driver

- Are results taking too long?

- Beware of returning large amounts of data to the driver, for example with `collect()`
- Process data on the workers, not the driver
- Save large results to HDFS

Watch for  
disproportionate result  
serialization times

| Tasks      |         |         |                |           |                     |          |         |                   |        |
|------------|---------|---------|----------------|-----------|---------------------|----------|---------|-------------------|--------|
| Task Index | Task ID | Status  | Locality Level | Executor  | Launch Time         | Duration | GC Time | Result Ser Time ▲ | Errors |
| 153        | 215     | SUCCESS | PROCESS_LOCAL  | localhost | 2014/06/06 08:17:23 | 27 ms    | 85 ms   | 94 ms             |        |
| 87         | 149     | SUCCESS | PROCESS_LOCAL  | localhost | 2014/06/06 08:17:21 | 35 ms    | 72 ms   | 87 ms             |        |
| 116        | 178     | SUCCESS | PROCESS_LOCAL  | localhost | 2014/06/06 08:17:22 | 44 ms    | 54 ms   | 77 ms             |        |

# Performance Analysis and Monitoring

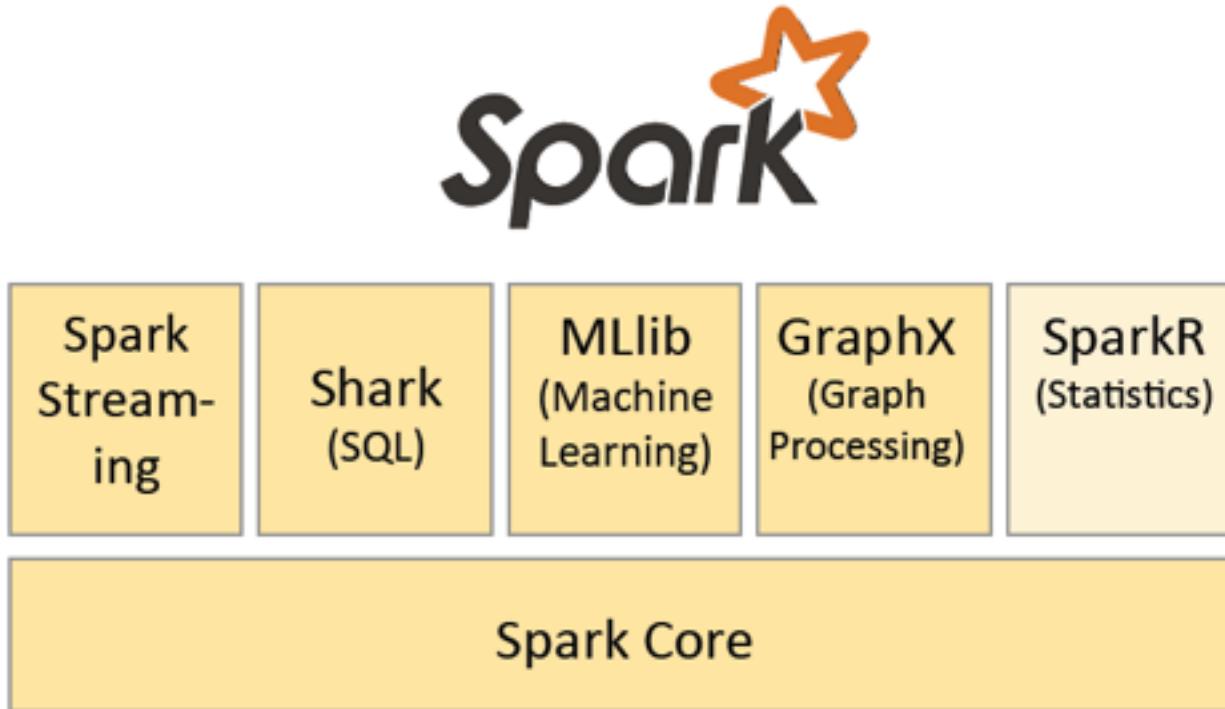
- Spark supports integration with other performance tools
  - Configurable metrics system built on the Coda Hale Metrics Library
  - Metrics can be
    - Saved to files
    - Output to the console
    - Viewed in the JMX console
    - Sent to reporting tools like Graphite or Ganglia

Chapter 12

# SPARK, HADOOP, AND THE ENTERPRISE DATA CENTER

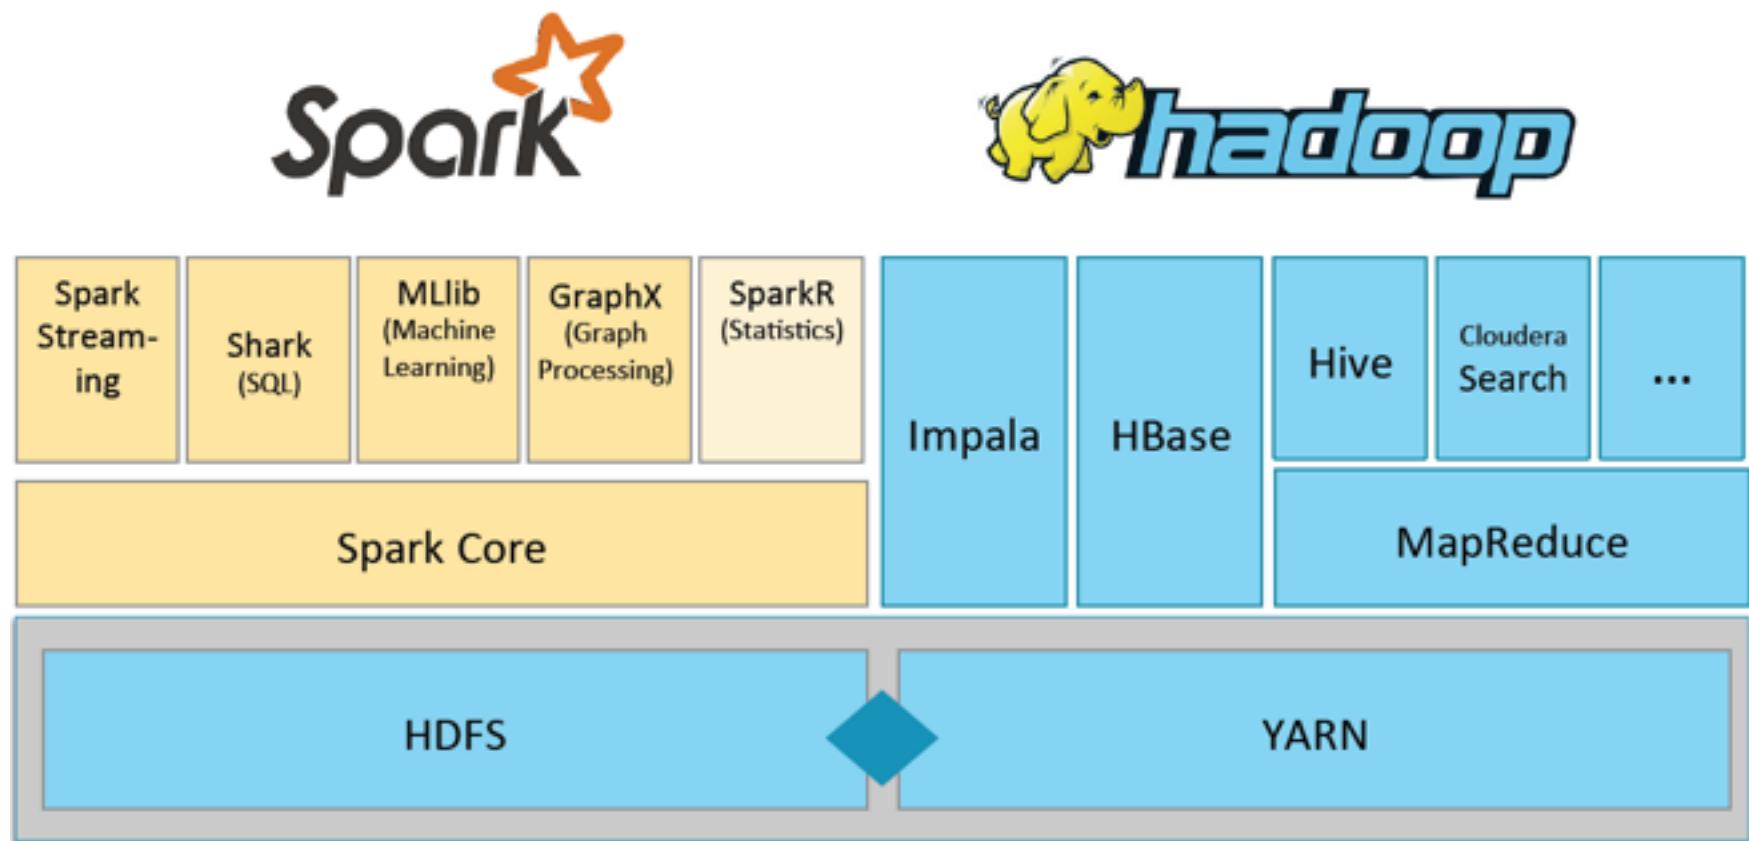
# The Spark Stack

- In addition to the core Spark engine, there are an ever-growing number of related projects
- Sometimes called the Berkeley Data Analytics Stack (BDAS)



# Spark and Hadoop

- Spark was created to complement, not replace, Hadoop

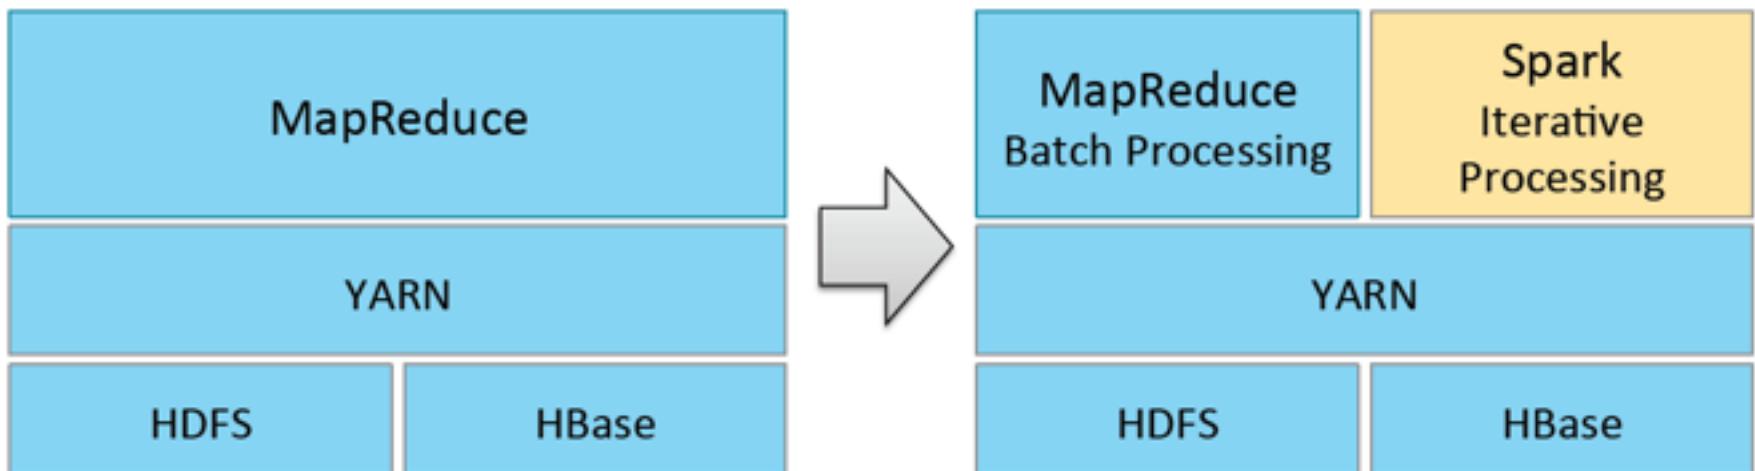


# Spark and Hadoop (2)

- Spark uses HDFS
  - Can use any Hadoop data source
  - Uses Hadoop InputFormats and OutputFormats
    - This means it can manipulate e.g., Avro files and SequenceFiles
- Spark runs on YARN
  - Can run on the same cluster with MapReduce jobs, Impala, etc.
- Spark works with the Hadoop ecosystem
  - Flume
  - Sqoop
  - HBase

# Yahoo

- Example use-case: Yahoo is a major user of Hadoop
  - Uses Hadoop for personalization, collaborative filtering, ad analytics...
- MapReduce couldn't keep up
  - Highly iterative machine learning algorithms
- Moved iterative processing to Spark

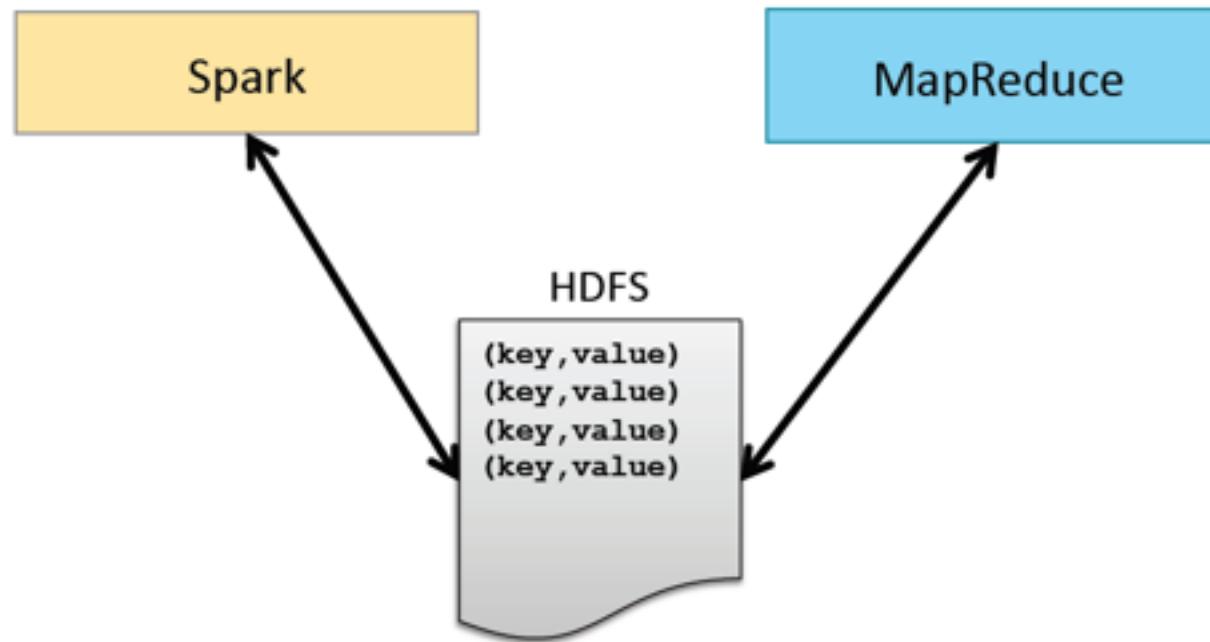


# Spark vs Hadoop MapReduce

- Hadoop MapReduce
  - Widely used, huge investment already made
  - Supports and supported by many complementary tools
  - Mature, stable, well-tested technology
  - Skilled developers available
- Spark
  - Flexible
  - Elegant
  - Fast
  - Changing rapidly

# Sharing Data Between Spark and MapReduce Jobs

- Apache Avro is a binary file format for saving datasets
- Hadoop SequenceFiles are similar; used by many existing Hadoop data centers
- Both are supported by Spark (Scala now, Python soon)



# The Hadoop Ecosystem

- In addition to HDFS and MapReduce, the Hadoop Ecosystem includes many additional components
- Some that may be of particular interest to Spark developers
  - Data Storage: HBase
  - Data Analysis: Hive and Impala
  - Data Integration: Flume and Sqoop

# Data Storage: Hbase - The Hadoop Database

- HBase: database layered on top of HDFS
  - Provides interactive access to data
- Stores massive amounts of data
  - Petabytes+
- High throughput
  - Thousands of writes per second (per node)
- Handles sparse data well
  - No wasted space for a row with empty columns
- Limited access model
  - Optimized for lookup of a row by key rather than full queries
    - No transactions: single row operations only



# Data Analysis: Hive



- **What is Hive?**

- Open source Apache project
- Built on Hadoop MapReduce
- **HiveQL:** An SQL-like interface to Hadoop

```
SELECT * FROM purchases WHERE price > 10000 ORDER BY
storeid
```

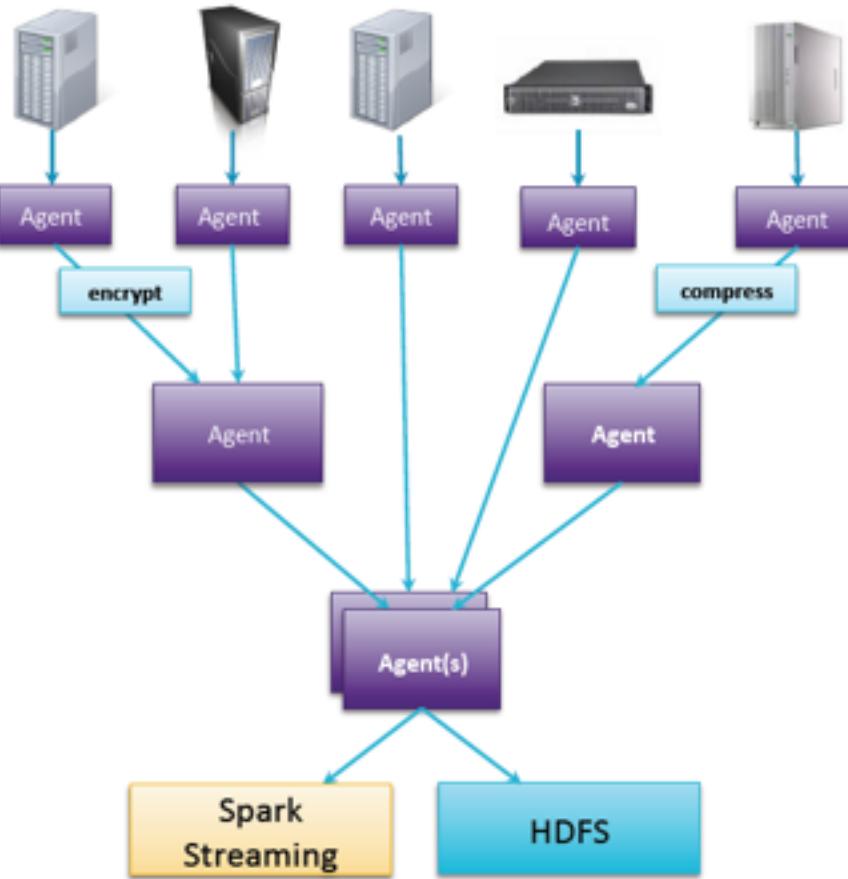
# Data Analysis: Flume

- **What is Flume?**
  - A service to move large amounts of data in real time
  - Example: storing log files in HDFS
- **Flume is**
  - Distributed
  - Reliable and available
  - Horizontally scalable
  - Extensible
- **Spark Streaming is integrated with Flume**



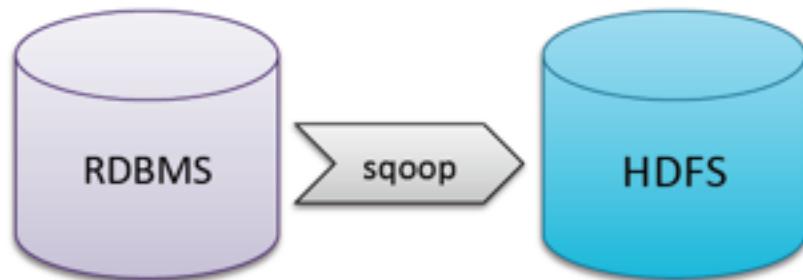
# Data Analysis: Flume

- Collect data as it is produced
  - Files, syslogs, stdout or custom source
- Process in place
  - e.g., encrypt, compress
- Pre-process data before storing
  - e.g., transform, scrub, enrich
- Write in parallel
  - Scalable throughput
- Store in any format
  - Text, compressed, binary, or custom sink



# Data Integration: Sqoop - SQL to Hadoop

- Typical scenario: data stored in an RDBMS is needed in a Spark application
  - Lookup tables
- Possible to read directly from an RDBMS in your Spark application
  - Can lead to the equivalent of a distributed denial of service (DDoS) attack on your RDBMS
  - In practice – don't do it!
- Better idea: use Sqoop to import the data into HDFS beforehand



# Data Integration: Squeryl - SQL to Hadoop



- Squeryl: open source tool
  - Now a top-level Apache Software Foundation project
- Imports tables from an RDBMS into HDFS
  - Just one table, all tables, or portions of a table
  - Uses MapReduce to actually import the data
- Uses a JDBC interface
  - Works with virtually any JDBC-compatable database
- Imports data to HDFS as delimited text files or SequenceFiles
  - Default is comma-delimited text files
- Can be used for incremental data imports
  - First import retrieves all rows in a table
  - Subsequent imports retrieve just rows created since the last import

# Sqoop: Basic Syntax

- Standard syntax:

```
$ sqoop tool-name [tool-options]
```

- Tools include:

- `import`
  - `import-all-tables`
  - `list-tables`

- Options include:

- `--connect`
  - `--username`
  - `--password`

# Sqoop: Example

- Example: import a table called employees from a database called

```
$ sqoop import --username fred --password derf \
 --connect jdbc:mysql://database.example.com/personnel \
 --table employees
```

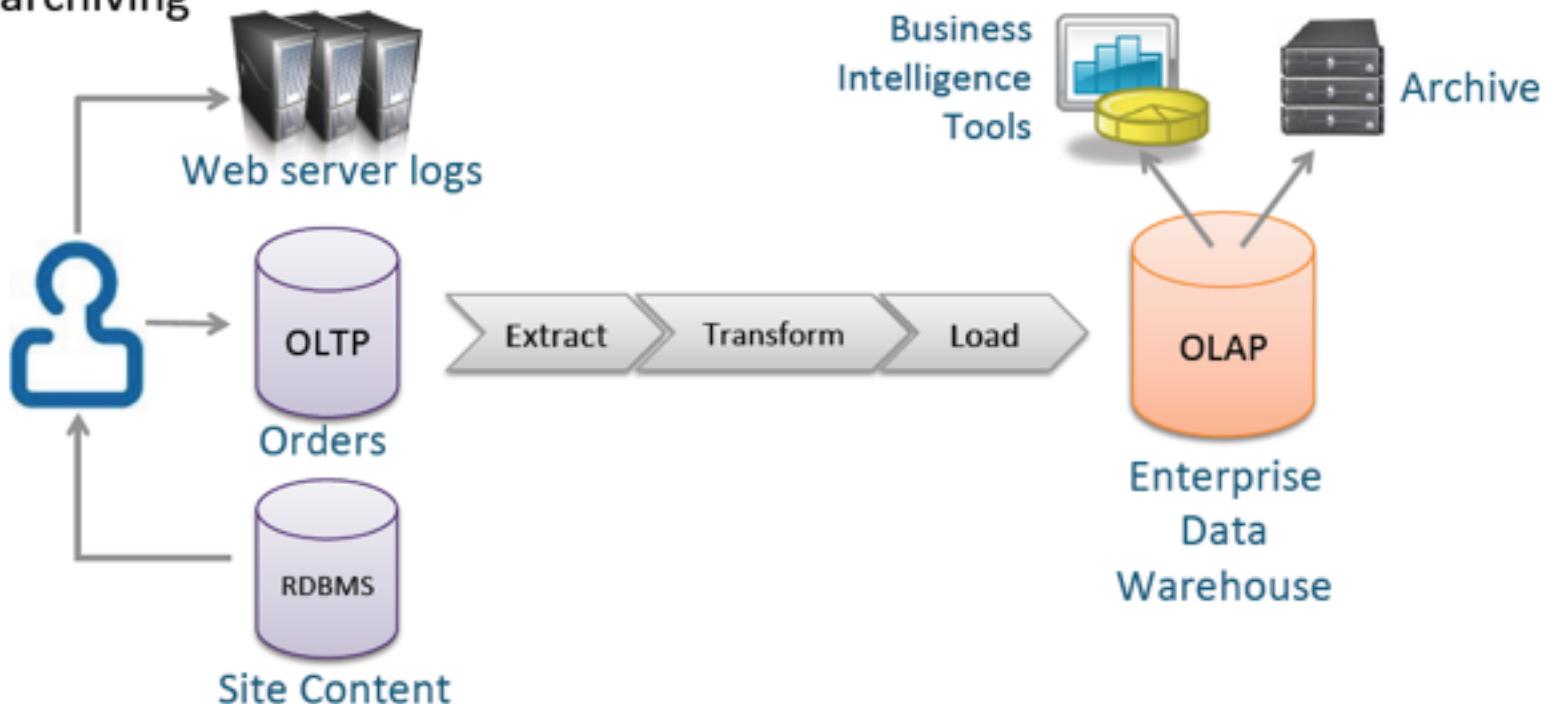
- Example: as above, but only records with an ID greater than 1000

```
$ sqoop import --username fred --password derf \
 --connect jdbc:mysql://database.example.com/personnel \
 --table employees \
 --where "id > 1000"
```

# Typical RDBMS Scenario

- **Typical scenario:**

- Interactive RDBMS serves queries from a web site
- Data is extracted and loaded into a data warehouse for processing and archiving

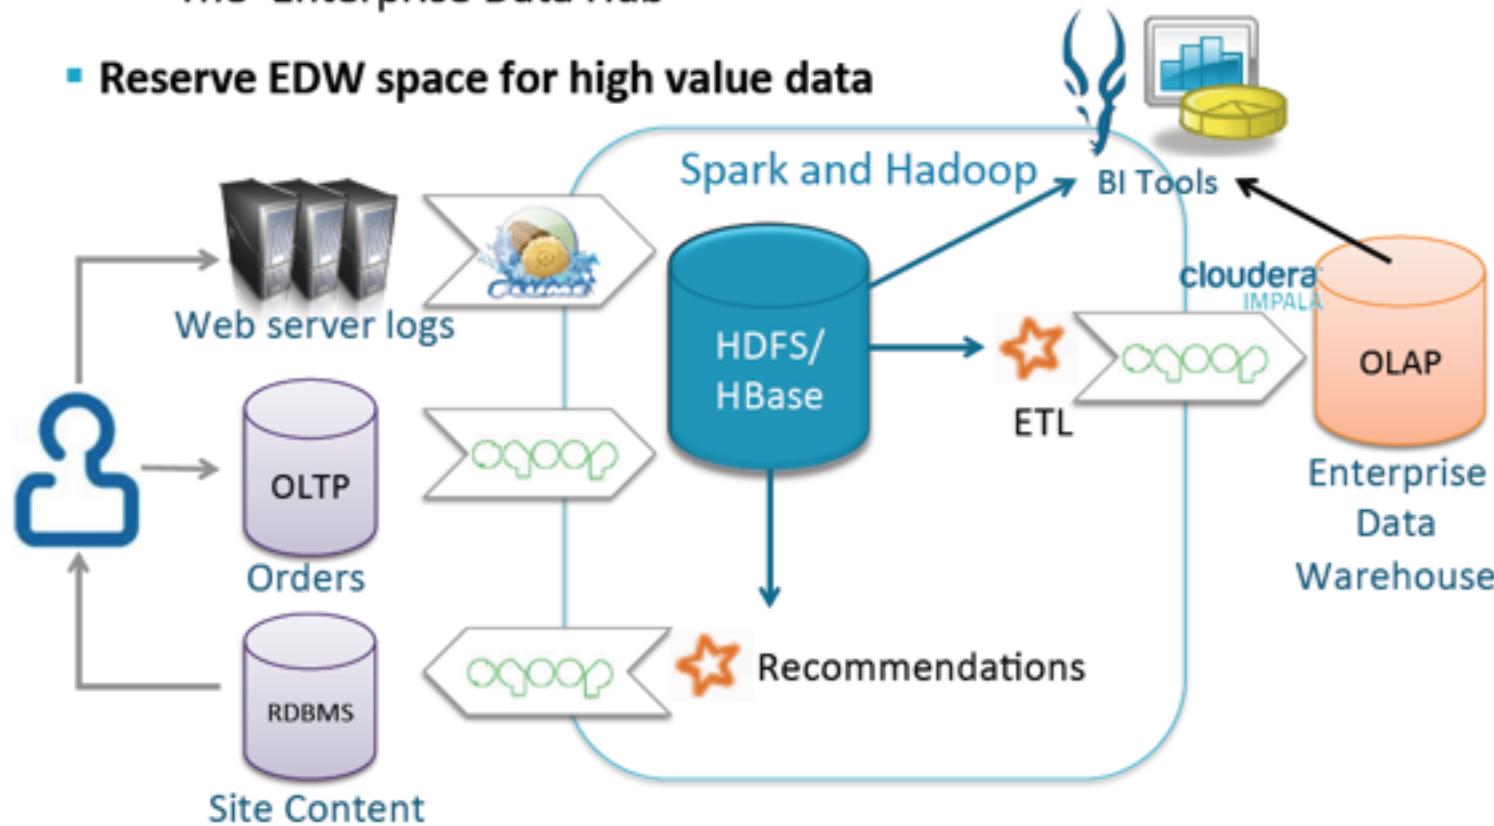


OLTP: Online Transaction Processing

OLAP: Online Analytical Processing

# Using Spark and Hadoop to Augment Existing Databases

- With Spark and Hadoop you can store and process *all* your data
  - The ‘Enterprise Data Hub’
- Reserve EDW space for high value data



# Benefits of Spark and Hadoop Over RDBMSs

- Processing power scales with data storage
  - As you add more nodes for storage, you get more processing power ‘for free’
- Views do not need prematerialization
  - Ad-hoc full or partial dataset queries are possible
- Total query size can be multiple petabytes

# Traditional High-Performance File Servers

- Enterprise data is often held on large fileservers, such as products from
  - NetApp
  - EMC
- Advantages
  - Fast random access
  - Many concurrent clients
- Disadvantages
  - High cost per terabyte of storage

# File Servers and HDFS

- Choice of storage depends on the expected access patterns
  - Sequentially read, append-only data: HDFS
  - Random access: file server
- HDFS can crunch sequential data faster
- Offloading data to HDFS leaves more room on file servers for ‘interactive’ data
- Use the right tool for the job!

Hands-On Exercise: Importing RDBMS Data Into Spark

## HANDS-ON EXERCISE: IMPORTING RDBMS DATA INTO SPARK

# THANK YOU!